

e__commerce

August 30, 2026

1 1. Problem Definition

1.1 1.1 Application Objective

This application uses the Brazilian E-Commerce Public Dataset by Olist to analyze customer purchasing behavior and discover customer interests.

The application has two prediction outputs:

1. Customer behavior
2. Customer product-category interest

After predicting the customer's interest, the system recommends the Top 5 products that are most relevant to that predicted category.

The main supervised-learning task is:

Predict the product category of a customer's target purchase based on the customer's previous purchasing history.

The learning problem is defined as:

- X = customer behavioral features from previous purchases
- y = product category of the target purchase

The final application follows:

Customer History -> Customer Profile -> Feature Vector -> Machine Learning Model -> Predicted Behavior -> Predicted Interest -> Top 5 Product Recommendations

1.2 1.2 Why This Formulation?

Each customer may have multiple orders.

Instead of using an entire customer's history to predict the same history, we simulate a prediction scenario:

Previous orders -> input features Later target order -> prediction target

This prevents the target order from directly becoming part of the input representation.

1.3 1.3 Main Questions

1.3.1 What does one observation represent?

One observation represents one customer who has at least two delivered orders:

- previous orders are used as input;
- the customer's final delivered order is used to create the target.

1.3.2 What is X?

X is a customer-level behavioral feature vector.

1.3.3 What is y?

y is the product category of the target purchase.

1.3.4 What is the final objective?

The objective is not only to train a classifier. The objective is to build a complete deployable intelligent system:

Data -> Clean -> Represent -> Train -> Evaluate -> Persist -> Web -> Mobile

```
[41]: # =====
# CELL 1 - Environment, imports and global configuration
# =====

import os
import json
import random
import warnings
from pathlib import Path

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from IPython.display import display

from sklearn.base import clone
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import (
    OneHotEncoder,
    StandardScaler,
    FunctionTransformer
)
from sklearn.impute import SimpleImputer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
```

```

from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC

from sklearn.metrics import (
    accuracy_score,
    precision_score,
    recall_score,
    f1_score,
    confusion_matrix,
    classification_report,
    roc_auc_score
)

import joblib

warnings.filterwarnings("ignore")

# -----
# Reproducibility
# -----

RANDOM_SEED = 42

np.random.seed(RANDOM_SEED)
random.seed(RANDOM_SEED)

# -----
# Locate DATA folder
# Notebook location:
# assignment 2/A2_submit/e-commerce/e_commerce.ipynb
#
# DATA folder:
# assignment 2/DATA/
# -----

cwd = Path.cwd()

candidate_data_dirs = [
    cwd / "../.. /DATA",
    cwd.parent.parent / "DATA",
    cwd.parent / "DATA",
    cwd / "DATA"
]

DATA_DIR = None

for candidate in candidate_data_dirs:

```

```

    resolved = candidate.resolve()
    if resolved.exists() and resolved.is_dir():
        DATA_DIR = resolved
        break

if DATA_DIR is None:
    raise FileNotFoundError(
        "DATA folder was not found.\n"
        f"Current working directory: {cwd}\n"
        "Please check the notebook location."
    )

print("Current working directory:")
print(cwd.resolve())

print("\nDATA directory:")
print(DATA_DIR)

# -----
# Dataset paths
# -----

DATA_FILES = {
    "reviews": DATA_DIR / "olist_order_reviews_dataset.csv",
    "orders": DATA_DIR / "olist_orders_dataset.csv",
    "products": DATA_DIR / "olist_products_dataset.csv",
    "sellers": DATA_DIR / "olist_sellers_dataset.csv",
    "translation": DATA_DIR / "product_category_name_translation.csv",
    "customers": DATA_DIR / "olist_customers_dataset.csv",
    "geolocation": DATA_DIR / "olist_geolocation_dataset.csv",
    "order_items": DATA_DIR / "olist_order_items_dataset.csv",
    "payments": DATA_DIR / "olist_order_payments_dataset.csv",
}

# -----
# Model output directory
# Save artifacts outside DATA.
# -----

PROJECT_DIR = DATA_DIR.parent

MODEL_DIR = (
    PROJECT_DIR
    / "A2_submit"
    / "e-commerce"
    / "model"
)

```

```
MODEL_DIR.mkdir(parents=True, exist_ok=True)

print("\nModel directory:")
print(MODEL_DIR)

print("\nRandom seed:", RANDOM_SEED)
```

Current working directory:

C:\Users\ADMIN\Downloads\Intelligent System\assignment 2\A2_submit\e-commerce

DATA directory:

C:\Users\ADMIN\Downloads\Intelligent System\assignment 2\DATA

Model directory:

C:\Users\ADMIN\Downloads\Intelligent System\assignment
2\A2_submit\e-commerce\model

Random seed: 42

2 2. Dataset Source

2.1 Kaggle Dataset

Dataset name:

Brazilian E-Commerce Public Dataset by Olist

Kaggle URL:

<https://www.kaggle.com/datasets/olistbr/brazilian-ecommerce>

The Olist dataset is a relational e-commerce dataset containing several related tables.

The main tables are:

- Customers
- Orders
- Order Items
- Products
- Payments
- Reviews
- Sellers
- Geolocation
- Product Category Translation

2.2 Tables Used in the Main Application

The core customer-interest system uses:

1. Customers
2. Orders

3. Order Items
4. Products
5. Payments
6. Reviews
7. Product Category Translation

Sellers and geolocation are inspected but are not required for the main customer interest prediction pipeline.

2.3 Why a Relational Dataset?

A customer's behavior cannot be represented correctly from a single table.

The required representation therefore follows:

Orders -> Order Items -> Products

and:

Orders -> Payments -> Reviews

Finally:

Multiple transactions -> Customer Profile -> Feature Vector

```
[42]: # =====
# CELL 2 - Dataset metadata
# =====

dataset_metadata = {
    "Dataset": "Brazilian E-Commerce Public Dataset by Olist",
    "Source": "Kaggle",
    "Kaggle URL": "https://www.kaggle.com/datasets/olistbr/brazilian-ecommerce",
    "Application": "E-commerce Customer Behavior and Interest Discovery",
    "Main task": "Customer product-category interest classification",
    "Secondary task": "Customer behavior/value classification",
    "Random seed": RANDOM_SEED,
}

metadata_df = pd.DataFrame(
    list(dataset_metadata.items()),
    columns=["Property", "Value"]
)

display(metadata_df)
```

	Property	Value
0	Dataset	Brazilian E-Commerce Public Dataset by Olist
1	Source	Kaggle
2	Kaggle URL	https://www.kaggle.com/datasets/olistbr/brazil...
3	Application	E-commerce Customer Behavior and Interest Disc...
4	Main task	Customer product-category interest classification

5	Secondary task	Customer behavior/value classification
6	Random seed	42

3. Dataset Loading

The Olist dataset consists of multiple CSV files.

The purpose of this section is to load every available table once and keep the DataFrames in memory so that later sections do not repeatedly read the CSV files.

The main relationships are:

customer_id -> orders -> order_items -> product_id -> products

and:

order_id -> payments

order_id -> reviews

The stable customer identifier used for customer-level analysis is:

customer_unique_id

This is important because the same real customer can be associated with more than one customer_id across different orders.

```
[43]: # =====
# CELL 3 - Load all Olist CSV files ONCE
# =====

reviews = pd.read_csv(
    DATA_FILES["reviews"],
    parse_dates=[
        "review_creation_date",
        "review_answer_timestamp"
    ]
)

orders = pd.read_csv(
    DATA_FILES["orders"],
    parse_dates=[
        "order_purchase_timestamp",
        "order_approved_at",
        "order_delivered_carrier_date",
        "order_delivered_customer_date",
        "order_estimated_delivery_date"
    ]
)

products = pd.read_csv(DATA_FILES["products"])
```

```

sellers = pd.read_csv(DATA_FILES["sellers"])

translation = pd.read_csv(
    DATA_FILES["translation"]
)

customers = pd.read_csv(
    DATA_FILES["customers"]
)

geolocation = pd.read_csv(
    DATA_FILES["geolocation"]
)

order_items = pd.read_csv(
    DATA_FILES["order_items"],
    parse_dates=["shipping_limit_date"]
)

payments = pd.read_csv(
    DATA_FILES["payments"]
)

dataframes = {
    "reviews": reviews,
    "orders": orders,
    "products": products,
    "sellers": sellers,
    "translation": translation,
    "customers": customers,
    "geolocation": geolocation,
    "order_items": order_items,
    "payments": payments
}

print("Loaded tables:")
print()

for name, df in dataframes.items():
    print(f"{name:15s} -> {df.shape}")

```

Loaded tables:

```

reviews      -> (99224, 7)
orders       -> (99441, 8)
products     -> (32951, 9)
sellers      -> (3095, 4)
translation  -> (71, 2)

```

```
customers      -> (99441, 5)
geolocation    -> (1000163, 5)
order_items    -> (112650, 7)
payments       -> (103886, 5)
```

4 4. Dataset Inspection

Assignment 02 requires the following inspection operations:

- `df.shape`
- `df.head()`
- `df.info()`
- `df.describe()`
- `df.isna().sum()`
- `df.duplicated().sum()`

The inspection is performed before cleaning.

This is important because data quality problems should be documented before they are modified or removed.

4.1 Questions

4.1.1 What does one row represent?

The meaning depends on the table:

- customer table -> one customer record
- order table -> one order
- order_items table -> one product line within an order
- products table -> one product
- payments table -> one payment record
- reviews table -> one review

4.1.2 Why is inspection necessary?

The ML model will eventually require one consistent customer-level observation, but the raw dataset contains multiple levels of information.

```
[44]: # =====
# CELL 4 - Required dataset inspection
# =====

inspection_tables = {
    "customers": customers,
    "orders": orders,
    "order_items": order_items,
    "products": products,
    "payments": payments,
    "reviews": reviews,
```

```

}

for name, df in inspection_tables.items():

    print("=" * 90)
    print(f"TABLE: {name.upper()}")
    print("=" * 90)

    print("\nShape:")
    print(df.shape)

    print("\nHead:")
    display(df.head())

    print("\nInfo:")
    df.info()

    print("\nDescribe:")
    display(
        df.describe(
            include="all"
        ).T.head(40)
    )

    print("\nMissing values:")
    display(
        df.isna().sum()
        .sort_values(ascending=False)
        .head(20)
        .to_frame("missing_count")
    )

    print("\nDuplicated rows:")
    print(df.duplicated().sum())

    print()

```

```

=====
=====
TABLE: CUSTOMERS
=====
=====

```

Shape:
(99441, 5)

Head:

	customer_id	customer_unique_id \
0	06b8999e2fba1a1fbc88172c00ba8bc7	861eff4711a542e4b93843c6dd7febb0
1	18955e83d337fd6b2def6b18a428ac77	290c77bc529b7ac935b93aa66c333dc3
2	4e7b3e00288586ebd08712fdd0374a03	060e732b5b29e8181a18229c7b0b2b5e
3	b2b6027bc5c5109e529d4dc6358b12c3	259dac757896d24d7702b9acbbff3f3c
4	4f2d8ab171c80ec8364f7c12e35b23ad	345ecd01c38d18a9036ed96c73b8d066

	customer_zip_code_prefix	customer_city	customer_state
0	14409	franca	SP
1	9790	sao bernardo do campo	SP
2	1151	sao paulo	SP
3	8775	mogi das cruzeiras	SP
4	13056	campinas	SP

Info:

<class 'pandas.DataFrame'>

RangeIndex: 99441 entries, 0 to 99440

Data columns (total 5 columns):

#	Column	Non-Null Count	Dtype
0	customer_id	99441 non-null	str
1	customer_unique_id	99441 non-null	str
2	customer_zip_code_prefix	99441 non-null	int64
3	customer_city	99441 non-null	str
4	customer_state	99441 non-null	str

dtypes: int64(1), str(4)

memory usage: 11.0 MB

Describe:

	count	unique	top \
customer_id	99441	99441	06b8999e2fba1a1fbc88172c00ba8bc7
customer_unique_id	99441	96096	8d50f5eadf50201ccdcdfb9e2ac8455
customer_zip_code_prefix	99441.0	NaN	NaN
customer_city	99441	4119	sao paulo
customer_state	99441	27	SP

	freq	mean	std	min	25% \
customer_id	1	NaN	NaN	NaN	NaN
customer_unique_id	17	NaN	NaN	NaN	NaN
customer_zip_code_prefix	NaN	35137.474583	29797.938996	1003.0	11347.0
customer_city	15540	NaN	NaN	NaN	NaN
customer_state	41746	NaN	NaN	NaN	NaN

	50%	75%	max
customer_id	NaN	NaN	NaN
customer_unique_id	NaN	NaN	NaN
customer_zip_code_prefix	24416.0	58900.0	99990.0

customer_city	NaN	NaN	NaN
customer_state	NaN	NaN	NaN

Missing values:

	missing_count
customer_id	0
customer_unique_id	0
customer_zip_code_prefix	0
customer_city	0
customer_state	0

Duplicated rows:

0

```
=====
=====
TABLE: ORDERS
=====
=====
```

Shape:

(99441, 8)

Head:

	order_id	customer_id \
0	e481f51cbdc54678b7cc49136f2d6af7	9ef432eb6251297304e76186b10a928d
1	53cdb2fc8bc7dce0b6741e2150273451	b0830fb4747a6c6d20dea0b8c802d7ef
2	47770eb9100c2d0c44946d9cf07ec65d	41ce2a54c0b03bf3443c3d931a367089
3	949d5b44dbf5de918fe9c16f97b45f8a	f88197465ea7920adcbec7375364d82
4	ad21c59c0840e6cb83a9ceb5573f8159	8ab97904e6daea8866dbdbc4fb7aad2c

	order_status	order_purchase_timestamp	order_approved_at \
0	delivered	2017-10-02 10:56:33	2017-10-02 11:07:15
1	delivered	2018-07-24 20:41:37	2018-07-26 03:24:27
2	delivered	2018-08-08 08:38:49	2018-08-08 08:55:23
3	delivered	2017-11-18 19:28:06	2017-11-18 19:45:59
4	delivered	2018-02-13 21:18:39	2018-02-13 22:20:29

	order_delivered_carrier_date	order_delivered_customer_date \
0	2017-10-04 19:55:00	2017-10-10 21:25:13
1	2018-07-26 14:31:00	2018-08-07 15:27:45
2	2018-08-08 13:50:00	2018-08-17 18:06:29
3	2017-11-22 13:39:59	2017-12-02 00:28:42
4	2018-02-14 19:46:34	2018-02-16 18:17:02

order_estimated_delivery_date

```

0          2017-10-18
1          2018-08-13
2          2018-09-04
3          2017-12-15
4          2018-02-26

```

Info:

```
<class 'pandas.DataFrame'>
```

```
RangeIndex: 99441 entries, 0 to 99440
```

```
Data columns (total 8 columns):
```

#	Column	Non-Null Count	Dtype
0	order_id	99441 non-null	str
1	customer_id	99441 non-null	str
2	order_status	99441 non-null	str
3	order_purchase_timestamp	99441 non-null	datetime64[us]
4	order_approved_at	99281 non-null	datetime64[us]
5	order_delivered_carrier_date	97658 non-null	datetime64[us]
6	order_delivered_customer_date	96476 non-null	datetime64[us]
7	order_estimated_delivery_date	99441 non-null	datetime64[us]

```
dtypes: datetime64[us](5), str(3)
```

```
memory usage: 13.0 MB
```

Describe:

	count	unique	top \
order_id	99441	99441	e481f51cbdc54678b7cc49136f2d6af7
customer_id	99441	99441	9ef432eb6251297304e76186b10a928d
order_status	99441	8	delivered
order_purchase_timestamp	99441	NaN	NaN
order_approved_at	99281	NaN	NaN
order_delivered_carrier_date	97658	NaN	NaN
order_delivered_customer_date	96476	NaN	NaN
order_estimated_delivery_date	99441	NaN	NaN

	freq	mean \
order_id	1	NaN
customer_id	1	NaN
order_status	96478	NaN
order_purchase_timestamp	NaN	2017-12-31 08:43:12.776581
order_approved_at	NaN	2017-12-31 18:35:24.098800
order_delivered_carrier_date	NaN	2018-01-04 21:49:48.138278
order_delivered_customer_date	NaN	2018-01-14 12:09:19.035542
order_estimated_delivery_date	NaN	2018-01-24 03:08:37.730111

	min \
order_id	NaN
customer_id	NaN

order_status	NaN
order_purchase_timestamp	2016-09-04 21:15:19
order_approved_at	2016-09-15 12:16:38
order_delivered_carrier_date	2016-10-08 10:34:01
order_delivered_customer_date	2016-10-11 13:46:32
order_estimated_delivery_date	2016-09-30 00:00:00

	25%	\
order_id	NaN	
customer_id	NaN	
order_status	NaN	
order_purchase_timestamp	2017-09-12 14:46:19	
order_approved_at	2017-09-12 23:24:16	
order_delivered_carrier_date	2017-09-15 22:28:50.250000	
order_delivered_customer_date	2017-09-25 22:07:22.250000	
order_estimated_delivery_date	2017-10-03 00:00:00	

	50%	\
order_id	NaN	
customer_id	NaN	
order_status	NaN	
order_purchase_timestamp	2018-01-18 23:04:36	
order_approved_at	2018-01-19 11:36:13	
order_delivered_carrier_date	2018-01-24 16:10:58	
order_delivered_customer_date	2018-02-02 19:28:10.500000	
order_estimated_delivery_date	2018-02-15 00:00:00	

	75%		max
order_id	NaN		NaN
customer_id	NaN		NaN
order_status	NaN		NaN
order_purchase_timestamp	2018-05-04 15:42:16	2018-10-17 17:30:18	
order_approved_at	2018-05-04 20:35:10	2018-09-03 17:40:06	
order_delivered_carrier_date	2018-05-08 13:37:45	2018-09-11 19:48:28	
order_delivered_customer_date	2018-05-15 22:48:52.250000	2018-10-17 13:22:46	
order_estimated_delivery_date	2018-05-25 00:00:00	2018-11-12 00:00:00	

Missing values:

	missing_count
order_delivered_customer_date	2965
order_delivered_carrier_date	1783
order_approved_at	160
order_id	0
order_purchase_timestamp	0
order_status	0
customer_id	0
order_estimated_delivery_date	0

Duplicated rows:

0

```
=====
=====
TABLE: ORDER_ITEMS
=====
=====
```

Shape:

(112650, 7)

Head:

```

              order_id  order_item_id  \
0  00010242fe8c5a6d1ba2dd792cb16214      1
1  00018f77f2f0320c557190d7a144bdd3      1
2  000229ec398224ef6ca0657da4fc703e      1
3  00024acbcd0a6daa1e931b038114c75      1
4  00042b26cf59d7ce69dfabb4e55b4fd9      1

              product_id              seller_id  \
0  4244733e06e7ecb4970a6e2683c13e61  48436dade18ac8b2bce089ec2a041202
1  e5f2d52b802189ee658865ca93d83a8f  dd7ddc04e1b6c2c614352b383efe2d36
2  c777355d18b72b67abbeef9df44fd0fd  5b51032eddd242adc84c38acab88f23d
3  7634da152a4610f1595efa32f14722fc  9d7a1d34a5052409006425275ba1c2b4
4  ac6c3623068f30de03045865e4e10089  df560393f3a51e74553ab94004ba5c87

shipping_limit_date  price  freight_value
0  2017-09-19 09:45:35   58.90           13.29
1  2017-05-03 11:05:13  239.90           19.93
2  2018-01-18 14:48:30  199.00           17.87
3  2018-08-15 10:10:18   12.99           12.79
4  2017-02-13 13:57:51  199.90           18.14
```

Info:

<class 'pandas.DataFrame'>

RangeIndex: 112650 entries, 0 to 112649

Data columns (total 7 columns):

#	Column	Non-Null Count	Dtype
0	order_id	112650 non-null	str
1	order_item_id	112650 non-null	int64
2	product_id	112650 non-null	str
3	seller_id	112650 non-null	str
4	shipping_limit_date	112650 non-null	datetime64[us]
5	price	112650 non-null	float64

```

6    freight_value          112650 non-null float64
dtypes: datetime64[us](1), float64(2), int64(1), str(3)
memory usage: 16.3 MB

```

Describe:

	count	unique	top	freq	\
order_id	112650	98666	8272b63d03f5f79c56e9e4120aec44ef	21	
order_item_id	112650.0	NaN	NaN	NaN	
product_id	112650	32951	aca2eb7d00ea1a7b8ebd4e68314663af	527	
seller_id	112650	3095	6560211a19b47992c3666cc44a7e94c0	2033	
shipping_limit_date	112650	NaN	NaN	NaN	
price	112650.0	NaN	NaN	NaN	
freight_value	112650.0	NaN	NaN	NaN	

	mean	min	\
order_id	NaN	NaN	
order_item_id	1.197834	1.0	
product_id	NaN	NaN	
seller_id	NaN	NaN	
shipping_limit_date	2018-01-07 15:36:52.192685	2016-09-19 00:15:34	
price	120.653739	0.85	
freight_value	19.99032	0.0	

	25%	50%	\
order_id	NaN	NaN	
order_item_id	1.0	1.0	
product_id	NaN	NaN	
seller_id	NaN	NaN	
shipping_limit_date	2017-09-20 20:57:27.500000	2018-01-26 13:59:35	
price	39.9	74.99	
freight_value	13.08	16.26	

	75%	max	\
order_id	NaN	NaN	
order_item_id	1.0	21.0	
product_id	NaN	NaN	
seller_id	NaN	NaN	
shipping_limit_date	2018-05-10 14:34:00.750000	2020-04-09 22:35:08	
price	134.9	6735.0	
freight_value	21.15	409.68	

	std
order_id	NaN
order_item_id	0.705124
product_id	NaN
seller_id	NaN
shipping_limit_date	NaN

```
price          183.633928
freight_value  15.806405
```

Missing values:

```
missing_count
order_id      0
order_item_id 0
product_id    0
seller_id     0
shipping_limit_date 0
price         0
freight_value 0
```

Duplicated rows:

0

```
=====
=====
TABLE: PRODUCTS
=====
=====
```

Shape:

(32951, 9)

Head:

```
product_id  product_category_name \
0  1e9e8ef04dbcff4541ed26657ea517e5      perfumaria
1  3aa071139cb16b67ca9e5dea641aaa2f      artes
2  96bd76ec8810374ed1b65e291975717f      esporte_lazer
3  cef67bcfe19066a932b7673e239eb23d      bebes
4  9dc1a7de274444849c219cff195d0b71  utilidades_domesticas
```

```
product_name_lenght  product_description_lenght  product_photos_qty \
0          40.0          287.0          1.0
1          44.0          276.0          1.0
2          46.0          250.0          1.0
3          27.0          261.0          1.0
4          37.0          402.0          4.0
```

```
product_weight_g  product_length_cm  product_height_cm  product_width_cm
0          225.0          16.0          10.0          14.0
1         1000.0          30.0          18.0          20.0
2          154.0          18.0           9.0          15.0
3          371.0          26.0           4.0          26.0
4          625.0          20.0          17.0          13.0
```

Info:

<class 'pandas.DataFrame'>

RangeIndex: 32951 entries, 0 to 32950

Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
0	product_id	32951 non-null	str
1	product_category_name	32341 non-null	str
2	product_name_lenght	32341 non-null	float64
3	product_description_lenght	32341 non-null	float64
4	product_photos_qty	32341 non-null	float64
5	product_weight_g	32949 non-null	float64
6	product_length_cm	32949 non-null	float64
7	product_height_cm	32949 non-null	float64
8	product_width_cm	32949 non-null	float64

dtypes: float64(7), str(2)

memory usage: 3.7 MB

Describe:

	count	unique	top
product_id	32951	32951	1e9e8ef04dbcff4541ed26657ea517e5
product_category_name	32341	73	cama_mesa_banho
product_name_lenght	32341.0	NaN	NaN
product_description_lenght	32341.0	NaN	NaN
product_photos_qty	32341.0	NaN	NaN
product_weight_g	32949.0	NaN	NaN
product_length_cm	32949.0	NaN	NaN
product_height_cm	32949.0	NaN	NaN
product_width_cm	32949.0	NaN	NaN

	freq	mean	std	min	25%	50%
product_id	1	NaN	NaN	NaN	NaN	NaN
product_category_name	3029	NaN	NaN	NaN	NaN	NaN
product_name_lenght	NaN	48.476949	10.245741	5.0	42.0	51.0
product_description_lenght	NaN	771.495285	635.115225	4.0	339.0	595.0
product_photos_qty	NaN	2.188986	1.736766	1.0	1.0	1.0
product_weight_g	NaN	2276.472488	4282.038731	0.0	300.0	700.0
product_length_cm	NaN	30.815078	16.914458	7.0	18.0	25.0
product_height_cm	NaN	16.937661	13.637554	2.0	8.0	13.0
product_width_cm	NaN	23.196728	12.079047	6.0	15.0	20.0

	75%	max
product_id	NaN	NaN
product_category_name	NaN	NaN
product_name_lenght	57.0	76.0
product_description_lenght	972.0	3992.0

product_photos_qty	3.0	20.0
product_weight_g	1900.0	40425.0
product_length_cm	38.0	105.0
product_height_cm	21.0	105.0
product_width_cm	30.0	118.0

Missing values:

	missing_count
product_category_name	610
product_description_lenght	610
product_name_lenght	610
product_photos_qty	610
product_weight_g	2
product_height_cm	2
product_length_cm	2
product_width_cm	2
product_id	0

Duplicated rows:

0

```
=====
=====
TABLE: PAYMENTS
=====
=====
```

Shape:

(103886, 5)

Head:

	order_id	payment_sequential	payment_type	\
0	b81ef226f3fe1789b1e8b2acac839d17	1	credit_card	
1	a9810da82917af2d9aefd1278f1dcfa0	1	credit_card	
2	25e8ea4e93396b6fa0d3dd708e76c1bd	1	credit_card	
3	ba78997921bbcdc1373bb41e913ab953	1	credit_card	
4	42fdf880ba16b47b59251dd489d4441a	1	credit_card	

	payment_installments	payment_value
0	8	99.33
1	1	24.39
2	1	65.71
3	8	107.78
4	2	128.45

Info:

<class 'pandas.DataFrame'>

RangeIndex: 103886 entries, 0 to 103885

Data columns (total 5 columns):

#	Column	Non-Null Count	Dtype
0	order_id	103886 non-null	str
1	payment_sequential	103886 non-null	int64
2	payment_type	103886 non-null	str
3	payment_installments	103886 non-null	int64
4	payment_value	103886 non-null	float64

dtypes: float64(1), int64(2), str(2)

memory usage: 8.1 MB

Describe:

	count	unique	top \
order_id	103886	99440	fa65dad1b0e818e3ccc5cb0e39231352
payment_sequential	103886.0	NaN	NaN
payment_type	103886	5	credit_card
payment_installments	103886.0	NaN	NaN
payment_value	103886.0	NaN	NaN

	freq	mean	std	min	25%	50% \
order_id	29	NaN	NaN	NaN	NaN	NaN
payment_sequential	NaN	1.092679	0.706584	1.0	1.0	1.0
payment_type	76795	NaN	NaN	NaN	NaN	NaN
payment_installments	NaN	2.853349	2.687051	0.0	1.0	1.0
payment_value	NaN	154.10038	217.494064	0.0	56.79	100.0

	75%	max
order_id	NaN	NaN
payment_sequential	1.0	29.0
payment_type	NaN	NaN
payment_installments	4.0	24.0
payment_value	171.8375	13664.08

Missing values:

	missing_count
order_id	0
payment_sequential	0
payment_type	0
payment_installments	0
payment_value	0

Duplicated rows:

0

```
=====
TABLE: REVIEWS
=====
```

```
Shape:
(99224, 7)
```

```
Head:
```

```

      review_id      order_id \
0  7bc2406110b926393aa56f80a40eba40  73fc7af87114b39712e6da79b0a377eb
1  80e641a11e56f04c1ad469d5645fdfde  a548910a1c6147796b98fdf73dbeba33
2  228ce5500dc1d8e020d8d1322874b6f0  f9e4b658b201a9f2ecdecbb34bed034b
3  e64fb393e7b32834bb789ff8bb30750e  658677c97b385a9be170737859d3511b
4  f7c4243c7fe1938f181bec41a392bdeb  8e6bfb81e283fa7e4f11123a3fb894f1

      review_score review_comment_title \
0                4                NaN
1                5                NaN
2                5                NaN
3                5                NaN
4                5                NaN

      review_comment_message review_creation_date \
0                NaN                2018-01-18
1                NaN                2018-03-10
2                NaN                2018-02-17
3      Recebi bem antes do prazo estipulado.    2017-04-21
4  Parabéns lojas lannister adorei comprar pela I... 2018-03-01

      review_answer_timestamp
0      2018-01-18 21:46:59
1      2018-03-11 03:05:13
2      2018-02-18 14:36:24
3      2017-04-21 22:02:06
4      2018-03-02 10:26:53
```

```
Info:
```

```
<class 'pandas.DataFrame'>
```

```
RangeIndex: 99224 entries, 0 to 99223
```

```
Data columns (total 7 columns):
```

```

#      Column      Non-Null Count  Dtype
---  -
0      review_id    99224 non-null  str
1      order_id     99224 non-null  str
```

```

2   review_score          99224 non-null  int64
3   review_comment_title   11568 non-null  str
4   review_comment_message 40977 non-null  str
5   review_creation_date    99224 non-null  datetime64[us]
6   review_answer_timestamp 99224 non-null  datetime64[us]
dtypes: datetime64[us](2), int64(1), str(4)
memory usage: 14.2 MB

```

Describe:

	count	unique	top \
review_id	99224	98410	c444278834184f72b1484dfe47de7f97
order_id	99224	98673	c88b1d1b157a9999ce368f218a407141
review_score	99224.0	NaN	NaN
review_comment_title	11568	4527	Recomendo
review_comment_message	40977	36159	Muito bom
review_creation_date	99224	NaN	NaN
review_answer_timestamp	99224	NaN	NaN

	freq	mean	min \
review_id	3	NaN	NaN
order_id	3	NaN	NaN
review_score	NaN	4.086421	1.0
review_comment_title	423	NaN	NaN
review_comment_message	230	NaN	NaN
review_creation_date	NaN	2018-01-12 20:49:23.948238	2016-10-02 00:00:00
review_answer_timestamp	NaN	2018-01-16 00:23:56.977939	2016-10-07 18:32:28

	25% \
review_id	NaN
order_id	NaN
review_score	4.0
review_comment_title	NaN
review_comment_message	NaN
review_creation_date	2017-09-23 00:00:00
review_answer_timestamp	2017-09-27 01:53:27.250000

	50% \
review_id	NaN
order_id	NaN
review_score	5.0
review_comment_title	NaN
review_comment_message	NaN
review_creation_date	2018-02-02 00:00:00
review_answer_timestamp	2018-02-04 22:41:47.500000

	75%	max \
review_id	NaN	NaN

order_id		NaN	NaN
review_score		5.0	5.0
review_comment_title		NaN	NaN
review_comment_message		NaN	NaN
review_creation_date	2018-05-16 00:00:00	2018-08-31 00:00:00	
review_answer_timestamp	2018-05-20 12:11:21.500000	2018-10-29 12:27:35	

	std
review_id	NaN
order_id	NaN
review_score	1.347579
review_comment_title	NaN
review_comment_message	NaN
review_creation_date	NaN
review_answer_timestamp	NaN

Missing values:

	missing_count
review_comment_title	87656
review_comment_message	58247
review_id	0
review_score	0
order_id	0
review_creation_date	0
review_answer_timestamp	0

Duplicated rows:

0

5 5. Data-Quality Analysis

The Olist dataset contains transactional information, so data quality must be checked from both technical and business perspectives.

Important problems to investigate include:

1. Invalid order statuses
2. Invalid or impossible prices
3. Invalid freight values
4. Invalid payment values
5. Missing product categories
6. Duplicate transactional records
7. Missing review comments
8. Extremely large purchases

5.1 Business Rules

For the main customer behavior model:

- only delivered orders are treated as completed purchases;
- product price must be greater than zero;
- freight value must not be negative;
- payment value must be greater than zero;
- a target purchase must have an identifiable product category.

Unusual but valid purchases are not automatically deleted. They are analyzed separately as possible outliers.

5.2 Why?

Deleting high-value customers simply because they are unusual could remove exactly the customer behavior that the system is intended to understand.

```
[45]: # =====  
# CELL 5 - General data quality checks  
# =====  
  
print("ORDER STATUS")  
display(  
    orders["order_status"]  
    .value_counts(dropna=False)  
    .to_frame("count")  
)  
  
print("\nORDER ITEM PRICE")  
print("Minimum:", order_items["price"].min())  
print("Zero or negative:", (order_items["price"] <= 0).sum())  
  
print("\nFREIGHT VALUE")  
print("Minimum:", order_items["freight_value"].min())  
print("Negative:", (order_items["freight_value"] < 0).sum())  
  
print("\nPAYMENT VALUE")  
print("Minimum:", payments["payment_value"].min())  
print("Zero or negative:", (payments["payment_value"] <= 0).sum())  
  
print("\nPRODUCT CATEGORY")  
print(  
    "Missing:",  
    products["product_category_name"].isna().sum()  
)  
  
print(  
    "Unique categories:",
```

```

    products["product_category_name"].nunique()
)

print("\nREVIEW COMMENT")
print(
    "Missing comment:",
    reviews["review_comment_message"].isna().sum()
)

print(
    "Total reviews:",
    len(reviews)
)

```

ORDER STATUS

	count
order_status	
delivered	96478
shipped	1107
canceled	625
unavailable	609
invoiced	314
processing	301
created	5
approved	2

ORDER ITEM PRICE

Minimum: 0.85
Zero or negative: 0

FREIGHT VALUE

Minimum: 0.0
Negative: 0

PAYMENT VALUE

Minimum: 0.0
Zero or negative: 9

PRODUCT CATEGORY

Missing: 610
Unique categories: 73

REVIEW COMMENT

Missing comment: 58247
Total reviews: 99224

6 6. Missing-Value Analysis

Missing values have different meanings depending on the variable.

Examples:

- Missing review comment can mean that the customer did not provide textual feedback.
- Missing product dimensions do not necessarily mean that the product is invalid.
- Missing categorical customer information can be handled by imputation.
- Missing target category cannot be accepted because there is no valid target.

6.1 Strategy

Numerical features:

Missing values -> Median imputation

Categorical features:

Missing values -> Most frequent value

Text:

Missing comment -> Empty string

Target:

Missing target category -> Remove from supervised-learning dataset

The same preprocessing logic will later be stored and reused during inference.

```
[46]: # =====  
# CELL 6 - Missing value report  
# =====  
  
def missing_report(df):  
    result = pd.DataFrame({  
        "missing_count": df.isna().sum()  
    })  
  
    result["missing_percent"] = (  
        result["missing_count"]  
        / len(df)  
        * 100  
    )  
  
    return (  
        result  
        .sort_values(  
            "missing_count",  
            ascending=False  
        )  
    )
```

```

for name, df in inspection_tables.items():

    print("=" * 70)
    print(name.upper())
    print("=" * 70)

    report = missing_report(df)

    nonzero = report[
        report["missing_count"] > 0
    ]

    if len(nonzero) == 0:
        print("No missing values.")
    else:
        display(nonzero.head(20))

```

```
=====
CUSTOMERS
=====
```

```
No missing values.
=====
```

```
ORDERS
=====
```

	missing_count	missing_percent
order_delivered_customer_date	2965	2.981668
order_delivered_carrier_date	1783	1.793023
order_approved_at	160	0.160899

```
=====
ORDER_ITEMS
=====
```

```
No missing values.
=====
```

```
PRODUCTS
=====
```

	missing_count	missing_percent
product_category_name	610	1.851234
product_description_lenght	610	1.851234
product_name_lenght	610	1.851234
product_photos_qty	610	1.851234
product_weight_g	2	0.006070
product_height_cm	2	0.006070
product_length_cm	2	0.006070
product_width_cm	2	0.006070

```
=====
PAYMENTS
=====
No missing values.
=====
REVIEWS
=====
```

	missing_count	missing_percent
review_comment_title	87656	88.341530
review_comment_message	58247	58.702532

7 7. Duplicate Analysis

Duplicate detection is performed using both:

1. complete-row duplication;
2. business-key duplication.

A business key is more meaningful than simply comparing entire rows.

Examples:

- order_id identifies one order;
- product_id identifies one product;
- review_id identifies one review;
- order_id + order_item_id identifies one item line;
- order_id + payment_sequential identifies one payment record.

7.1 Important

Multiple rows are not automatically duplicates.

For example, one order can contain several different products, so multiple rows in order_items are expected.

Therefore, duplicate removal is performed only when the business key indicates an actual duplicate record.

```
[47]: # =====
# CELL 7 - Duplicate analysis
# =====

duplicate_summary = pd.DataFrame({
    "table": [
        "customers",
        "orders",
        "order_items",
        "products",
        "payments",
        "reviews"
    ]
})
```

```

],

"exact_duplicate_rows": [
    customers.duplicated().sum(),
    orders.duplicated().sum(),
    order_items.duplicated().sum(),
    products.duplicated().sum(),
    payments.duplicated().sum(),
    reviews.duplicated().sum()
],

"business_key_duplicates": [
    customers.duplicated(
        subset=["customer_id"]
    ).sum(),

    orders.duplicated(
        subset=["order_id"]
    ).sum(),

    order_items.duplicated(
        subset=[
            "order_id",
            "order_item_id"
        ]
    ).sum(),

    products.duplicated(
        subset=["product_id"]
    ).sum(),

    payments.duplicated(
        subset=[
            "order_id",
            "payment_sequential"
        ]
    ).sum(),

    reviews.duplicated(
        subset=["review_id"]
    ).sum()
]
})

display(duplicate_summary)

```

	table	exact_duplicate_rows	business_key_duplicates
0	customers	0	0

1	orders	0	0
2	order_items	0	0
3	products	0	0
4	payments	0	0
5	reviews	0	814

8 8. Invalid-Value Analysis

Invalid values are values that contradict the meaning of the business variable.

The cleaning rules are:

8.0.1 Orders

Only delivered orders are used as completed purchases.

8.0.2 Product price

Price ≤ 0 is considered invalid.

8.0.3 Freight

Freight value < 0 is invalid.

8.0.4 Payment

Payment value ≤ 0 is invalid for monetary aggregation.

8.0.5 Target category

A target order without a valid product category cannot be used for supervised interest classification.

8.1 Why not remove every unusual value?

A very expensive order can be a valid transaction. A customer with many purchases can be a valid high-value customer.

Therefore:

Invalid value -> Remove or correct

Extreme but valid value -> Analyze as outlier

```
[48]: # =====
# CELL 8 - Create cleaned transactional tables
# =====

# Delivered orders only
orders_clean = orders[
    orders["order_status"].eq("delivered")
].copy()
```

```

# Valid order items
order_items_clean = order_items[
    (order_items["price"] > 0) &
    (order_items["freight_value"] >= 0)
].copy()

# Valid payments
payments_clean = payments[
    payments["payment_value"] > 0
].copy()

# Product categories
products_clean = products.copy()

products_clean["product_category_name"] = (
    products_clean["product_category_name"]
    .fillna("unknown")
)

# Review comments
reviews_clean = reviews.copy()

reviews_clean["review_comment_message"] = (
    reviews_clean["review_comment_message"]
    .fillna("")
)

# Drop exact duplicate records only
orders_clean = orders_clean.drop_duplicates(
    subset=["order_id"]
)

order_items_clean = order_items_clean.drop_duplicates(
    subset=[
        "order_id",
        "order_item_id"
    ]
)

products_clean = products_clean.drop_duplicates(
    subset=["product_id"]
)

payments_clean = payments_clean.drop_duplicates(
    subset=[
        "order_id",
        "payment_sequential"
    ]
)

```

```

    ]
)

reviews_clean = reviews_clean.drop_duplicates(
    subset=["review_id"]
)

print("Cleaned table sizes:")
print()

for name, df in {
    "orders_clean": orders_clean,
    "order_items_clean": order_items_clean,
    "products_clean": products_clean,
    "payments_clean": payments_clean,
    "reviews_clean": reviews_clean
}.items():
    print(f"{name:20s} -> {df.shape}")

```

Cleaned table sizes:

```

orders_clean      -> (96478, 8)
order_items_clean -> (112650, 7)
products_clean    -> (32951, 9)
payments_clean    -> (103877, 5)
reviews_clean     -> (98410, 7)

```

9.9. Outlier Analysis

Customer behavior naturally contains extreme values.

Examples:

- very high spending;
- many orders;
- unusually expensive products;
- large freight cost.

These observations can represent legitimate customer behavior.

Therefore the main strategy is:

1. visualize the distribution;
2. quantify high percentiles;
3. use transformations where appropriate;
4. allow tree-based models to capture nonlinear behavior.

Outliers will not be deleted purely because they are statistically unusual.

This is particularly important for e-commerce because high-value customers are potentially important business segments.

```

[49]: # =====
# CELL 9 - Transaction-level outlier analysis
# =====

outlier_columns = [
    "price",
    "freight_value"
]

fig, axes = plt.subplots(
    1,
    2,
    figsize=(14, 5)
)

for ax, column in zip(
    axes,
    outlier_columns
):

    values = (
        order_items_clean[column]
        .dropna()
    )

    ax.boxplot(values)

    ax.set_title(
        f"Boxplot - {column}"
    )

    ax.set_ylabel(column)

plt.tight_layout()
plt.show()

# Quantile report
for column in outlier_columns:

    print("=" * 70)
    print(column)


    display(
        order_items_clean[column]
        .quantile(
            [
                0,

```

```

        0.01,
        0.25,
        0.50,
        0.75,
        0.90,
        0.95,
        0.99,
        1
    ]
)
.to_frame("value")
)

```



```
=====
price
```

	value
0.00	0.85
0.01	9.99
0.25	39.90
0.50	74.99
0.75	134.90
0.90	229.80
0.95	349.90
0.99	890.00
1.00	6735.00

```
=====
freight_value
```

	value
0.00	0.0000
0.01	4.4198
0.25	13.0800
0.50	16.2600

0.75	21.1500
0.90	34.0410
0.95	45.1200
0.99	84.5200
1.00	409.6800

10 10. Exploratory Data Analysis

EDA is used to understand the business behavior before building the model.

The visualization strategy is intentionally diverse.

We will examine:

1. Order-status distribution
2. Monthly order trend
3. Product-category frequency
4. Product-price distribution
5. Customer order-frequency distribution
6. Customer spending distribution
7. Frequency versus spending
8. Review-score distribution
9. Customer-state distribution
10. Correlation between important behavioral variables

For major visualizations we answer three questions:

10.0.1 Observation

What does the figure show?

10.0.2 Interpretation

What does that pattern mean?

10.0.3 ML implication

Why is the pattern relevant to feature engineering or model performance?

For the correlation matrix, numerical values are displayed directly inside the cells.

```
[50]: # =====  
# CELL 10A - Order status distribution  
# =====  
  
status_counts = orders[  
    "order_status"  
].value_counts()  
  
plt.figure(figsize=(10, 5))
```

```

plt.bar(
    status_counts.index,
    status_counts.values
)

plt.title(
    "Order Status Distribution"
)

plt.xlabel(
    "Order Status"
)

plt.ylabel(
    "Number of Orders"
)

plt.xticks(
    rotation=45
)


plt.tight_layout()
plt.show()

print(
    "Observation: Delivered orders dominate the usable transaction history."
)

print(
    "Interpretation: Most completed purchasing behavior can be represented,
    ↳using"
    " delivered orders."
)

print(
    "ML implication: The main supervised dataset should focus on completed,
    ↳purchases."
)

```



Observation: Delivered orders dominate the usable transaction history.

Interpretation: Most completed purchasing behavior can be represented using delivered orders.

ML implication: The main supervised dataset should focus on completed purchases.

```
[51]: # =====
# CELL 10B - Monthly purchase trend
# =====

monthly_orders = (
    orders_clean
    .set_index("order_purchase_timestamp")
    .resample("ME")
    .size()
)

plt.figure(figsize=(13, 5))

plt.plot(
    monthly_orders.index,
    monthly_orders.values,
    marker="o"
)

plt.title(
    "Monthly Delivered Orders"
)

plt.xlabel(
```

```


    "Month"
)

plt.ylabel(
    "Number of Orders"
)

plt.xticks(
    rotation=45
)

plt.tight_layout()
plt.show()

```



```

[52]: # =====
# CELL 10C - Product category distribution
# =====

category_translation = translation.set_index(
    "product_category_name"
)[
    "product_category_name_english"
].to_dict()

eda_items = (
    order_items_clean
    .merge(
        products_clean[
            [
                "product_id",
                "product_category_name"
            ]

```

```

        ],
        on="product_id",
        how="left"
    )
)

eda_items["category_en"] = (
    eda_items["product_category_name"]
    .map(category_translation)
    .fillna(
        eda_items["product_category_name"]
    )
    .fillna("unknown")
)

top_categories_eda = (
    eda_items["category_en"]
    .value_counts()
    .head(15)
)

plt.figure(figsize=(11, 7))

plt.barh(
    top_categories_eda.index[::-1],
    top_categories_eda.values[::-1]
)


plt.title(
    "Top 15 Product Categories"
)

plt.xlabel(
    "Number of Items Sold"
)

plt.ylabel(
    "Product Category"
)

plt.tight_layout()
plt.show()

```



```
[53]: # =====
# CELL 10D - Price distribution
# =====

price_limit = (
    order_items_clean["price"]
    .quantile(0.99)
)

plt.figure(figsize=(10, 5))

plt.hist(
    order_items_clean[
        order_items_clean["price"] <= price_limit
    ]["price"],
    bins=60
)

plt.title(
    "Product Price Distribution - Up to 99th Percentile"
)

plt.xlabel(
    "Price"
```


```

)

plt.ylabel(
    "Frequency"
)

plt.tight_layout()
plt.show()

```



```

[54]: # =====
# CELL 10E - Review score distribution
# =====

review_score_counts = (
    reviews_clean["review_score"]
    .value_counts()
    .sort_index()
)

plt.figure(figsize=(8, 5))

plt.bar(
    review_score_counts.index.astype(str),
    review_score_counts.values
)

plt.title(
    "Review Score Distribution"

```

```

)

plt.xlabel(
    "Review Score"
)


plt.ylabel(
    "Number of Reviews"
)

plt.tight_layout()
plt.show()

print(
    "Observation: Review scores are concentrated at several score levels."
)

print(
    "ML implication: Average review score can provide a useful behavioral_
    ↪feature,"
    " while missing review scores must be handled carefully."
)

```



Observation: Review scores are concentrated at several score levels.
 ML implication: Average review score can provide a useful behavioral feature,

while missing review scores must be handled carefully.

11 11. Feature Types

The final customer representation contains three major feature groups.

11.1 Numerical features

Examples:

- previous order count
- total spending
- average order value
- total purchased items
- average item price
- average freight value
- average purchase interval
- category purchase counts
- category purchase shares
- average review score
- review count
- average installments

11.2 Categorical features

Examples:

- customer state
- dominant payment type

Categorical values are not treated as ordered numbers.

They are encoded using One-Hot Encoding.

11.3 Text features

Customer review comments are represented separately.

The transformation is:

Review text -> cleaned text -> TF-IDF numerical representation

11.4 Target

The main target is:

target_category

which represents the product category of the target purchase.

```
[55]: # =====  
# CELL 11 - Utility for text cleaning  
# =====
```

```

import re

def clean_text(text):
    """
    Basic deterministic text cleaning.
    """
    if pd.isna(text):
        return ""

    text = str(text).lower()

    text = re.sub(
        r"\s+",
        " ",
        text
    )

    text = re.sub(
        r"[\w\s]",
        " ",
        text
    )

    text = re.sub(
        r"\s+",
        " ",
        text
    ).strip()

    return text

```

12. Data Representation

This section connects the Olist dataset directly to Lecture 02.

12.1 Raw Representation

The raw data is stored as multiple relational CSV tables.

Example:

CSV tables -> records -> transactions -> products -> reviews

12.2 Customer Representation

Multiple transactions are aggregated into a customer profile.

Transactions -> Customer Profile -> Feature Vector

A simplified customer vector is:

x = [frequency, monetary_value, average_order_value, total_items, category_activity, payment_behavior, review_score, customer_state]

12.3 Matrix Representation

The customer dataset becomes:

X shape = (N, d)

where:

- N = number of customer-level observations
- d = number of input features after preprocessing

12.4 Text Representation

For customer comments:

Comment -> cleaned text -> tokens -> numerical vector using TF-IDF

The important principle is:

CSV file is not ML representation is not model input

The exact representation supplied to a model is determined by the preprocessing pipeline.

```
[56]: # =====  
# CELL 12 - Build order-level category information  
# =====  
  
# Attach product category to every valid order item  
item_category = (  
    order_items_clean  
    .merge(  
        products_clean[  
            [  
                "product_id",  
                "product_category_name"  
            ]  
        ],  
        on="product_id",  
        how="left"  
    )  
)  
  
# Number of items and monetary value by order/category  
order_category_statistics = (  
    item_category  
    .groupby(  
        [  

```

```

        "order_id",
        "product_category_name"
    ],
    dropna=False
)
.agg(
    item_count=("order_item_id", "count"),
    category_value=("price", "sum")
)
.reset_index()
)

# Dominant category:
# highest item count, then highest category value
dominant_order_category = (
    order_category_statistics
    .sort_values(
        [
            "order_id",
            "item_count",
            "category_value"
        ],
        ascending=[
            True,
            False,
            False
        ]
    )
    .drop_duplicates(
        subset=["order_id"]
    )
    [
        [
            "order_id",
            "product_category_name"
        ]
    ]
)

display(
    dominant_order_category.head(10)
)

```

	order_id	product_category_name
0	00010242fe8c5a6d1ba2dd792cb16214	cool_stuff
1	00018f77f2f0320c557190d7a144bdd3	pet_shop
2	000229ec398224ef6ca0657da4fc703e	moveis_decoracao
3	00024acbcd0a6daa1e931b038114c75	perfumaria

4	00042b26cf59d7ce69dfabb4e55b4fd9	ferramentas_jardim
5	00048cc3ae777c65dbb7d2a0634bc1ea	utilidades_domesticas
6	00054e8431b9d7675808bcb819fb4a32	telefonica
7	000576fe39319847cbb9d288c5617fa6	ferramentas_jardim
8	0005a1a1728c9d785b8e2b08b904576c	beleza_saude
9	0005f50442cb953dcd1d21e1fb923495	livros_tecnicos

13. Feature Engineering

The main feature-engineering goal is:

Transactions -> Customer History -> Customer Profile -> Feature Vector

13.1 Customer-Level Behavioral Features

The feature vector contains:

13.1.1 RFM-style features

- prior_order_count
- total_spending
- average_order_value
- max_order_value
- recency
- average_days_between_orders

13.1.2 Purchase features

- total_items
- average_items_per_order
- average_item_price
- average_freight_value

13.1.3 Category features

- category purchase counts
- category purchase proportions

13.1.4 Payment behavior

- average installments
- average payment value
- dominant payment type

13.1.5 Review behavior

- average review score
- review count
- text review count
- aggregated previous review text

13.1.6 Customer profile

- customer state

13.2 Leakage Prevention

The following variables are explicitly excluded from X:

- target order ID
- target order timestamp
- target product category
- target order value
- customer ID

The target order is used only to construct y.

Therefore:

Previous customer history -> X

Target order -> y

```
[57]: # =====  
# CELL 13A - Create customer-order history  
# =====  
  
# Add customer_unique_id and customer_state  
delivered_orders = (  
    orders_clean[  
        [  
            "order_id",  
            "customer_id",  
            "order_purchase_timestamp"  
        ]  
    ]  
    .merge(  
        customers[  
            [  
                "customer_id",  
                "customer_unique_id",  
                "customer_state"  
            ]  
        ],  
        on="customer_id",  
        how="left"  
    )  
)  
  
# Attach dominant category of each delivered order  
delivered_orders = delivered_orders.merge(  
    dominant_order_category,
```

```

        on="order_id",
        how="left"
    )

    # Aggregate payment amount per order
    payment_by_order = (
        payments_clean
        .groupby("order_id")[
            "payment_value"
        ]
        .sum()
        .rename("order_payment_value")
    )

    delivered_orders = delivered_orders.merge(
        payment_by_order,
        on="order_id",
        how="left"
    )

    delivered_orders["order_payment_value"] = (
        delivered_orders["order_payment_value"]
        .fillna(0)
    )

    # Sort chronologically per customer
    delivered_orders = (
        delivered_orders
        .sort_values(
            [
                "customer_unique_id",
                "order_purchase_timestamp"
            ]
        )
        .reset_index(drop=True)
    )

    # Customers must have >= 2 delivered orders
    customer_order_counts = (
        delivered_orders
        .groupby(
            "customer_unique_id"
        )["order_id"]
        .nunique()
    )

    eligible_customer_ids = (

```

```

        customer_order_counts[
            customer_order_counts >= 2
        ]
        .index
    )

eligible_orders = delivered_orders[
    delivered_orders[
        "customer_unique_id"
    ].isin(
        eligible_customer_ids
    )
].copy()

print(
    "Delivered orders:",
    len(delivered_orders)
)

print(
    "Customers with at least 2 delivered orders:",
    len(eligible_customer_ids)
)

print(
    "Eligible delivered orders:",
    len(eligible_orders)
)

```

Delivered orders: 96478
 Customers with at least 2 delivered orders: 2801
 Eligible delivered orders: 5921

```

[58]: # =====
# CELL 13B - Create supervised customer snapshots
# =====

snapshot_records = []

for customer_unique_id, group in eligible_orders.groupby(
    "customer_unique_id"
):

    group = (
        group
        .sort_values(
            "order_purchase_timestamp"

```

```

    )
    .reset_index(drop=True)
)

# Last delivered order = target
target_order = group.iloc[-1]

# All previous orders = historical input
history_orders = group.iloc[:-1]

snapshot_records.append({
    "customer_unique_id":
        customer_unique_id,

    "target_order_id":
        target_order["order_id"],

    "target_order_ts":
        target_order[
            "order_purchase_timestamp"
        ],

    "target_category":
        target_order[
            "product_category_name"
        ],

    "target_order_value":
        target_order[
            "order_payment_value"
        ],

    "history_order_count":
        len(history_orders)
})

snapshots = pd.DataFrame(
    snapshot_records
)

# Target category must exist
snapshots = snapshots.dropna(
    subset=["target_category"]
).reset_index(drop=True)

print(
    "Number of supervised customer snapshots:",

```

```

len(snapshots)
)

display(
    snapshots.head(10)
)

```

Number of supervised customer snapshots: 2801

	customer_unique_id	target_order_id \
0	004288347e5e88a27ded2bb23747066c	08204559bebd39e09ee52dcb56d8faa2
1	00a39521eb40f7012db50455bf083460	cea3e6c11eb60acb9d8d4d51694832f8
2	00cc12a6d8b578b8ebd21ea4e2ae8b27	d61b915b69851aec8a8865f36cfd793e
3	011575986092c30523ecb71ff10cb473	ed0a6265b5b0a292cd537214f401f1fc
4	011b4adcd54683b480c4d841250a987f	93b35affb86435b71ca01a6dd1eab2b1
5	012452d40dafae4df401bcd74cdb490	e49d0607832db7c3b1343764225d8df6
6	012a218df8995d3ec3bb221828360c86	207258d556f9b9784b54de837ae49286
7	013ef03e0f3f408dd9bf555e4edcdc0a	00cee8a4dbb042385cbbd84699e78b2e
8	013f4353d26bb05dc6652f1269458d8d	9eead354475b8b8fde3811f26b2599fc
9	015557c9912277312b9073947804a7ba	48d9d778ca005971949d2ccad7e99d70

	target_order_ts	target_category	target_order_value \
0	2018-01-14 07:36:54	malas_acessorios	103.28
1	2018-06-03 10:12:57	eletronicos	26.78
2	2017-03-21 19:25:23	cama_mesa_banho	81.78
3	2018-04-18 21:58:08	instrumentos_musicais	80.32
4	2018-02-15 11:40:57	alimentos	149.88
5	2018-05-14 12:12:45	beleza_saude	340.07
6	2018-06-18 13:08:38	automotivo	1435.97
7	2018-07-20 04:13:54	instrumentos_musicais	69.07
8	2017-11-28 13:30:58	esporte_lazer	164.04
9	2017-05-01 14:48:33	utilidades_domesticas	72.90

	history_order_count
0	1
1	1
2	1
3	1
4	1
5	1
6	1
7	1
8	1
9	1

```

[59]: # =====
# CELL 13C - Build behavioral features from previous orders
# =====

```

```

# Add target timestamp to every history order
history_orders = (
    eligible_orders
    .merge(
        snapshots[
            [
                "customer_unique_id",
                "target_order_ts"
            ]
        ],
        on="customer_unique_id",
        how="inner"
    )
)

# Only information before the target order is allowed
history_orders = history_orders[
    history_orders["order_purchase_timestamp"]
    <
    history_orders["target_order_ts"]
].copy()

# -----
# Order-level features
# -----

order_features = (
    history_orders
    .groupby("customer_unique_id")
    .agg(
        prior_order_count=(
            "order_id",
            "nunique"
        ),

        first_history_purchase=(
            "order_purchase_timestamp",
            "min"
        ),

        last_history_purchase=(
            "order_purchase_timestamp",
            "max"
        ),

        total_spending=(

```

```

        "order_payment_value",
        "sum"
    ),

    average_order_value=(
        "order_payment_value",
        "mean"
    ),

    max_order_value=(
        "order_payment_value",
        "max"
    )
)
.reset_index()
)

# Recency:
# target date - last previous purchase
order_features = order_features.merge(
    snapshots[
        [
            "customer_unique_id",
            "target_order_ts"
        ]
    ],
    on="customer_unique_id",
    how="left"
)

order_features["recency_days"] = (
    order_features["target_order_ts"]
    -
    order_features["last_history_purchase"]
).dt.total_seconds() / 86400

order_features["history_duration_days"] = (
    order_features["last_history_purchase"]
    -
    order_features["first_history_purchase"]
).dt.total_seconds() / 86400

# -----
# Purchase interval statistics
# -----

history_sorted = (

```

```

history_orders
.sort_values(
    [
        "customer_unique_id",
        "order_purchase_timestamp"
    ]
)
.copy()
)

history_sorted["days_since_previous_order"] = (
    history_sorted
    .groupby("customer_unique_id")
    ["order_purchase_timestamp"]
    .diff()
    .dt.total_seconds()
    / 86400
)

interval_features = (
    history_sorted
    .groupby("customer_unique_id")
    .agg(
        average_days_between_orders=(
            "days_since_previous_order",
            "mean"
        ),

        min_days_between_orders=(
            "days_since_previous_order",
            "min"
        ),

        max_days_between_orders=(
            "days_since_previous_order",
            "max"
        )
    )
    .reset_index()
)

order_features = order_features.merge(
    interval_features,
    on="customer_unique_id",
    how="left"
)

```

```

[60]: # =====
# CELL 13D - Item-level behavioral features
# =====

history_items = (
    order_items_clean
    .merge(
        products_clean[
            [
                "product_id",
                "product_category_name"
            ]
        ],
        on="product_id",
        how="left"
    )
    .merge(
        history_orders[
            [
                "order_id",
                "customer_unique_id"
            ]
        ],
        on="order_id",
        how="inner"
    )
)

item_features = (
    history_items
    .groupby("customer_unique_id")
    .agg(
        total_items=(
            "order_item_id",
            "count"
        ),
        average_item_price=(
            "price",
            "mean"
        ),
        average_freight_value=(
            "freight_value",
            "mean"
        )
    )
)

```

```

        .reset_index()
    )

    # Correct customer-level ratio
    item_features = item_features.merge(
        order_features[
            [
                "customer_unique_id",
                "prior_order_count"
            ]
        ],
        on="customer_unique_id",
        how="left"
    )

    item_features["average_items_per_order"] = (
        item_features["total_items"]
        /
        item_features["prior_order_count"]
        .replace(0, np.nan)
    )

    item_features = item_features.drop(
        columns=["prior_order_count"]
    )

```

```

[61]: # =====
# CELL 13E - Category behavioral features
# =====

category_counts = (
    history_items
    .groupby(
        [
            "customer_unique_id",
            "product_category_name"
        ],
        dropna=False
    )
    .size()
    .reset_index(
        name="category_item_count"
    )
)

# Ignore unknown category for behavioral category features
category_counts = category_counts[

```

```

        category_counts[
            "product_category_name"
        ].notna()
    ].copy()

# Top categories based on historical purchase volume
top_behavior_categories = (
    category_counts[
        "product_category_name"
    ]
    .value_counts()
    .head(12)
    .index
    .tolist()
)

category_counts = category_counts[
    category_counts[
        "product_category_name"
    ].isin(
        top_behavior_categories
    )
].copy()

category_pivot = (
    category_counts
    .pivot(
        index="customer_unique_id",
        columns="product_category_name",
        values="category_item_count"
    )
    .fillna(0)
)

# Rename count columns
category_pivot.columns = [
    f"category_count_{name}"
    for name in category_pivot.columns
]

category_pivot = category_pivot.reset_index()

# -----
# Category share
# -----

count_columns = [

```

```

        column
    for column in category_pivot.columns
    if column.startswith("category_count__")
]

total_category_items = (
    category_pivot[count_columns]
    .sum(axis=1)
    .replace(0, 1)
)

for count_column in count_columns:

    category_name = count_column.replace(
        "category_count__",
        ""
    )

    share_column = (
        f"category_share_{category_name}"
    )

    category_pivot[share_column] = (
        category_pivot[count_column]
        / total_category_items
    )

```

```

[62]: # =====
# CELL 13F - Payment behavioral features
# =====

history_payments = (
    payments_clean
    .merge(
        history_orders[
            [
                "order_id",
                "customer_unique_id"
            ]
        ],
        on="order_id",
        how="inner"
    )
)

payment_features = (
    history_payments

```

```

.groupby("customer_unique_id")
.agg(
    average_installments=(
        "payment_installments",
        "mean"
    ),

    average_payment_value=(
        "payment_value",
        "mean"
    )
)
.reset_index()
)

# Dominant payment type
dominant_payment = (
    history_payments
    .groupby(
        [
            "customer_unique_id",
            "payment_type"
        ]
    )
    .size()
    .reset_index(
        name="payment_count"
    )
    .sort_values(
        [
            "customer_unique_id",
            "payment_count",
            "payment_type"
        ],
        ascending=[
            True,
            False,
            True
        ]
    )
    .drop_duplicates(
        subset=["customer_unique_id"]
    )
    [
        [
            "customer_unique_id",
            "payment_type"

```

```

    ]
]
.rename(
    columns={
        "payment_type":
            "dominant_payment_type"
    }
)
)

payment_features = payment_features.merge(
    dominant_payment,
    on="customer_unique_id",
    how="left"
)

```

```

[63]: # =====
# CELL 13G - Review behavior features
# =====

history_reviews = (
    reviews_clean[
        [
            "order_id",
            "review_score",
            "review_comment_message",
            "review_creation_date"
        ]
    ]
    .merge(
        history_orders[
            [
                "order_id",
                "customer_unique_id",
                "target_order_ts"
            ]
        ],
        on="order_id",
        how="inner"
    )
)

# Only reviews available before the target timestamp
history_reviews = history_reviews[
    history_reviews["review_creation_date"].isna()
    |
    (

```

```

        history_reviews["review_creation_date"]
        <=
        history_reviews["target_order_ts"]
    )
].copy()

history_reviews["review_text_clean"] = (
    history_reviews[
        "review_comment_message"
    ]
    .fillna("")
    .map(clean_text)
)

review_features = (
    history_reviews
    .groupby(
        "customer_unique_id"
    )
    .agg(
        average_review_score=(
            "review_score",
            "mean"
        ),
        review_count=(
            "order_id",
            "nunique"
        ),
        review_text=(
            "review_text_clean",
            lambda values:
                " ".join(
                    value
                    for value in values
                    if value
                )
        )
    )
    .reset_index()
)

review_features["text_review_count"] = (
    review_features["review_text"]
    .str.len()
    .gt(0)
)

```

```
        .astype(int)
    )
```

```
[64]: # =====
# CELL 13H - Customer-level profile
# =====

customer_state_features = (
    customers[
        [
            "customer_unique_id",
            "customer_state"
        ]
    ]
    .drop_duplicates(
        subset=["customer_unique_id"]
    )
)

customer_features = (
    snapshots[
        [
            "customer_unique_id",
            "target_order_id",
            "target_order_ts",
            "target_category",
            "target_order_value"
        ]
    ]
    .copy()
)

feature_tables = [
    order_features.drop(
        columns=["target_order_ts"]
    ),

    item_features,

    category_pivot,

    payment_features,

    review_features,

    customer_state_features
]
```

```

for table in feature_tables:

    customer_features = customer_features.merge(
        table,
        on="customer_unique_id",
        how="left"
    )

# Fill text and categorical missing values
customer_features["review_text"] = (
    customer_features["review_text"]
    .fillna("")
)

customer_features["customer_state"] = (
    customer_features["customer_state"]
    .fillna("unknown")
)

# Numerical missing values will be handled by
# the preprocessing pipeline.
# We do not manually fill them here.

print(
    "Customer-level feature dataset shape:",
    customer_features.shape
)

display(
    customer_features.head()
)

```

Customer-level feature dataset shape: (2801, 52)

	customer_unique_id	target_order_id \
0	004288347e5e88a27ded2bb23747066c	08204559bebd39e09ee52dcb56d8faa2
1	00a39521eb40f7012db50455bf083460	cea3e6c11eb60acb9d8d4d51694832f8
2	00cc12a6d8b578b8ebd21ea4e2ae8b27	d61b915b69851aec8a8865f36cfd793e
3	011575986092c30523ecb71ff10cb473	ed0a6265b5b0a292cd537214f401f1fc
4	011b4adcd54683b480c4d841250a987f	93b35affb86435b71ca01a6dd1eab2b1

	target_order_ts	target_category	target_order_value \
0	2018-01-14 07:36:54	malas_acessorios	103.28
1	2018-06-03 10:12:57	eletronicos	26.78
2	2017-03-21 19:25:23	cama_mesa_banho	81.78
3	2018-04-18 21:58:08	instrumentos_musicais	80.32
4	2018-02-15 11:40:57	alimentos	149.88

```

prior_order_count first_history_purchase last_history_purchase \
0          1.0      2017-07-27 14:13:03   2017-07-27 14:13:03
1          1.0      2018-05-23 20:14:21   2018-05-23 20:14:21
2          1.0      2017-03-21 19:25:22   2017-03-21 19:25:22
3          1.0      2018-02-17 15:54:49   2018-02-17 15:54:49
4          1.0      2017-08-22 12:51:29   2017-08-22 12:51:29

total_spending average_order_value ... category_share__telefonia \
0          251.09          251.09 ...                NaN
1           96.47           96.47 ...                0.0
2           44.42           44.42 ...                NaN
3          134.58          134.58 ...                0.0
4           86.42           86.42 ...                0.0

category_share__utilidades_domesticas average_installments \
0                                NaN                2.0
1                                0.0                1.0
2                                NaN                1.0
3                                0.0                2.0
4                                0.0                8.0

average_payment_value dominant_payment_type average_review_score \
0          251.09          credit_card          5.0
1           96.47          credit_card          NaN
2           44.42             boleto          NaN
3          134.58          credit_card          3.0
4           86.42          credit_card          5.0

review_count                                review_text \
0          1.0                                amei a bolsa ótima qualidade e lindas
1          NaN
2          NaN
3          1.0      produto foi entregue no prazo previsto ainda n...
4          1.0

text_review_count customer_state
0          1.0          RJ
1          NaN          DF
2          NaN          SP
3          1.0          MG
4          0.0          BA

```

[5 rows x 52 columns]

```

[65]: # =====
      # CELL 13I - Feature selection and leakage audit

```

```
# =====

IDENTIFIER_COLUMNS = [
    "customer_unique_id",
    "target_order_id",
    "target_order_ts"
]

TARGET_COLUMNS = [
    "target_category",
    "target_order_value"
]

TEXT_COLUMNS = [
    "review_text"
]

CATEGORICAL_FEATURES = [
    "customer_state",
    "dominant_payment_type"
]

CATEGORICAL_FEATURES = [
    column
    for column in CATEGORICAL_FEATURES
    if column in customer_features.columns
]

NUMERICAL_FEATURES = [
    column
    for column in customer_features.columns
    if (
        column not in IDENTIFIER_COLUMNS
        and column not in TARGET_COLUMNS
        and column not in TEXT_COLUMNS
        and column not in CATEGORICAL_FEATURES
        and pd.api.types.is_numeric_dtype(
            customer_features[column]
        )
    )
]

print("Numerical features:")
for feature in NUMERICAL_FEATURES:
    print(" -", feature)

print("\nCategorical features:")
```

```

for feature in CATEGORICAL_FEATURES:
    print(" -", feature)

print("\nText features:")
for feature in TEXT_COLUMNS:
    print(" -", feature)

print("\nExcluded identifiers:")
for feature in IDENTIFIER_COLUMNS:
    print(" -", feature)

print("\nTargets:")
for feature in TARGET_COLUMNS:
    print(" -", feature)

# -----
# Formal leakage audit
# -----

for column in [
    "target_category",
    "target_order_value",
    "target_order_ts",
    "target_order_id",
    "customer_unique_id"
]:
    if column in NUMERICAL_FEATURES + CATEGORICAL_FEATURES + TEXT_COLUMNS:
        raise ValueError(
            f"Potential leakage column in X: {column}"
        )

print(
    "\nLeakage audit passed: target and identifier columns "
    "are excluded from model input."
)

```

Numerical features:

- prior_order_count
- total_spending
- average_order_value
- max_order_value
- recency_days
- history_duration_days
- average_days_between_orders
- min_days_between_orders
- max_days_between_orders
- total_items
- average_item_price

- average_freight_value
- average_items_per_order
- category_count__beleza_saude
- category_count__brinquedos
- category_count__cama_mesa_banho
- category_count__esporte_lazer
- category_count__fashion_bolsas_e_acessorios
- category_count__ferramentas_jardim
- category_count__informatica_acessorios
- category_count__moveis_decoracao
- category_count__perfumaria
- category_count__relogios_presentes
- category_count__telefonia
- category_count__utilidades_domesticas
- category_share__beleza_saude
- category_share__brinquedos
- category_share__cama_mesa_banho
- category_share__esporte_lazer
- category_share__fashion_bolsas_e_acessorios
- category_share__ferramentas_jardim
- category_share__informatica_acessorios
- category_share__moveis_decoracao
- category_share__perfumaria
- category_share__relogios_presentes
- category_share__telefonia
- category_share__utilidades_domesticas
- average_installments
- average_payment_value
- average_review_score
- review_count
- text_review_count

Categorical features:

- customer_state
- dominant_payment_type

Text features:

- review_text

Excluded identifiers:

- customer_unique_id
- target_order_id
- target_order_ts

Targets:

- target_category
- target_order_value

Leakage audit passed: target and identifier columns are excluded from model input.

14 14. Train / Validation / Test Split

Assignment 02 requires a 70% / 15% / 15% split.

The data is divided into:

- 70% training
- 15% validation
- 15% testing

14.1 Purpose

14.1.1 Training set

Used to fit preprocessing and ML models.

14.1.2 Validation set

Used for:

- model comparison;
- model selection;
- experiment comparison.

14.1.3 Test set

Used only for final evaluation.

The test set is never used to:

- fit scalers;
- fit encoders;
- fit TF-IDF;
- select the final model.

14.2 Stratification

The target is multiclass and potentially imbalanced.

Rare categories are grouped into “other” for the purpose of stable stratified splitting.

This does not change the raw target permanently. It only prevents very small classes from making the split statistically unstable.

```
[66]: # =====  
# CELL 14 - 70 / 15 / 15 split  
# =====  
  
model_data = customer_features.copy()
```

```

# -----
# Prepare a stable target for splitting.
# Keep categories with enough observations.
# -----

raw_target_counts = (
    model_data["target_category"]
    .value_counts()
)

# A category should have enough samples for
# both splitting stages.
stable_categories = (
    raw_target_counts[
        raw_target_counts >= 4
    ]
    .head(10)
    .index
    .tolist()
)

model_data["split_target"] = (
    model_data["target_category"]
    .where(
        model_data["target_category"]
        .isin(stable_categories),
        "other"
    )
)

split_counts = (
    model_data["split_target"]
    .value_counts()
)

print("Stratification classes:")
display(
    split_counts.to_frame("count")
)

# -----
# First split:
# 70% train
# 30% temporary
# -----

train_df, temp_df = train_test_split(

```

```

    model_data,
    test_size=0.30,
    random_state=RANDOM_SEED,
    stratify=model_data["split_target"]
)

# -----
# Second split:
# 15% validation
# 15% test
# -----

val_df, test_df = train_test_split(
    temp_df,
    test_size=0.50,
    random_state=RANDOM_SEED,
    stratify=temp_df["split_target"]
)

# Remove helper column
train_df = train_df.drop(
    columns=["split_target"]
)

val_df = val_df.drop(
    columns=["split_target"]
)

test_df = test_df.drop(
    columns=["split_target"]
)

print("\nDataset sizes:")
print("Total:", len(model_data))
print("Train:", len(train_df))
print("Validation:", len(val_df))
print("Test:", len(test_df))

print("\nActual ratios:")

print(
    "Train:",
    round(
        len(train_df) / len(model_data),
        4
    )
)

```

```

print(
    "Validation:",
    round(
        len(val_df) / len(model_data),
        4
    )
)

print(
    "Test:",
    round(
        len(test_df) / len(model_data),
        4
    )
)

```

Stratification classes:

	count
split_target	
other	891
cama_mesa_banho	392
esporte_lazer	260
moveis_decoracao	257
beleza_saude	235
informatica_acessorios	198
utilidades_domesticas	164
relogios_presentes	135
fashion_bolsas_e_acessorios	99
brinquedos	89
automotivo	81

Dataset sizes:

Total: 2801

Train: 1960

Validation: 420

Test: 421

Actual ratios:

Train: 0.6998

Validation: 0.1499

Test: 0.1503

15. Preprocessing Pipeline

The preprocessing pipeline must be identical between training and inference.

15.1 Numerical features

Numerical preprocessing:

Missing value -> median imputation -> standard scaling

15.2 Categorical features

Categorical preprocessing:

Missing value -> most-frequent imputation -> One-Hot Encoding

15.3 Text features

Review comments:

Raw comment -> cleaned comment -> TF-IDF vector

15.4 Leakage Prevention

The preprocessing components are fitted only using training data.

Validation and test data are transformed using the already-fitted preprocessing pipeline.

The same saved preprocessing logic will later be used by:

- REST API
- Web application
- Mobile application

Therefore:

Training preprocessing = Deployment preprocessing

```
[67]: # =====  
# CELL 15 - Preprocessing utilities  
# =====  
  
# -----  
# X / y  
# -----  
  
MODEL_INPUT_COLUMNS = (  
    NUMERICAL_FEATURES  
    + CATEGORICAL_FEATURES  
    + TEXT_COLUMNS  
)  
  
X_train = train_df[  
    MODEL_INPUT_COLUMNS  
> ].copy()  
  
X_val = val_df[
```

```

    MODEL_INPUT_COLUMNS
].copy()

X_test = test_df[
    MODEL_INPUT_COLUMNS
].copy()

# -----
# Main target
# -----

# Determine the top target categories using TRAINING data only.
TARGET_TOP_K = 10

target_categories_train = (
    train_df["target_category"]
    .value_counts()
    .head(TARGET_TOP_K)
    .index
    .tolist()
)

def map_target_category(value):
    if value in target_categories_train:
        return value
    return "other"

y_train = (
    train_df["target_category"]
    .map(map_target_category)
)

y_val = (
    val_df["target_category"]
    .map(map_target_category)
)

y_test = (
    test_df["target_category"]
    .map(map_target_category)
)

INTEREST_LABELS = sorted(
    y_train.unique()
)

print("Final target classes:")

```

```

print(INTEREST_LABELS)

print("\nTraining target distribution:")
display(
    y_train.value_counts()
    .to_frame("count")
)

# -----
# OneHotEncoder compatibility
# -----

try:
    one_hot_encoder = OneHotEncoder(
        handle_unknown="ignore",
        sparse_output=True
    )
except TypeError:
    one_hot_encoder = OneHotEncoder(
        handle_unknown="ignore",
        sparse=True
    )

# Create a fresh encoder for each pipeline
def create_one_hot_encoder():
    try:
        return OneHotEncoder(
            handle_unknown="ignore",
            sparse_output=True
        )
    except TypeError:
        return OneHotEncoder(
            handle_unknown="ignore",
            sparse=True
        )

# -----
# Flatten text column
# -----

def flatten_text_column(X):
    return np.asarray(X).ravel()

# -----
# Preprocessor factory
# -----

```

```

def make_tabular_preprocessor(
    numeric_features,
    categorical_features
):

    numeric_pipeline = Pipeline([
        (
            "imputer",
            SimpleImputer(
                strategy="median"
            )
        ),
        (
            "scaler",
            StandardScaler()
        )
    ])

    categorical_pipeline = Pipeline([
        (
            "imputer",
            SimpleImputer(
                strategy="most_frequent"
            )
        ),
        (
            "onehot",
            create_one_hot_encoder()
        )
    ])

    return ColumnTransformer(
        transformers=[
            (
                "num",
                numeric_pipeline,
                numeric_features
            ),
            (
                "cat",
                categorical_pipeline,
                categorical_features
            )
        ],
        remainder="drop"
    )

```

```

def make_text_preprocessor(
    numeric_features,
    categorical_features
):

    tabular_preprocessor = (
        make_tabular_preprocessor(
            numeric_features,
            categorical_features
        )
    )

    text_pipeline = Pipeline([
        (
            "flatten",
            FunctionTransformer(
                flatten_text_column,
                validate=False
            )
        ),
        (
            "tfidf",
            TfidfVectorizer(
                max_features=1200,
                min_df=2,
                max_df=0.95,
                ngram_range=(1, 2),
                sublinear_tf=True
            )
        )
    ])

    return ColumnTransformer(
        transformers=[
            (
                "tabular",
                tabular_preprocessor,
                numeric_features
                + categorical_features
            ),
            (
                "text",
                text_pipeline,
                ["review_text"]
            )
        ]
    )

```

```

        )
    ],
    remainder="drop"
)

print(
    "Raw model input columns:",
    len(MODEL_INPUT_COLUMNS)
)

```

Final target classes:

```

['beleza_saude', 'brinquedos', 'cama_mesa_banho', 'esporte_lazer',
'fashion_bolsas_e_acessorios', 'informatica_acessorios', 'moveis_decoracao',
'other', 'relogios_presentes', 'telefonica', 'utilidades_domesticas']

```

Training target distribution:

target_category	count
other	622
cama_mesa_banho	274
esporte_lazer	182
moveis_decoracao	180
beleza_saude	164
informatica_acessorios	139
utilidades_domesticas	115
relogios_presentes	94
fashion_bolsas_e_acessorios	69
brinquedos	62
telefonica	59

Raw model input columns: 45

16. Baseline Model

Before training machine-learning algorithms, a simple baseline is required.

The baseline uses:

DummyClassifier(strategy="most_frequent")

This baseline does not learn meaningful relationships between customer features and the target. It simply predicts the most frequent class in the training data.

16.1 Why use a dummy baseline?

The purpose of the baseline is to establish the minimum reference performance.

A useful ML model should perform better than this trivial strategy.

Therefore:

Dummy Baseline -> reference performance

Machine Learning Models -> should improve upon the baseline

The DummyClassifier is NOT counted as one of the six required ML models.

The six required models will be trained and compared in the next section.

```
[68]: # =====  
# CELL 16A - Dummy baseline  
# =====  
  
from sklearn.dummy import DummyClassifier  
  
# -----  
# Dummy baseline  
# -----  
  
baseline_model = DummyClassifier(  
    strategy="most_frequent"  
)  
  
print("Training DummyClassifier baseline...")  
  
baseline_model.fit(  
    X_train,  
    y_train  
)  
  
baseline_prediction = (  
    baseline_model.predict(  
        X_val  
    )  
)  
  
baseline_accuracy = accuracy_score(  
    y_val,  
    baseline_prediction  
)  
  
baseline_precision = precision_score(  
    y_val,  
    baseline_prediction,  
    average="macro",  
    zero_division=0  
)
```

```

baseline_recall = recall_score(
    y_val,
    baseline_prediction,
    average="macro",
    zero_division=0
)

baseline_f1 = f1_score(
    y_val,
    baseline_prediction,
    average="macro",
    zero_division=0
)

baseline_results = pd.DataFrame([
    {
        "Model": "DummyClassifier",
        "Strategy": "most_frequent",
        "Accuracy": baseline_accuracy,
        "Precision Macro": baseline_precision,
        "Recall Macro": baseline_recall,
        "F1 Macro": baseline_f1
    }
])

display(
    baseline_results
)

print(
    "Baseline predicted class:",
    baseline_model.classes_[
        0
    ]
    if hasattr(
        baseline_model,
        "classes_"
    )
    else "N/A"
)

```

Training DummyClassifier baseline...

	Model	Strategy	Accuracy	Precision Macro	Recall Macro	\
0	DummyClassifier	most_frequent	0.319048	0.029004	0.090909	

	F1 Macro
0	0.043978

Baseline predicted class: beleza_saude

17 17. Model Training

Application 3 requires six machine-learning models to be compared.

The six models are:

1. Logistic Regression
2. Decision Tree
3. Random Forest
4. Support Vector Machine
5. Text-based Linear Classifier
6. Gradient Boosting Classifier

17.1 Representation

Models 1, 2, 3, 4 and 6 use structured customer behavioral features:

Customer history -> Customer profile -> Numerical + categorical features -> Preprocessing -> ML model

Model 5 uses customer review text:

Review comments -> Clean text -> TF-IDF -> Linear classifier

This distinction is important because the assignment explicitly asks the student to consider both behavioral/tabular representation and customer comments when text data is available.

17.2 Training policy

Each model is trained only once during this experiment.

The validation set is used for comparison.

The test set remains untouched until the final evaluation.

```
[69]: # =====  
# CELL 17A - Preprocessing utilities for Tabular + Text  
# =====  
  
# -----  
# Dense OneHotEncoder  
# -----  
  
def create_dense_one_hot_encoder():  
    try:  
        return OneHotEncoder(  
            handle_unknown="ignore",  
            sparse_output=False
```

```

    )

    except TypeError:

        return OneHotEncoder(
            handle_unknown="ignore",
            sparse=False
        )

# -----
# Tabular preprocessing
# -----

def make_tabular_preprocessor(
    numeric_features,
    categorical_features
):

    numeric_pipeline = Pipeline([
        (
            "imputer",
            SimpleImputer(
                strategy="median"
            )
        ),

        (
            "scaler",
            StandardScaler()
        )
    ])

    categorical_pipeline = Pipeline([
        (
            "imputer",
            SimpleImputer(
                strategy="most_frequent"
            )
        ),

        (
            "onehot",
            create_dense_one_hot_encoder()
        )
    ])

```

```

return ColumnTransformer([
    (
        "num",
        numeric_pipeline,
        numeric_features
    ),

    (
        "cat",
        categorical_pipeline,
        categorical_features
    )
])

# -----
# Text helper
# -----

def text_to_id(X):

    return (
        np.asarray(X)
        .ravel()
        .astype(str)
    )

# -----
# Combined Tabular + Text preprocessor
# -----

def make_combined_preprocessor():

    tabular_preprocessor = (
        make_tabular_preprocessor(
            NUMERICAL_FEATURES,
            CATEGORICAL_FEATURES
        )
    )

# -----
# Word-level TF-IDF
# -----

```

```

word_tfidf = TfidfVectorizer(
    analyzer="word",
    max_features=3000,
    min_df=1,
    max_df=0.98,
    ngram_range=(1, 2),
    sublinear_tf=True
)

# -----
# Character-level TF-IDF
# -----

char_tfidf = TfidfVectorizer(
    analyzer="char_wb",
    max_features=3000,
    min_df=1,
    max_df=0.98,
    ngram_range=(3, 5),
    sublinear_tf=True
)

# -----
# Convert one text column to 1-D
# Then build two text representations:
#
# word TF-IDF
# +
# character TF-IDF
# -----

word_pipeline = Pipeline([
    (
        "flatten",
        FunctionTransformer(
            text_to_1d,
            validate=False
        )
    ),
    (
        "tfidf",
        word_tfidf
    )
])

```

```

])

char_pipeline = Pipeline([
    (
        "flatten",
        FunctionTransformer(
            text_to_id,
            validate=False
        )
    ),

    (
        "tfidf",
        char_tfidf
    )
])

# -----
# Combine:
#
# Tabular
# + Word TF-IDF
# + Character TF-IDF
# -----

return ColumnTransformer(
    transformers=[

        (
            "tabular",
            tabular_preprocessor,
            NUMERICAL_FEATURES
            + CATEGORICAL_FEATURES
        ),

        (
            "word_text",
            word_pipeline,
            ["review_text"]
        ),

        (
            "char_text",
            char_pipeline,
            ["review_text"]
        )
    ]
)

```

```

        )
    ],

    remainder="drop"
)

print(
    "Combined preprocessing functions updated."
)

```

Combined preprocessing functions updated.

```

[70]: # =====
# CELL 17B - Prepare training, validation and test inputs
# =====

# -----
# Main tabular inputs
# -----

MODEL_INPUT_COLUMNS = (
    NUMERICAL_FEATURES
    + CATEGORICAL_FEATURES
    + TEXT_COLUMNS
)

X_train = train_df[
    MODEL_INPUT_COLUMNS
].copy()

X_val = val_df[
    MODEL_INPUT_COLUMNS
].copy()

X_test = test_df[
    MODEL_INPUT_COLUMNS
].copy()

# -----
# Text-only inputs
#
# Model 5:
# Review Text -> TF-IDF -> Linear Classifier
# -----

```

```

X_train_text = X_train[
    ["review_text"]
].copy()

X_val_text = X_val[
    ["review_text"]
].copy()

X_test_text = X_test[
    ["review_text"]
].copy()


# -----
# Main target
# -----

y_train = (
    train_df[
        "target_category"
    ]
    .map(
        map_target_category
    )
)

y_val = (
    val_df[
        "target_category"
    ]
    .map(
        map_target_category
    )
)

y_test = (
    test_df[
        "target_category"
    ]
    .map(
        map_target_category
    )
)


# -----
# Basic validation

```

```

# -----

print("=" * 70)
print("MODEL INPUT SUMMARY")
print("=" * 70)

print(
    "\nNumber of raw input columns:",
    len(MODEL_INPUT_COLUMNS)
)

print(
    "\nX_train shape:",
    X_train.shape
)

print(
    "X_val shape:",
    X_val.shape
)

print(
    "X_test shape:",
    X_test.shape
)

print(
    "\nX_train_text shape:",
    X_train_text.shape
)

print(
    "X_val_text shape:",
    X_val_text.shape
)

print(
    "X_test_text shape:",
    X_test_text.shape
)

# -----
# Target summary
# -----

print(

```

```

        "\nTarget classes:"
    )

    for index, label in enumerate(
        sorted(
            y_train.unique()
        ),
        start=1
    ):
        print(
            f"{index}. {label}"
        )

    print(
        "\nTraining target distribution:"
    )

    display(
        y_train
        .value_counts()
        .to_frame(
            "count"
        )
    )

# -----
# Missing values in model inputs
# -----

    print(
        "\nMissing values in X_train:"
    )

    missing_input_summary = (
        X_train
        .isna()
        .sum()
        .sort_values(
            ascending=False
        )
    )

    display(
        missing_input_summary[
            missing_input_summary > 0
        ]
    )

```

```

    ]
    .to_frame(
        "missing_count"
    )
)

# -----
# Text statistics
# -----

print(
    "\nReview text statistics:"
)

text_lengths = (
    X_train_text[
        "review_text"
    ]
    .fillna("")
    .astype(str)
    .str.len()
)

print(
    "Customers with non-empty review text:",
    (
        text_lengths > 0
    ).sum()
)

print(
    "Customers with empty review text:",
    (
        text_lengths == 0
    ).sum()
)

print(
    "Average review text length:",
    round(
        text_lengths.mean(),
        2
    )
)

print(

```

```
"\nCell 17B completed successfully."
)
```

```
=====
MODEL INPUT SUMMARY
=====
```

Number of raw input columns: 45

X_train shape: (1960, 45)

X_val shape: (420, 45)

X_test shape: (421, 45)

X_train_text shape: (1960, 1)

X_val_text shape: (420, 1)

X_test_text shape: (421, 1)

Target classes:

1. beleza_saude
2. brinquedos
3. cama_mesa_banho
4. esporte_lazer
5. fashion_bolsas_e_acessorios
6. informatica_acessorios
7. moveis_decoracao
8. other
9. relorios_presentes
10. telefonia
11. utilidades_domesticas

Training target distribution:

target_category	count
other	622
cama_mesa_banho	274
esporte_lazer	182
moveis_decoracao	180
beleza_saude	164
informatica_acessorios	139
utilidades_domesticas	115
relorios_presentes	94
fashion_bolsas_e_acessorios	69
brinquedos	62
telefonia	59

Missing values in X_train:

	missing_count
average_days_between_orders	1801
min_days_between_orders	1801
max_days_between_orders	1801
review_count	708
text_review_count	708
average_review_score	708
category_count__esporte_lazer	590
category_count__brinquedos	590
category_count__fashion_bolsas_e_acessorios	590
category_count__beleza_saude	590
category_count__cama_mesa_banho	590
category_count__relogios_presentes	590
category_count__perfumaria	590
category_share__esporte_lazer	590
category_share__cama_mesa_banho	590
category_share__brinquedos	590
category_share__fashion_bolsas_e_acessorios	590
category_share__perfumaria	590
category_share__ferramentas_jardim	590
category_share__informatica_acessorios	590
category_share__moveis_decoracao	590
category_share__utilidades_domesticas	590
category_share__telefonica	590
category_count__telefonica	590
category_count__utilidades_domesticas	590
category_share__beleza_saude	590
category_count__ferramentas_jardim	590
category_count__informatica_acessorios	590
category_count__moveis_decoracao	590
category_share__relogios_presentes	590
average_freight_value	151
average_items_per_order	151
total_items	151
history_duration_days	151
average_order_value	151
max_order_value	151
recency_days	151
prior_order_count	151
total_spending	151
average_item_price	151
average_payment_value	151
average_installments	151
dominant_payment_type	151

Review text statistics:

Customers with non-empty review text: 547

Customers with empty review text: 1413
Average review text length: 17.62

Cell 17B completed successfully.

```
[71]: # =====  
# CELL 17B - Text-only preprocessing  
#  
# Review Text  
# -> flatten  
# -> TF-IDF  
# =====  
  
def make_text_only_preprocessor():  
  
    text_pipeline = Pipeline([  
  
        (  
            "flatten",  
            FunctionTransformer(  
                text_to_1d,  
                validate=False  
            )  
        ),  
  
        (  
            "tfidf",  
            TfidfVectorizer(  
                analyzer="word",  
  
                # Keep a reasonably large vocabulary  
                max_features=3000,  
  
                # Keep terms that appear at least once  
                # in the training documents  
                min_df=1,  
  
                # Remove extremely common terms  
                max_df=0.98,  
  
                # Unigrams + bigrams  
                ngram_range=(1, 2),  
  
                # Helps reduce the effect of very frequent terms  
                sublinear_tf=True  
            )  
        )  
    ])
```

```

    ])

    return ColumnTransformer(
        transformers=[
            (
                "text",
                text_pipeline,
                ["review_text"]
            )
        ],
        remainder="drop"
    )

print(
    "Text-only preprocessing function created."
)

```

Text-only preprocessing function created.

```

[72]: # =====
# CELL 17C - Define six required models
# =====
from sklearn.ensemble import (
    RandomForestClassifier,
    GradientBoostingClassifier
)
interest_models = {}

# -----
# MODEL 1 - Logistic Regression
# -----

interest_models[
    "Logistic Regression"
] = Pipeline([
    (
        "preprocessor",
        make_tabular_preprocessor(
            NUMERICAL_FEATURES,
            CATEGORICAL_FEATURES
        )
    ),

    (
        "classifier",
        LogisticRegression(

```

```

        max_iter=3000,
        class_weight="balanced",
        random_state=RANDOM_SEED
    )
)
])

# -----
# MODEL 2 - Decision Tree
# -----

interest_models[
    "Decision Tree"
] = Pipeline([
    (
        "preprocessor",
        make_tabular_preprocessor(
            NUMERICAL_FEATURES,
            CATEGORICAL_FEATURES
        )
    ),

    (
        "classifier",
        DecisionTreeClassifier(
            max_depth=10,
            min_samples_leaf=5,
            class_weight="balanced",
            random_state=RANDOM_SEED
        )
    )
])

# -----
# MODEL 3 - Random Forest
# -----

interest_models[
    "Random Forest"
] = Pipeline([
    (
        "preprocessor",
        make_tabular_preprocessor(
            NUMERICAL_FEATURES,
            CATEGORICAL_FEATURES
        )
    ),

```

```

        (
            "classifier",
            RandomForestClassifier(
                n_estimators=400,
                max_depth=15,
                min_samples_leaf=2,
                class_weight="balanced_subsample",
                random_state=RANDOM_SEED,
                n_jobs=-1
            )
        )
    ])

# -----
# MODEL 4 - SVM
# -----

interest_models[
    "SVM"
] = Pipeline([
    (
        "preprocessor",
        make_tabular_preprocessor(
            NUMERICAL_FEATURES,
            CATEGORICAL_FEATURES
        )
    ),

    (
        "classifier",
        SVC(
            kernel="rbf",
            C=2.0,
            class_weight="balanced",
            probability=True,
            random_state=RANDOM_SEED
        )
    )
])

# -----
# MODEL 5 - Text-based Linear Classifier
#
# Review text -> TF-IDF -> Logistic Regression
# -----

```

```

interest_models[
    "Text-based Linear Classifier"
] = Pipeline([
    (
        "preprocessor",
        make_text_only_preprocessor()
    ),
    (
        "classifier",
        LogisticRegression(
            max_iter=3000,
            class_weight="balanced",
            random_state=RANDOM_SEED
        )
    )
])

# -----
# MODEL 6 - Gradient Boosting Classifier
# -----

interest_models[
    "Gradient Boosting"
] = Pipeline([
    (
        "preprocessor",
        make_tabular_preprocessor(
            NUMERICAL_FEATURES,
            CATEGORICAL_FEATURES
        )
    ),
    (
        "classifier",
        GradientBoostingClassifier(
            n_estimators=200,
            learning_rate=0.05,
            max_depth=3,
            min_samples_leaf=5,
            random_state=RANDOM_SEED
        )
    )
])

print(
    "Six required models:"

```

```

)

for index, model_name in enumerate(
    interest_models.keys(),
    start=1
):

    print(
        f"{index}. {model_name}"
    )

```

Six required models:

1. Logistic Regression
2. Decision Tree
3. Random Forest
4. SVM
5. Text-based Linear Classifier
6. Gradient Boosting

```

[73]: # =====
# CELL 17D - Train all six models exactly once
# =====

for model_name, model in interest_models.items():

    print("=" * 70)
    print(
        f"Training {model_name}"
    )
    print("=" * 70)

    if (
        model_name
        ==
        "Text-based Linear Classifier"
    ):

        model.fit(
            X_train_text,
            y_train
        )

    else:

        model.fit(
            X_train,
            y_train

```

```

    )

print(
    "\nTraining completed for all six required models."
)

```

```

=====
Training Logistic Regression
=====

```

```

=====
Training Decision Tree
=====

```

```

=====
Training Random Forest
=====

```

```

=====
Training SVM
=====

```

```

=====
Training Text-based Linear Classifier
=====

```

```

=====
Training Gradient Boosting
=====

```

```

Training completed for all six required models.

```

18. Model Comparison

The six required models are compared using the validation set.

The reported metrics are:

- Accuracy
- Precision
- Recall
- Macro F1
- Weighted F1
- ROC-AUC where appropriate

18.1 Primary metric

Macro F1 is the primary selection metric.

The target contains multiple product categories and may be imbalanced.

Macro F1 gives each class equal importance.

18.2 Baseline comparison

The DummyClassifier establishes the minimum reference performance.

A useful ML model should perform better than the dummy baseline.

The test set is not used for model selection.

```
[74]: # =====  
# CELL 18A - Evaluate six models on validation data  
# =====  
  
comparison_results = []  
  
for model_name, model in interest_models.items():  
  
    if (  
        model_name  
        ==  
        "Text-based Linear Classifier"  
    ):  
  
        evaluation_X = X_val_text  
  
    else:  
  
        evaluation_X = X_val  
  
    predictions = model.predict(  
        evaluation_X  
    )  
  
    result = {  
        "Model":  
            model_name,  
  
        "Accuracy":  
            accuracy_score(  
                y_val,  
                predictions  
            ),  
  
        "Precision Macro":  
            precision_score(  
                y_val,  
                predictions,  
                average="macro",  
                zero_division=0  
            ),  
    }
```

```

    "Recall Macro":
        recall_score(
            y_val,
            predictions,
            average="macro",
            zero_division=0
        ),

    "F1 Macro":
        f1_score(
            y_val,
            predictions,
            average="macro",
            zero_division=0
        ),

    "F1 Weighted":
        f1_score(
            y_val,
            predictions,
            average="weighted",
            zero_division=0
        )
}

if hasattr(
    model,
    "predict_proba"
):

    try:

        probabilities = (
            model
            .predict_proba(
                evaluation_X
            )
        )

        result[
            "ROC-AUC OVR Macro"
        ] = roc_auc_score(
            y_val,
            probabilities,
            multi_class="ovr",
            average="macro"

```

```

    )

    except Exception:

        result[
            "ROC-AUC OVR Macro"
        ] = np.nan

    else:

        result[
            "ROC-AUC OVR Macro"
        ] = np.nan

    comparison_results.append(
        result
    )

validation_results = (
    pd.DataFrame(
        comparison_results
    )
    .sort_values(
        "F1 Macro",
        ascending=False
    )
    .reset_index(drop=True)
)

display(
    validation_results
)

```

	Model	Accuracy	Precision Macro	Recall Macro	\
0	Random Forest	0.497619	0.485873	0.461863	
1	SVM	0.488095	0.470895	0.449461	
2	Logistic Regression	0.442857	0.411063	0.432246	
3	Decision Tree	0.464286	0.441558	0.403759	
4	Gradient Boosting	0.454762	0.436318	0.350934	
5	Text-based Linear Classifier	0.111905	0.120359	0.112756	

	F1 Macro	F1 Weighted	ROC-AUC OVR Macro
0	0.470516	0.494243	0.748889
1	0.454889	0.488049	0.714408
2	0.413791	0.448854	0.734432
3	0.408025	0.465292	0.709791
4	0.379435	0.438258	0.742687
5	0.064915	0.072338	0.521204

```

[75]: # =====
# CELL 18B - Add baseline to comparison visualization
# =====

baseline_comparison = pd.DataFrame([
    {
        "Model": "DummyClassifier",
        "F1 Macro": baseline_f1
    }
])

six_model_f1 = validation_results[
    [
        "Model",
        "F1 Macro"
    ]
]

comparison_plot_df = pd.concat(
    [
        baseline_comparison,
        six_model_f1
    ],
    ignore_index=True
)

plt.figure(figsize=(12, 6))

plt.bar(
    comparison_plot_df[
        "Model"
    ],
    comparison_plot_df[
        "F1 Macro"
    ]
)

plt.title(
    "Validation Macro F1 - Baseline vs Six ML Models"
)

plt.xlabel(
    "Model"
)

plt.ylabel(
    "Macro F1"
)

```

```


)

plt.ylim(
    0,
    1
)

plt.xticks(
    rotation=30,
    ha="right"
)

plt.tight_layout()
plt.show()

```



```

[76]: # =====
# CELL 18C - Compare all major metrics
# =====

metric_columns = [
    "Accuracy",
    "Precision Macro",
    "Recall Macro",
    "F1 Macro"
]

display(
    validation_results[

```

```

        [
            "Model"
        ]
        + metric_columns
    ]
)

# -----
# Highlight best model
# -----

best_model_name = (
    validation_results
    .iloc[0]
    [
        "Model"
    ]
)

print(
    "Best model according to validation Macro F1:"
)

print(
    best_model_name
)

print(
    "\nDummy baseline Macro F1:",
    round(
        baseline_f1,
        4
    )
)

print(
    "Best ML model Macro F1:",
    round(
        validation_results
        .iloc[0]
        [
            "F1 Macro"
        ],
        4
    )
)

```

```

print(
    "\nImprovement over baseline:",
    round(
        validation_results
        .iloc[0]
        [
            "F1 Macro"
        ]
        - baseline_f1,
        4
    )
)

```

	Model	Accuracy	Precision Macro	Recall Macro	\
0	Random Forest	0.497619	0.485873	0.461863	
1	SVM	0.488095	0.470895	0.449461	
2	Logistic Regression	0.442857	0.411063	0.432246	
3	Decision Tree	0.464286	0.441558	0.403759	
4	Gradient Boosting	0.454762	0.436318	0.350934	
5	Text-based Linear Classifier	0.111905	0.120359	0.112756	

```

F1 Macro
0 0.470516
1 0.454889
2 0.413791
3 0.408025
4 0.379435
5 0.064915

```

Best model according to validation Macro F1:
Random Forest

Dummy baseline Macro F1: 0.044
Best ML model Macro F1: 0.4705

Improvement over baseline: 0.4265

19 19. Evaluation

The selected model is now evaluated on the untouched test set.

19.1 Accuracy

Accuracy measures the overall percentage of correct classifications.

A high accuracy indicates that many customer categories were predicted correctly, but it may hide poor performance on smaller classes.

19.2 Precision

Precision answers:

When the model predicts a customer interest category, how often is the prediction correct?

For recommendation-oriented applications, precision matters because incorrect interest predictions can lead to irrelevant product recommendations.

19.3 Recall

Recall answers:

Of customers who actually belong to a category, how many were detected?

High recall helps discover more potential customer interests.

19.4 F1-score

F1 balances precision and recall.

Macro F1 is the primary metric because each product-interest category receives equal importance.

19.5 ROC-AUC

Multiclass One-vs-Rest ROC-AUC measures how effectively the model separates each category from the remaining categories.

19.6 Confusion Matrix

The confusion matrix identifies specific category-level errors.

It helps explain which product categories are difficult to distinguish.

```
[77]: # =====  
# CELL 19A - Final test evaluation  
# =====  
  
selected_interest_model = (  
    interest_models[  
        best_model_name  
    ]  
)  
  
if (  
    best_model_name  
    ==  
    "Text-based Linear Classifier"  
):  
  
    evaluation_X_test = X_test_text
```

```

else:

    evaluation_X_test = X_test

final_test_prediction = (
    selected_interest_model
    .predict(
        evaluation_X_test
    )
)

final_test_metrics = {
    "Accuracy":
        accuracy_score(
            y_test,
            final_test_prediction
        ),

    "Precision Macro":
        precision_score(
            y_test,
            final_test_prediction,
            average="macro",
            zero_division=0
        ),

    "Recall Macro":
        recall_score(
            y_test,
            final_test_prediction,
            average="macro",
            zero_division=0
        ),

    "F1 Macro":
        f1_score(
            y_test,
            final_test_prediction,
            average="macro",
            zero_division=0
        ),

    "F1 Weighted":
        f1_score(
            y_test,
            final_test_prediction,

```

```

        average="weighted",
        zero_division=0
    )
}

final_test_probability = None

if hasattr(
    selected_interest_model,
    "predict_proba"
):

    try:

        final_test_probability = (
            selected_interest_model
            .predict_proba(
                evaluation_X_test
            )
        )

        final_test_metrics[
            "ROC-AUC OVR Macro"
        ] = roc_auc_score(
            y_test,
            final_test_probability,
            multi_class="ovr",
            average="macro"
        )

    except Exception:

        final_test_metrics[
            "ROC-AUC OVR Macro"
        ] = np.nan

else:

    final_test_metrics[
        "ROC-AUC OVR Macro"
    ] = np.nan

display(
    pd.DataFrame(
        [final_test_metrics]
    )
)

```

	Accuracy	Precision Macro	Recall Macro	F1 Macro	F1 Weighted \
0	0.470309	0.420042	0.393125	0.401735	0.467071

	ROC-AUC	OVR Macro
0		0.72624

```
[78]: # =====
# CELL 19C - Confusion matrix with actual numbers
# =====

test_labels = sorted(
    y_train.unique()
)

cm = confusion_matrix(
    y_test,
    final_test_prediction,
    labels=test_labels
)

fig, ax = plt.subplots(
    figsize=(11, 9)
)

image = ax.imshow(
    cm,
    interpolation="nearest"
)

ax.set_title(
    "Confusion Matrix - Final Interest Model"
)

ax.set_xlabel(
    "Predicted Category"
)

ax.set_ylabel(
    "Actual Category"
)

ax.set_xticks(
    range(len(test_labels))
)

ax.set_yticks(
    range(len(test_labels))
)
```

```

)

ax.set_xticklabels(
    test_labels,
    rotation=90,
    fontsize=8
)


ax.set_yticklabels(
    test_labels,
    fontsize=8
)

# Display numbers in cells
for i in range(
    cm.shape[0]
):
    for j in range(
        cm.shape[1]
    ):
        ax.text(
            j,
            i,
            str(
                cm[i, j]
            ),
            ha="center",
            va="center",
            fontsize=8
        )

fig.colorbar(
    image,
    ax=ax
)

plt.tight_layout()
plt.show()

```



```
[79]: # =====
# CELL 19D - Per-class evaluation
# =====

classification_dict = (
    classification_report(
        y_test,
        final_test_prediction,
        labels=test_labels,
        output_dict=True,
        zero_division=0
    )
)

per_class_metrics = (
```

```

pd.DataFrame(
    classification_dict
)
.T
)

display(
    per_class_metrics.loc[
        test_labels,
        [
            "precision",
            "recall",
            "f1-score",
            "support"
        ]
    ]
    .sort_values(
        "f1-score",
        ascending=False
    )
)

```

	precision	recall	f1-score	support
beleza_saude	0.628571	0.611111	0.619718	36.0
cama_mesa_banho	0.600000	0.610169	0.605042	59.0
other	0.477419	0.544118	0.508591	136.0
esporte_lazer	0.463415	0.487179	0.475000	39.0
relogios_presentes	0.466667	0.350000	0.400000	20.0
moveis_decoracao	0.394737	0.384615	0.389610	39.0
fashion_bolsas_e_acessorios	0.454545	0.333333	0.384615	15.0
informatica_acessorios	0.450000	0.300000	0.360000	30.0
utilidades_domesticas	0.315789	0.250000	0.279070	24.0
telefonica	0.187500	0.300000	0.230769	10.0
brinquedos	0.181818	0.153846	0.166667	13.0

20. Error Analysis

Aggregate metrics do not explain why a prediction fails.

This section analyzes:

1. Category-level performance.
2. The most frequent confusion pairs.
3. Prediction accuracy according to the amount of customer history.
4. Whether customers with very limited history are harder to classify.

20.1 Interpretation

A customer with more previous orders can provide stronger evidence about:

- preferred categories;
- purchase frequency;
- spending behavior;
- category diversity;
- review behavior.

A customer with very little history provides much less information.

Therefore, history length itself can be important when interpreting model errors.

The objective of error analysis is not simply to count errors, but to understand what kind of information the model is missing.

```
[80]: # =====
# CELL 20A - Validation error analysis
#
# Uses the best model from Experiment 1
# (six required models - tabular/text baseline comparison)
# =====

# -----
# Select best model from Experiment 1
# -----

best_experiment1_model_name = (
    validation_results
        .iloc[0]["Model"]
)

best_experiment1_model = (
    interest_models[
        best_experiment1_model_name
    ]
)

print(
    "Model used for validation error analysis:",
    best_experiment1_model_name
)

# -----
# Select correct validation input
#
# Text-based Linear Classifier uses X_val_text.
# Other models use X_val.
# -----

if (
```

```

    best_experiment1_model_name
    ==
    "Text-based Linear Classifier"
):

    error_analysis_X = X_val_text

else:

    error_analysis_X = X_val

# -----
# Generate validation predictions
# -----

validation_prediction = (
    best_experiment1_model
    .predict(
        error_analysis_X
    )
)

# -----
# Build error-analysis dataframe
# -----

error_df = pd.DataFrame({

    "actual_class":
        y_val.reset_index(
            drop=True
        ),

    "predicted_class":
        pd.Series(
            validation_prediction
        ).reset_index(
            drop=True
        )
})

# -----
# Correct / incorrect
# -----

```

```

error_df["correct"] = (
    error_df["actual_class"]
    ==
    error_df["predicted_class"]
)

# -----
# Overall accuracy
# -----

overall_accuracy = (
    error_df["correct"].mean()
)

print(
    "\nValidation accuracy:",
    round(
        overall_accuracy,
        4
    )
)

# -----
# Category-level performance
# -----

category_performance = (
    error_df
    .groupby(
        "actual_class"
    )
    .agg(
        customers=(
            "correct",
            "size"
        ),
        correct_predictions=(
            "correct",
            "sum"
        ),
        accuracy=(

```

```

        "correct",
        "mean"
    )
)

.sort_values(
    "accuracy",
    ascending=False
)

display(
    category_performance
)

```

Model used for validation error analysis: Random Forest

Validation accuracy: 0.4976

	customers	correct_predictions	accuracy
actual_class			
other	134	80	0.597015
informatica_acessorios	29	17	0.586207
telefonica	11	6	0.545455
beleza_saude	35	19	0.542857
esporte_lazer	39	21	0.538462
relogios_presentes	21	10	0.476190
cama_mesa_banho	59	26	0.440678
utilidades_domesticas	25	10	0.400000
brinquedos	14	5	0.357143
fashion_bolsas_e_acessorios	15	5	0.333333
moveis_decoracao	38	10	0.263158

```

[81]: # =====
# CELL 20B - Most common confusion pairs
# =====

confusion_pairs = (
    error_df[
        ~error_df["correct"]
    ]
    .groupby(
        [
            "actual_class",
            "predicted_class"
        ]
    )
)

```

```

.size()
.reset_index(
    name="count"
)
.sort_values(
    "count",
    ascending=False
)
)

print(
    "Most frequent incorrect category pairs:"
)

display(
    confusion_pairs.head(20)
)

```

Most frequent incorrect category pairs:

	actual_class	predicted_class	count
13	cama_mesa_banho	other	19
35	moveis_decoracao	other	16
20	esporte_lazer	other	13
40	other	esporte_lazer	12
33	moveis_decoracao	cama_mesa_banho	9
37	other	beleza_saude	8
12	cama_mesa_banho	moveis_decoracao	8
30	informatica_acessorios	other	7
60	utilidades_domesticas	other	7
39	other	cama_mesa_banho	6
43	other	moveis_decoracao	6
3	beleza_saude	other	6
41	other	fashion_bolsas_e_acessorios	5
7	brinquedos	other	5
25	fashion_bolsas_e_acessorios	other	5
50	relogios_presentes	other	5
46	other	utilidades_domesticas	4
45	other	telefonica	4
0	beleza_saude	cama_mesa_banho	3
38	other	brinquedos	3

```

[82]: # =====
# CELL 20C - Prediction performance by customer history length
# =====

history_analysis = (
    test_df.iloc[

```

```

        :len(X_val)
    ][
        ["prior_order_count"]
    ]
    .reset_index(
        drop=True
    )
)

history_analysis["correct"] = (
    error_df["correct"]
)

def history_group(value):

    if value == 1:
        return "1 prior order"

    if value == 2:
        return "2 prior orders"

    if value <= 4:
        return "3-4 prior orders"

    return "5+ prior orders"

history_analysis["history_group"] = (
    history_analysis[
        ["prior_order_count"]
    ]
    .map(
        history_group
    )
)

history_performance = (
    history_analysis
    .groupby(
        ["history_group"]
    )
    .agg(
        customers=(
            "correct",
            "size"
        ),

```

```

        correct_predictions=(
            "correct",
            "sum"
        ),

        accuracy=(
            "correct",
            "mean"
        )
    )
    .reset_index()
)

print(
    "Prediction performance by customer history length:"
)

display(
    history_performance
)

```

Prediction performance by customer history length:

	history_group	customers	correct_predictions	accuracy
0	1 prior order	344	176	0.511628
1	2 prior orders	22	11	0.500000
2	3-4 prior orders	5	2	0.400000
3	5+ prior orders	49	20	0.408163

21 21. Model Selection

The final model decision is based on the validation set.

The test set is not used to select the final model.

21.1 Model Selection Criteria

The following factors are considered:

1. Macro F1
2. Precision / Recall balance
3. ROC-AUC where appropriate
4. Robustness
5. Computational cost
6. Probability/confidence availability
7. Deployment simplicity

21.2 Experiment 1 - Six Model Comparison

The six required models are:

1. Logistic Regression
2. Decision Tree
3. Random Forest
4. SVM
5. Text-based Linear Classifier
6. Gradient Boosting

21.3 Experiment 2 - Customer Behavior Prediction

A separate behavior classifier predicts the expected value segment of the customer's target order:

- Low
- Medium
- High

This provides the second application output:

Predicted Customer Behavior

21.4 Experiment 3 - Text Enhancement

The same Logistic Regression algorithm is tested with:

A. Tabular customer behavior features

versus

B. Tabular customer behavior features + review text

This experiment answers an important question:

Does customer review text improve prediction beyond structured behavioral features?

If the combined representation performs better, it can be selected as the final representation.

```
[83]: # =====
# CELL 21A - EXPERIMENT 2
# Customer Behavior Prediction
#
# Target:
# Low / Medium / High target-order value
# =====

# -----
# Thresholds are learned from TRAINING DATA ONLY.
# -----

low_threshold, high_threshold = (
    train_df[
        "target_order_value"
    ]
    .quantile(
        [1 / 3, 2 / 3]
```

```

    )
    .tolist()
)

def make_behavior_target(values):

    return pd.cut(
        values,

        bins=[
            -np.inf,
            low_threshold,
            high_threshold,
            np.inf
        ],

        labels=[
            "Low",
            "Medium",
            "High"
        ]
    ).astype(str)

behavior_y_train = make_behavior_target(
    train_df[
        "target_order_value"
    ]
)

behavior_y_val = make_behavior_target(
    val_df[
        "target_order_value"
    ]
)

behavior_y_test = make_behavior_target(
    test_df[
        "target_order_value"
    ]
)

# -----
# Behavior model
# -----

```

```

behavior_model = Pipeline([
    (
        "preprocessor",
        make_tabular_preprocessor(
            NUMERICAL_FEATURES,
            CATEGORICAL_FEATURES
        )
    ),
    (
        "classifier",
        RandomForestClassifier(
            n_estimators=400,
            max_depth=12,
            min_samples_leaf=3,
            class_weight="balanced",
            random_state=RANDOM_SEED,
            n_jobs=-1
        )
    )
])

print(
    "Training auxiliary behavior model..."
)

behavior_model.fit(
    X_train,
    behavior_y_train
)

behavior_val_prediction = (
    behavior_model.predict(
        X_val
    )
)

behavior_test_prediction = (
    behavior_model.predict(
        X_test
    )
)

behavior_validation_results = {
    "Dataset":
    "Validation",

```

```

    "Accuracy":
        accuracy_score(
            behavior_y_val,
            behavior_val_prediction
        ),

    "Precision Macro":
        precision_score(
            behavior_y_val,
            behavior_val_prediction,
            average="macro",
            zero_division=0
        ),

    "Recall Macro":
        recall_score(
            behavior_y_val,
            behavior_val_prediction,
            average="macro",
            zero_division=0
        ),

    "F1 Macro":
        f1_score(
            behavior_y_val,
            behavior_val_prediction,
            average="macro",
            zero_division=0
        )
}

behavior_test_results = {
    "Dataset":
        "Test",

    "Accuracy":
        accuracy_score(
            behavior_y_test,
            behavior_test_prediction
        ),

    "Precision Macro":
        precision_score(
            behavior_y_test,
            behavior_test_prediction,
            average="macro",
            zero_division=0
        )
}

```

```

    ),

    "Recall Macro":
        recall_score(
            behavior_y_test,
            behavior_test_prediction,
            average="macro",
            zero_division=0
        ),

    "F1 Macro":
        f1_score(
            behavior_y_test,
            behavior_test_prediction,
            average="macro",
            zero_division=0
        )
}

display(
    pd.DataFrame([
        behavior_validation_results,
        behavior_test_results
    ])
)

```

Training auxiliary behavior model...

	Dataset	Accuracy	Precision Macro	Recall Macro	F1 Macro
0	Validation	0.461905	0.466505	0.461910	0.463508
1	Test	0.470309	0.476710	0.470127	0.472345

```

[84]: # =====
# CELL 21B - EXPERIMENT 3
# Compare multiple classifiers using:
#
# Tabular Features + Review Text
# =====

# -----
# Text helper
# -----

def text_to_id(X):
    return (
        np.asarray(X)
        .ravel()
        .astype(str)
    )

```

```

)

# -----
# Combined preprocessing
# -----

def make_combined_preprocessor():

    # =====
    # TABULAR
    # =====

    tabular_preprocessor = (
        make_tabular_preprocessor(
            NUMERICAL_FEATURES,
            CATEGORICAL_FEATURES
        )
    )

    # =====
    # TEXT
    # =====

    # Word-level representation
    word_text_pipeline = Pipeline([

        (
            "flatten",
            FunctionTransformer(
                text_to_id,
                validate=False
            )
        ),

        (
            "tfidf",
            TfidfVectorizer(
                analyzer="word",
                max_features=3000,
                min_df=1,
                max_df=0.98,
                ngram_range=(1, 2),
                sublinear_tf=True
            )
        )
    ])

```

```

])

# Character-level representation
char_text_pipeline = Pipeline([

    (
        "flatten",
        FunctionTransformer(
            text_to_1d,
            validate=False
        )
    ),

    (
        "tfidf",
        TfidfVectorizer(
            analyzer="char_wb",
            max_features=3000,
            min_df=1,
            max_df=0.98,
            ngram_range=(3, 5),
            sublinear_tf=True
        )
    )
])

# =====
# COMBINE TABULAR + TEXT
# =====

return ColumnTransformer(

    transformers=[

        (
            "tabular",
            tabular_preprocessor,
            NUMERICAL_FEATURES
            + CATEGORICAL_FEATURES
        ),

        (
            "word_text",
            word_text_pipeline,
            ["review_text"]
        )
    ]
)

```

```

        ),
        (
            "char_text",
            char_text_pipeline,
            ["review_text"]
        )
    ],
    remainder="drop"
)

# =====
# CLASSIFIERS
# =====

combined_classifier_templates = {

    "Logistic Regression + Text":

        LogisticRegression(
            max_iter=3000,
            class_weight="balanced",
            random_state=RANDOM_SEED
        ),

    "Decision Tree + Text":

        DecisionTreeClassifier(
            max_depth=10,
            min_samples_leaf=5,
            class_weight="balanced",
            random_state=RANDOM_SEED
        ),

    "Random Forest + Text":

        RandomForestClassifier(
            n_estimators=400,
            max_depth=15,
            min_samples_leaf=2,
            class_weight="balanced_subsample",
            random_state=RANDOM_SEED,
            n_jobs=-1
        )
    }

```

```

    ),

    "SVM + Text":

        SVC(
            kernel="rbf",
            C=2.0,
            class_weight="balanced",
            probability=True,
            random_state=RANDOM_SEED
        )
}

# =====
# TRAIN + VALIDATE
# =====

combined_models_fitted = {}

combined_results = []

for model_name, classifier in (
    combined_classifier_templates.items()
):

    print("=" * 80)

    print(
        f"Training {model_name}"
    )

    print("=" * 80)

    combined_pipeline = Pipeline([

        (
            "preprocessor",
            make_combined_preprocessor()
        ),

        (
            "classifier",
            classifier

```

```

    )
])

# -----
# TRAIN ONLY ON TRAIN SET
# -----

combined_pipeline.fit(
    X_train,
    y_train
)

combined_models_fitted[
    model_name
] = combined_pipeline

# -----
# VALIDATION
# -----

predictions = (
    combined_pipeline
    .predict(
        X_val
    )
)

probabilities = (
    combined_pipeline
    .predict_proba(
        X_val
    )
)

# -----
# METRICS
# -----

result = {
    "Model":
        model_name,

```

```

"Representation":
    "Tabular + Review Text",

"Accuracy":
    accuracy_score(
        y_val,
        predictions
    ),

"Precision Macro":
    precision_score(
        y_val,
        predictions,
        average="macro",
        zero_division=0
    ),

"Recall Macro":
    recall_score(
        y_val,
        predictions,
        average="macro",
        zero_division=0
    ),

"F1 Macro":
    f1_score(
        y_val,
        predictions,
        average="macro",
        zero_division=0
    ),

"F1 Weighted":
    f1_score(
        y_val,
        predictions,
        average="weighted",
        zero_division=0
    ),

"ROC-AUC OVR Macro":
    roc_auc_score(
        y_val,
        probabilities,
        multi_class="ovr",
        average="macro"
    )

```

```

    )
}

combined_results.append(
    result
)

# =====
# RESULT TABLE
# =====

combined_validation_results = (
    pd.DataFrame(
        combined_results
    )
    .sort_values(
        "F1 Macro",
        ascending=False
    )
    .reset_index(
        drop=True
    )
)

print(
    "\nCombined Tabular + Review Text models:"
)

display(
    combined_validation_results
)

```

```

=====
Training Logistic Regression + Text
=====
=====
Training Decision Tree + Text
=====
=====
Training Random Forest + Text
=====
=====
Training SVM + Text
=====

```

Combined Tabular + Review Text models:

	Model	Representation	Accuracy \
0	SVM + Text	Tabular + Review Text	0.488095
1	Random Forest + Text	Tabular + Review Text	0.447619
2	Logistic Regression + Text	Tabular + Review Text	0.450000
3	Decision Tree + Text	Tabular + Review Text	0.464286

	Precision Macro	Recall Macro	F1 Macro	F1 Weighted	ROC-AUC	OVR Macro
0	0.469716	0.449461	0.454386	0.488088		0.715368
1	0.433812	0.459676	0.436428	0.454152		0.734700
2	0.421968	0.447935	0.427673	0.454855		0.734434
3	0.438752	0.403759	0.406059	0.465069		0.709455

```
[85]: # =====
# CELL 21C - FINAL MODEL SELECTION
# =====

print("=" * 80)
print("FINAL MODEL SELECTION")
print("=" * 80)

# =====
# EXPERIMENT 1
# Best Tabular-only model
# =====

best_tabular_model_name = (
    validation_results
    .iloc[0]["Model"]
)

best_tabular_f1 = (
    validation_results
    .iloc[0]["F1 Macro"]
)

# =====
# EXPERIMENT 2
# Best Tabular + Review Text model
# =====

best_combined_model_name = (
    combined_validation_results
    .iloc[0]["Model"]
)
```

```

best_combined_f1 = (
    combined_validation_results
        .iloc[0]["F1 Macro"]
)

# =====
# SHOW COMPARISON
# =====

print(
    "\nBest Tabular-only model:"
)

print(
    best_tabular_model_name
)

print(
    "Validation Macro F1:",
    round(
        best_tabular_f1,
        4
    )
)

print(
    "\nBest Tabular + Review Text model:"
)

print(
    best_combined_model_name
)

print(
    "Validation Macro F1:",
    round(
        best_combined_f1,
        4
    )
)

# =====
# FINAL DECISION

```

```

#
# The final model MUST use:
# Tabular + Review Text
# =====

FINAL_MODEL_NAME = (
    best_combined_model_name
)

FINAL_REPRESENTATION = (
    "Tabular + Review Text"
)

FINAL_VALIDATION_F1 = (
    best_combined_f1
)

FINAL_SOURCE = (
    "Experiment 2 - Tabular + Review Text"
)

print(
    "\n" + "-" * 60
)

print(
    "FINAL DEPLOYMENT DECISION"
)

print(
    "-" * 60
)

print(
    "Final model:",
    FINAL_MODEL_NAME
)

print(
    "Representation:",
    FINAL_REPRESENTATION
)

print(
    "Validation Macro F1:",
    round(

```

```

        FINAL_VALIDATION_F1,
        4
    )
)

print(
    "Selection source:",
    FINAL_SOURCE
)

```

FINAL MODEL SELECTION

Best Tabular-only model:
Random Forest
Validation Macro F1: 0.4705

Best Tabular + Review Text model:
SVM + Text
Validation Macro F1: 0.4544

FINAL DEPLOYMENT DECISION

Final model: SVM + Text
Representation: Tabular + Review Text
Validation Macro F1: 0.4544
Selection source: Experiment 2 - Tabular + Review Text

22. Model Persistence

The final intelligent system must be reusable without retraining.

After the model and representation have been selected, Train and Validation are combined.

The final selected model is then fitted exactly once.

The Test set remains untouched during this final fitting.

22.1 Saved components

The deployment package contains:

1. Interest prediction pipeline
2. Interest preprocessing component
3. Customer behavior pipeline
4. Customer behavior preprocessing
5. Recommendation catalog
6. Feature metadata

7. Target category configuration
8. Behavior thresholds

The main pipeline includes preprocessing and the classifier so that inference applies the same transformation used during training.

This ensures:

23 Training preprocessing

Deployment preprocessing

No scaler, encoder, imputer, or TF-IDF vocabulary is fitted using user input.

```
[86]: # =====
# CELL 22A - FINAL TRAIN + VALIDATION REFIT
# =====

# =====
# COMBINE TRAIN + VALIDATION
# =====

train_val_df = pd.concat(
    [
        train_df,
        val_df
    ],
    ignore_index=True
)

# =====
# FINAL INPUT
# =====

X_train_val = (
    train_val_df[
        MODEL_INPUT_COLUMNS
    ]
    .copy()
)

# =====
# FINAL TARGET
# =====

y_train_val = (
    train_val_df[
```

```

        "target_category"
    ]
    .map(
        map_target_category
    )
)

# =====
# GET SELECTED CLASSIFIER
# =====

selected_combined_pipeline = (
    combined_models_fitted[
        FINAL_MODEL_NAME
    ]
)

selected_classifier = (
    selected_combined_pipeline
    .named_steps[
        "classifier"
    ]
)

# =====
# BUILD FINAL PIPELINE
#
# Tabular
# +
# Word TF-IDF
# +
# Character TF-IDF
# +
# Selected classifier
# =====

final_interest_model = Pipeline([

    (
        "preprocessor",
        make_combined_preprocessor()
    ),

    (
        "classifier",

```

```

        clone(
            selected_classifier
        )
    )
])

# =====
# FIT FINAL MODEL
# =====

print("=" * 80)
print("FINAL MODEL REFIT")
print("=" * 80)

print(
    "\nModel:",
    FINAL_MODEL_NAME
)

print(
    "Representation:",
    FINAL_REPRESENTATION
)

print(
    "Training + validation shape:",
    X_train_val.shape
)

print(
    "\nFitting final model..."
)

final_interest_model.fit(
    X_train_val,
    y_train_val
)

print(
    "\nFinal model fitted successfully."
)

```

```

# =====
# VERIFY FINAL PIPELINE
# =====

print(
    "\nFinal pipeline:"
)

print(
    final_interest_model
)

# =====
# SAVE FINAL MODEL
# =====

FINAL_MODEL_PATH = (
    MODEL_DIR
    / "interest_model_pipeline.joblib"
)

joblib.dump(
    final_interest_model,
    FINAL_MODEL_PATH
)

print(
    "\nSaved final model:"
)

print(
    FINAL_MODEL_PATH
)

```

```

=====
FINAL MODEL REFIT
=====

```

```

Model: SVM + Text
Representation: Tabular + Review Text
Training + validation shape: (2380, 45)

```

```

Fitting final model...

```

```

Final model fitted successfully.

```

```

Final pipeline:
Pipeline(steps=[('preprocessor',
                  ColumnTransformer(transformers=[('tabular',
ColumnTransformer(transformers=[('num',
    Pipeline(steps=[('imputer',
                     SimpleImputer(strategy='median')),
                     ('scaler',
                      StandardScaler())])),
    ['prior_order_count',
     'total_spending',
     'average_order_value',
     'max_order_value',
     'recency_days',
     'history_duration_days',
     'average_days_between_orders...',
     sublinear_tf=True))])),
                  ('review_text'],
                  ('char_text',
                   Pipeline(steps=[('flatten',
FunctionTransformer(func=<function text_to_1d at 0x000002AF711F68D0>)),
                  ('tfidf',
TfidfVectorizer(analyzer='char_wb',
                 max_df=0.98,
                 max_features=3000,
                 ngram_range=(3,
                              5),
                 sublinear_tf=True))])),
                  ('review_text'])))],
              ('classifier',
               SVC(C=2.0, class_weight='balanced', probability=True,
                  random_state=42)))]

```

Saved final model:

C:\Users\ADMIN\Downloads\Intelligent System\assignment
2\A2_submit\e-commerce\model\interest_model_pipeline.joblib

```

[87]: # =====
# CELL 22B - Recommendation catalog
#
# Built from historical purchases available to Train + Validation.
# =====

train_val_customer_ids = set(
    train_val_df[
        "customer_unique_id"
    ]
)

```

```

# -----
# Historical orders belonging to Train + Validation customers
# -----

catalog_history_orders = (
    history_orders[
        history_orders[
            "customer_unique_id"
        ].isin(
            train_val_customer_ids
        )
    ]
    .copy()
)

# -----
# Attach products
# -----

catalog_items = (
    order_items_clean

    .merge(
        products_clean[
            [
                "product_id",
                "product_category_name"
            ]
        ],
        on="product_id",
        how="left"
    )

    .merge(
        catalog_history_orders[
            [
                "order_id",
                "customer_unique_id"
            ]
        ],
        on="order_id",
        how="inner"
    )
)

# English category name

```

```

catalog_items["category_en"] = (
    catalog_items[
        "product_category_name"
    ]
    .map(
        category_translation
    )
    .fillna(
        catalog_items[
            "product_category_name"
        ]
    )
    .fillna(
        "unknown"
    )
)

```

```

# -----
# Product-level sales statistics
# -----

```

```

product_sales = (
    catalog_items
    .groupby(
        [
            "product_id",
            "product_category_name",
            "category_en"
        ]
    )
    .agg(
        sold_units=(
            "order_id",
            "count"
        ),

        unique_orders=(
            "order_id",
            "nunique"
        ),

        average_price=(
            "price",
            "mean"
        )
    )
    .reset_index()

```

```

)

# -----
# Product review statistics
# -----

product_reviews = (
    reviews_clean[
        [
            "order_id",
            "review_score"
        ]
    ]

    .merge(
        order_items_clean[
            [
                "order_id",
                "product_id"
            ]
        ],
        on="order_id",
        how="inner"
    )

    .merge(
        catalog_history_orders[
            [
                "order_id"
            ]
        ],
        on="order_id",
        how="inner"
    )

    .groupby(
        "product_id"
    )
    .agg(
        average_review_score=(
            "review_score",
            "mean"
        ),

        review_count=(
            "review_score",
            "count"
        )
    )

```

```

    )
    )
    .reset_index()
)

# -----
# Combine sales + reviews
# -----

recommendation_catalog = (
    product_sales
    .merge(
        product_reviews,
        on="product_id",
        how="left"
    )
)

recommendation_catalog[
    "average_review_score"
] = (
    recommendation_catalog[
        "average_review_score"
    ]
    .fillna(
        recommendation_catalog[
            "average_review_score"
        ].median()
    )
)

recommendation_catalog[
    "review_count"
] = (
    recommendation_catalog[
        "review_count"
    ]
    .fillna(0)
)

# -----
# Helper
# -----

def group_minmax(series):

    minimum = series.min()

```

```

maximum = series.max()

if (
    pd.isna(minimum)
    or
    pd.isna(maximum)
):

    return pd.Series(
        np.zeros(
            len(series)
        ),
        index=series.index
    )

if maximum == minimum:

    return pd.Series(
        np.ones(
            len(series)
        ),
        index=series.index
    )

return (
    (series - minimum)
    /
    (maximum - minimum)
)

# -----
# Score each product inside its category
# -----

recommendation_catalog[
    "sales_score"
] = (
    recommendation_catalog
    .groupby(
        "category_en"
    )[
        "sold_units"
    ]
    .transform(
        group_minmax
    )
)

```

```

recommendation_catalog[
    "order_score"
] = (
    recommendation_catalog
    .groupby(
        "category_en"
    )[
        "unique_orders"
    ]
    .transform(
        group_minmax
    )
)

recommendation_catalog[
    "rating_score"
] = (
    recommendation_catalog
    .groupby(
        "category_en"
    )[
        "average_review_score"
    ]
    .transform(
        group_minmax
    )
)

# -----
# Recommendation score
# -----

recommendation_catalog[
    "recommendation_score"
] = (
    0.55
    *
    recommendation_catalog[
        "sales_score"
    ]

    +

    0.30
    *
    recommendation_catalog[

```

```

        "rating_score"
    ]

    +

    0.15
    *
    recommendation_catalog[
        "order_score"
    ]
)

recommendation_catalog = (
    recommendation_catalog
    .sort_values(
        [
            "category_en",
            "recommendation_score",
            "sold_units"
        ],
        ascending=[
            True,
            False,
            False
        ]
    )
    .reset_index(drop=True)
)

print(
    "Recommendation catalog shape:",
    recommendation_catalog.shape
)

display(
    recommendation_catalog.head(15)
)

```

Recommendation catalog shape: (2182, 12)

	product_id	product_category_name \
0	3e3f442db862cb6fe99389a41b7acb84	agro_industria_e_comercio
1	a0fe1efb855f3e786f0650268cd77f44	agro_industria_e_comercio
2	1490307eb0bc9c499936ad716b79ca45	climatizacao
3	20082386a501cbd8c41765656a115a2f	climatizacao
4	5c821d665e58b08092e830f81c12cc37	climatizacao
5	7ad1818fa5f8a28faf8844aec3b3f540	climatizacao
6	7deb3aeb864ba8e77c6db5baac8dae5f	climatizacao

7	bcbeea161819eed81ab43bc9751e012c	climatizacao
8	d9c379ab5377be590da568f41cd3c49d	climatizacao
9	a3f986f6c4400aa6545d224be435f812	climatizacao
10	f2a1b32f85cad59ff2a8444154ac25f0	climatizacao
11	58f036d3b8f12d9378518f902fb4e226	artes
12	d3d10495159bcc41173e3d26a273529a	artes
13	ec90895370885d025efeb8e72e7fa699	artes_e_artesanato
14	a59fb60fddcc72a9878b7ed5cb69d8e4	audio

	category_en	sold_units	unique_orders	average_price	\
0	agro_industry_and_commerce	1	1	869.97	
1	agro_industry_and_commerce	1	1	21.99	
2	air_conditioning	1	1	169.90	
3	air_conditioning	1	1	249.90	
4	air_conditioning	1	1	89.00	
5	air_conditioning	1	1	229.99	
6	air_conditioning	1	1	35.90	
7	air_conditioning	1	1	24.90	
8	air_conditioning	1	1	199.90	
9	air_conditioning	1	1	125.90	
10	air_conditioning	1	1	190.00	
11	art	1	1	49.90	
12	art	1	1	48.65	
13	arts_and_craftmanship	1	1	129.90	
14	audio	2	2	17.40	

	average_review_score	review_count	sales_score	order_score	\
0	5.0	1.0	1.0	1.0	
1	4.0	1.0	1.0	1.0	
2	5.0	1.0	1.0	1.0	
3	5.0	1.0	1.0	1.0	
4	5.0	0.0	1.0	1.0	
5	5.0	1.0	1.0	1.0	
6	5.0	1.0	1.0	1.0	
7	5.0	1.0	1.0	1.0	
8	4.0	1.0	1.0	1.0	
9	1.0	1.0	1.0	1.0	
10	1.0	1.0	1.0	1.0	
11	5.0	1.0	1.0	1.0	
12	5.0	1.0	1.0	1.0	
13	5.0	1.0	1.0	1.0	
14	1.5	2.0	1.0	1.0	

	rating_score	recommendation_score
0	1.00	1.000
1	0.00	0.700
2	1.00	1.000
3	1.00	1.000

4	1.00	1.000
5	1.00	1.000
6	1.00	1.000
7	1.00	1.000
8	0.75	0.925
9	0.00	0.700
10	0.00	0.700
11	1.00	1.000
12	1.00	1.000
13	1.00	1.000
14	0.00	0.700

```
[88]: # =====
# CELL 22C - Persist all model artifacts
# =====

# -----
# Model directory
# -----

MODEL_DIR = (
    DATA_DIR.parent
    / "A2_submit"
    / "e-commerce"
    / "model"
)

MODEL_DIR.mkdir(
    parents=True,
    exist_ok=True
)

print(
    "Model directory:"
)

print(
    MODEL_DIR
)

# =====
# 1. FINAL INTEREST MODEL
# =====

interest_model_path = (
```

```

MODEL_DIR
/ "interest_model_pipeline.joblib"
)

joblib.dump(
    final_interest_model,
    interest_model_path
)

# =====
# 2. INTEREST PREPROCESSOR
# =====

interest_preprocessor = (
    final_interest_model
    .named_steps[
        "preprocessor"
    ]
)

interest_preprocessor_path = (
    MODEL_DIR
    / "interest_preprocessor.joblib"
)

joblib.dump(
    interest_preprocessor,
    interest_preprocessor_path
)

# =====
# 3. BEHAVIOR MODEL
# =====

behavior_model_path = (
    MODEL_DIR
    / "behavior_model_pipeline.joblib"
)

joblib.dump(
    behavior_model,
    behavior_model_path
)

```

```

# =====
# 4. BEHAVIOR PREPROCESSOR
# =====

behavior_preprocessor = (
    behavior_model
    .named_steps[
        "preprocessor"
    ]
)

behavior_preprocessor_path = (
    MODEL_DIR
    / "behavior_preprocessor.joblib"
)

joblib.dump(
    behavior_preprocessor,
    behavior_preprocessor_path
)

# =====
# 5. RECOMMENDATION CATALOG
# =====

recommendation_catalog_path = (
    MODEL_DIR
    / "recommendation_catalog.joblib"
)

joblib.dump(
    recommendation_catalog,
    recommendation_catalog_path
)

# =====
# 6. METADATA
# =====

metadata = {
    "dataset":
        "Brazilian E-Commerce Public Dataset by Olist",
    "kaggle_url":

```

```

    "https://www.kaggle.com/datasets/"
    "olistbr/brazilian-ecommerce",

    "random_seed":
        RANDOM_SEED,

    "split":
        {
            "train": 0.70,
            "validation": 0.15,
            "test": 0.15
        },

    "interest_task":
        "Predict customer product-category interest",

    "behavior_task":
        "Predict target-order value segment",

    "final_model":
        FINAL_MODEL_NAME,

    "final_representation":
        FINAL_REPRESENTATION,

    "model_selection_source":
        FINAL_SOURCE,

    "validation_macro_f1":
        float(
            FINAL_VALIDATION_F1
        ),

    "numerical_features":
        NUMERICAL_FEATURES,

    "categorical_features":
        CATEGORICAL_FEATURES,

    "text_features":
        [
            "review_text"
        ],

    "text_representation":
        {
            "word_tfidf":

```

```

        True,

        "character_tfidf":
            True,

        "word_max_features":
            3000,

        "character_max_features":
            3000,

        "word_ngram_range":
            [1, 2],

        "character_ngram_range":
            [3, 5]
    },

    "interest_target_categories":
        sorted(
            y_train_val.unique()
        ),

    "behavior_classes":
        [
            "Low",
            "Medium",
            "High"
        ],

    "behavior_low_threshold":
        float(
            low_threshold
        ),

    "behavior_high_threshold":
        float(
            high_threshold
        ),

    "recommendation_score":
        {
            "sales": 0.55,
            "rating": 0.30,
            "order": 0.15
        },

```

```

        "api_endpoint":
            "/predict"
    }

# =====
# SAVE METADATA
# =====

metadata_path = (
    MODEL_DIR
    / "metadata.json"
)

with open(
    metadata_path,
    "w",
    encoding="utf-8"
) as file:

    json.dump(
        metadata,
        file,
        indent=2,
        ensure_ascii=False
    )

# =====
# VERIFY
# =====

print(
    "\nPersisted artifacts:"
)

for path in sorted(
    MODEL_DIR.iterdir()
):

    print(
        f" - {path.name}"
    )

```

Model directory:

C:\Users\ADMIN\Downloads\Intelligent System\assignment
2\A2_submit\e-commerce\model

Persisted artifacts:

- behavior_model_pipeline.joblib
- behavior_preprocessor.joblib
- interest_model_pipeline.joblib
- interest_preprocessor.joblib
- metadata.json
- recommendation_catalog.joblib

```
[89]: # =====
# CELL 22E - VERIFY DEPLOYMENT ARTIFACTS
# =====

print("=" * 80)
print("VERIFYING DEPLOYMENT ARTIFACTS")
print("=" * 80)

# =====
# LOAD FINAL MODEL
# =====

check_interest_model = joblib.load(
    MODEL_DIR
    / "interest_model_pipeline.joblib"
)

# =====
# LOAD CATALOG
# =====

check_catalog = joblib.load(
    MODEL_DIR
    / "recommendation_catalog.joblib"
)

# =====
# LOAD METADATA
# =====

with open(
    MODEL_DIR
    / "metadata.json",
    "r",
    encoding="utf-8"
) as file:
```

```

    check_metadata = json.load(
        file
    )

# =====
# DISPLAY
# =====

print(
    "\nFinal model:"
)

print(
    check_metadata[
        "final_model"
    ]
)

print(
    "\nFinal representation:"
)

print(
    check_metadata[
        "final_representation"
    ]
)

print(
    "\nSelection source:"
)

print(
    check_metadata[
        "model_selection_source"
    ]
)

print(
    "\nValidation Macro F1:"
)

```

```

print(
    check_metadata[
        "validation_macro_f1"
    ]
)

print(
    "\nPipeline:"
)

print(
    check_interest_model
)

print(
    "\nRecommendation catalog shape:"
)

print(
    check_catalog.shape
)

```

VERIFYING DEPLOYMENT ARTIFACTS

Final model:

SVM + Text

Final representation:

Tabular + Review Text

Selection source:

Experiment 2 - Tabular + Review Text

Validation Macro F1:

0.45438607191499364

Pipeline:

```

Pipeline(steps=[('preprocessor',
                  ColumnTransformer(transformers=[('tabular',
ColumnTransformer(transformers=[('num',
    Pipeline(steps=[('imputer',
                      SimpleImputer(strategy='median')),
                      ('scaler',
                        StandardScaler())])),

```

```

['prior_order_count',
 'total_spending',
 'average_order_value',
 'max_order_value',
 'recency_days',
 'history_duration_days',
 'average_days_between_orders...',
 sublinear_tf=True))]],

                                ['review_text']],
                                ('char_text',
                                 Pipeline(steps=[('flatten',
FunctionTransformer(func=<function text_to_1d at 0x000002AF711F68D0>)),
                                ('tfidf',
TfidfVectorizer(analyzer='char_wb',
                 max_df=0.98,
                 max_features=3000,
                 ngram_range=(3,
                               5),
                 sublinear_tf=True))])),

                                ['review_text']]])),
                                ('classifier',
                                 SVC(C=2.0, class_weight='balanced', probability=True,
                                     random_state=42)))]))

```

Recommendation catalog shape:
(2182, 12)

24 23. Inference Test

The final section verifies the complete inference pipeline.

The deployed inference process is:

Customer Input -> Validation -> Same Preprocessing -> Interest Model -> Predicted Interest ->
Behavior Model -> Predicted Behavior -> Recommendation Engine -> Top 5 Products

24.1 Training vs Inference

Training:

Training Data -> Fit preprocessing -> Fit ML model

Inference:

New Customer Input -> Load saved preprocessing -> Transform input -> Load saved model ->
Prediction

The system never fits a new scaler, encoder, imputer, or TF-IDF vocabulary using the user input.

24.2 Final Output

The intelligent system returns:

- Predicted behavior
- Behavior confidence
- Predicted interest
- Interest confidence
- Top 5 recommended products

This output will become the response contract for the REST API.

The same API response can then be consumed by:

- Web application
- Android/Kotlin mobile application

```
[90]: # =====  
# CELL 23A - Top 5 recommendation function  
# =====  
  
def recommend_top_products(  
    predicted_interest,  
    catalog,  
    top_n=5  
):  
  
    normalized_interest = (  
        str(  
            predicted_interest  
        )  
        .strip()  
        .lower()  
    )  
  
    filtered = catalog[  
        catalog[  
            "category_en"  
        ]  
        .fillna("")  
        .str  
        .strip()  
        .str  
        .lower()  
        .eq(  
            normalized_interest  
        )  
    ].copy()  
  
    return (
```

```

        filtered
        .sort_values(
            [
                "recommendation_score",
                "sold_units"
            ],
            ascending=[
                False,
                False
            ]
        )
        .head(
            top_n
        )[
            [
                "product_id",
                "category_en",
                "sold_units",
                "average_price",
                "average_review_score",
                "recommendation_score"
            ]
        ]
    )

```

```

[91]: # =====
# CELL 23B - FINAL TEST EVALUATION
# =====

final_evaluation_X = (
    X_test
    .copy()
)

# =====
# PREDICTION
# =====

final_test_prediction = (
    final_interest_model
    .predict(
        final_evaluation_X
    )
)

```

```

# =====
# PROBABILITY
# =====

final_test_probability = (
    final_interest_model
    .predict_proba(
        final_evaluation_X
    )
)

# =====
# METRICS
# =====

final_test_metrics = {

    "Accuracy":
        accuracy_score(
            y_test,
            final_test_prediction
        ),

    "Precision Macro":
        precision_score(
            y_test,
            final_test_prediction,
            average="macro",
            zero_division=0
        ),

    "Recall Macro":
        recall_score(
            y_test,
            final_test_prediction,
            average="macro",
            zero_division=0
        ),

    "F1 Macro":
        f1_score(
            y_test,
            final_test_prediction,
            average="macro",
            zero_division=0
        ),
}

```

```

    "F1 Weighted":
        f1_score(
            y_test,
            final_test_prediction,
            average="weighted",
            zero_division=0
        ),

    "ROC-AUC OVR Macro":
        roc_auc_score(
            y_test,
            final_test_probability,
            multi_class="ovr",
            average="macro"
        )
}

# =====
# DISPLAY
# =====

display(
    pd.DataFrame(
        [final_test_metrics
        ]
    )
)

print(
    "\nFinal model:",
    FINAL_MODEL_NAME
)

print(
    "Final representation:",
    FINAL_REPRESENTATION
)

```

	Accuracy	Precision Macro	Recall Macro	F1 Macro	F1 Weighted \
0	0.463183	0.414376	0.385122	0.394294	0.45841

	ROC-AUC OVR Macro
0	0.701825

Final model: SVM + Text

Final representation: Tabular + Review Text

```
[92]: # =====
# CELL 23C - Final inference demo
# FINAL MODEL = TABULAR + REVIEW TEXT
# =====

example_index = 0

example_customer_id = (
    test_df
    .iloc[
        example_index
    ]
    ["customer_unique_id"]
)

example_input = (
    X_test
    .iloc[
        [example_index]
    ]
    .copy()
)

# -----
# Interest prediction
# -----

predicted_interest = (
    final_interest_model
    .predict(
        example_input
    )[0]
)

interest_probability = (
    final_interest_model
    .predict_proba(
        example_input
    )[0]
)

interest_classes = (
    final_interest_model
```

```

        .named_steps[
            "classifier"
        ]
        .classes_
    )

interest_best_index = (
    np.argmax(
        interest_probability
    )
)

interest_confidence = float(
    interest_probability[
        interest_best_index
    ]
)

# -----
# Behavior prediction
# -----

predicted_behavior = (
    behavior_model
    .predict(
        example_input
    )[0]
)

behavior_probability = (
    behavior_model
    .predict_proba(
        example_input
    )[0]
)

behavior_best_index = (
    np.argmax(
        behavior_probability
    )
)

behavior_confidence = float(
    behavior_probability[
        behavior_best_index
    ]
)

```

```

# -----
# Recommendations
# -----

top5_products = (
    recommend_top_products(
        predicted_interest,
        recommendation_catalog,
        top_n=5
    )
)

# -----
# Display
# -----

print("=" * 80)
print("FINAL INFERENCE RESULT")
print("=" * 80)

print(
    "\nCustomer ID:",
    example_customer_id
)

print(
    "\nReview Text:"
)

print(
    example_input[
        "review_text"
    ].iloc[0]
)

print(
    "\nPredicted Customer Behavior:",
    predicted_behavior
)

print(
    "Behavior Confidence:",
    round(
        behavior_confidence,
        4
    )
)

```

```

)

print(
    "\nPredicted Customer Interest:",
    predicted_interest
)

print(
    "Interest Confidence:",
    round(
        interest_confidence,
        4
    )
)

print(
    "\nTop 5 Recommended Products:"
)

display(
    top5_products
)

```

```

=====
FINAL INFERENCE RESULT
=====

```

Customer ID: 15a38a0c1bf789e06ec42d24e485e794

Review Text:

o forro interno é frágil e todo manchado de tinta o suporte de controle tem grandes falhas na pintura interna

Predicted Customer Behavior: High
Behavior Confidence: 0.436

Predicted Customer Interest: moveis_decoracao
Interest Confidence: 0.2742

Top 5 Recommended Products:

Empty DataFrame

Columns: [product_id, category_en, sold_units, average_price,
average_review_score, recommendation_score]

Index: []

```

[93]: # =====
      # CELL 23C_Test - REVIEW TEXT EFFECT TEST

```

```

# =====

# -----
# Same customer
# Same tabular features
# Only review_text changes
# -----

input_a = (
    X_test
    .iloc[
        0]
    ]
    .copy()
)

input_b = (
    X_test
    .iloc[
        0]
    ]
    .copy()
)

# -----
# Review A
# -----

input_a[
    "review_text"
] = (
    "Good product, excellent quality, "
    "fast delivery and very satisfied with the purchase."
)

# -----
# Review B
# -----

input_b[
    "review_text"
] = (
    "I hate fashion. Terrible product, "
    "bad quality and very disappointing."
)

```

```

# -----
# Transform ONLY
# -----

vector_a = (
    final_interest_model
    .named_steps[
        "preprocessor"
    ]
    .transform(
        input_a
    )
)

vector_b = (
    final_interest_model
    .named_steps[
        "preprocessor"
    ]
    .transform(
        input_b
    )
)

# -----
# Check whether representations are different
# -----

representation_difference = (
    np.max(
        np.abs(
            vector_a
            -
            vector_b
        )
    )
)

print(
    "Maximum representation difference:",
    representation_difference
)

```

```

# -----
# Prediction
# -----

probability_a = (
    final_interest_model
    .predict_proba(
        input_a
    )[0]
)

probability_b = (
    final_interest_model
    .predict_proba(
        input_b
    )[0]
)

prediction_a = (
    final_interest_model
    .predict(
        input_a
    )[0]
)

prediction_b = (
    final_interest_model
    .predict(
        input_b
    )[0]
)

classes = (
    final_interest_model
    .named_steps[
        "classifier"
    ]
    .classes_
)

# -----
# Comparison
# -----

```

```

comparison = pd.DataFrame({

    "Class":
        classes,

    "Review A Probability":
        probability_a,

    "Review B Probability":
        probability_b,

    "Difference":
        probability_b
        -
        probability_a
})

display(
    comparison
)

print(
    "\nPrediction A:",
    prediction_a
)

print(
    "Prediction B:",
    prediction_b
)

max_probability_difference = (
    np.max(
        np.abs(
            probability_a
            -
            probability_b
        )
    )
)

print(
    "\nMaximum probability difference:",
    max_probability_difference
)

```

```

# -----
# Final diagnosis
# -----

if representation_difference == 0:

    print(
        "\nERROR: The two review texts produced "
        "the same transformed representation."
    )

elif max_probability_difference == 0:

    print(
        "\nText representation changed, but "
        "the classifier probability did not."
    )

else:

    print(
        "\nSUCCESS: Review text changes the "
        "final model probability."
    )

```

Maximum representation difference: 0.21108049342756888

	Class	Review A Probability	Review B Probability \
0	beleza_saude	0.085910	0.085928
1	brinquedos	0.032240	0.032268
2	cama_mesa_banho	0.274980	0.274734
3	esporte_lazer	0.060850	0.060179
4	fashion_bolsas_e_acessorios	0.012077	0.012087
5	informatica_acessorios	0.045579	0.046101
6	moveis_decoracao	0.181883	0.181084
7	other	0.183463	0.184656
8	relogios_presentes	0.037908	0.037664
9	telefonica	0.006353	0.006461
10	utilidades_domesticas	0.078758	0.078837

	Difference
0	0.000018
1	0.000028
2	-0.000246
3	-0.000671
4	0.000010

```
5    0.000523
6   -0.000799
7    0.001193
8   -0.000244
9    0.000108
10   0.000080
```

Prediction A: moveis_decoracao

Prediction B: moveis_decoracao

Maximum probability difference: 0.001192705545763867

SUCCESS: Review text changes the final model probability.

```
[94]: # =====
# CELL 23D - Final API response contract
# =====

api_response = {

    "customer_id":
        str(
            example_customer_id
        ),

    "predicted_behavior":
        str(
            predicted_behavior
        ),

    "behavior_confidence":
        round(
            behavior_confidence,
            4
        ),

    "predicted_interest":
        str(
            predicted_interest
        ),

    "interest_confidence":
        round(
            interest_confidence,
            4
        ),
```

```

        "recommended_products":
            top5_products
            .to_dict(
                orient="records"
            )
    }

    print(
        json.dumps(
            api_response,
            indent=2,
            ensure_ascii=False
        )
    )

```

```

{
  "customer_id": "15a38a0c1bf789e06ec42d24e485e794",
  "predicted_behavior": "High",
  "behavior_confidence": 0.436,
  "predicted_interest": "moveis_decoracao",
  "interest_confidence": 0.2742,
  "recommended_products": []
}

```

```

[95]: # =====
# CELL 23E - Final data representation summary
# =====

print("=" * 80)
print("FINAL DATA REPRESENTATION SUMMARY")
print("=" * 80)

print(
    "\n1. Raw representation:"
)

print(
    "Multiple relational CSV tables"
)

print(
    "\n2. Customer-level representation:"
)

print(
    "Transactions -> Customer Profile -> Feature Vector"
)

```

```
print(
    "\n3. Raw customer dataset shape:"
)

print(
    customer_features.shape
)

print(
    "\n4. Training examples:"
)

print(
    len(train_df)
)

print(
    "\n5. Validation examples:"
)

print(
    len(val_df)
)

print(
    "\n6. Test examples:"
)

print(
    len(test_df)
)

print(
    "\n7. Numerical features:"
)

print(
    len(NUMERICAL_FEATURES)
)

print(
    "\n8. Categorical features:"
)

print(
    len(CATEGORICAL_FEATURES)
```

```

)

print(
    "\n9. Text representation:"
)

print(
    "Review comments -> cleaned text -> TF-IDF"
)

print(
    "\n10. Final representation:"
)

print(
    FINAL_REPRESENTATION
)

print(
    "\n11. Final model:"
)

print(
    FINAL_MODEL_NAME
)

print(
    "\n12. Final model input dimensionality:"
)

# Transform a small sample
if (
    FINAL_REPRESENTATION
    ==
    "Tabular + Review Text"
):

    transformed_sample = (
        final_interest_model
        .named_steps[
            "preprocessor"
        ]
        .transform(
            X_train_val.iloc[:5]
        )
    )

```

```

elif (
    FINAL_MODEL_NAME
    ==
    "Text-based Linear Classifier"
):

    transformed_sample = (
        final_interest_model
        .named_steps[
            "preprocessor"
        ]
        .transform(
            X_train_val[
                ["review_text"]
            ]
            .iloc[:5]
        )
    )

else:

    transformed_sample = (
        final_interest_model
        .named_steps[
            "preprocessor"
        ]
        .transform(
            X_train_val.iloc[:5]
        )
    )

print(
    transformed_sample.shape
)

print(
    "\n13. Data type:"
)

print(
    type(
        transformed_sample
    )
)

print(
    "\nThe final transformed representation is the numerical input supplied to"

```

```
" the machine-learning classifier."  
)
```

```
=====
```

FINAL DATA REPRESENTATION SUMMARY

```
=====
```

1. Raw representation:

Multiple relational CSV tables

2. Customer-level representation:

Transactions -> Customer Profile -> Feature Vector

3. Raw customer dataset shape:

(2801, 52)

4. Training examples:

1960

5. Validation examples:

420

6. Test examples:

421

7. Numerical features:

42

8. Categorical features:

2

9. Text representation:

Review comments -> cleaned text -> TF-IDF

10. Final representation:

Tabular + Review Text

11. Final model:

SVM + Text

12. Final model input dimensionality:

(5, 6073)

13. Data type:

<class 'scipy.sparse._csr.csr_matrix'>

The final transformed representation is the numerical input supplied to the machine-learning classifier.

[]: