

CNN_Assignment

September 22, 2026

1 Assignment 04 - Image Classification, CNN and Framework Comparison

1.1 0. Project overview

This notebook is organized as a complete AI report and implementation.

Required datasets:

1. MNIST - grayscale image classification.
2. CIFAR-10 - RGB image classification.
3. Diabetes Health Indicators - tabular classification.

Required implementations for image datasets:

- CNN from Scratch with NumPy.
- CNN with PyTorch.
- CNN with TensorFlow/Keras.

For the Diabetes dataset, this notebook follows the teaching logic of Lecture 04:

- Tabular data -> MLP is the natural neural-network baseline.
- Image data -> CNN is the natural architecture.

The goal is not to memorize APIs. The goal is to understand:

Data representation -> mathematical operations -> model structure -> training -> evaluation -> visualization -> comparison.

1.2 0.1 Table of contents

1.2.1 Part A - Project and data preparation

1. Project structure
2. Environment
3. Dataset inventory
4. Reproducibility
5. Common visualization and evaluation utilities

1.2.2 Part B - Dataset 1: MNIST

6. Load data
7. Explore data

- 8. Preprocess and split data
- 9. CNN architecture
- 10. CNN from Scratch
- 11. CNN with PyTorch
- 12. CNN with Keras
- 13. MNIST evaluation and visualization

1.2.3 Part C - Dataset 2: CIFAR-10

- 14. Load data
- 15. Explore data
- 16. Preprocess and split data
- 17. CNN from Scratch
- 18. CNN with PyTorch
- 19. CNN with Keras
- 20. CIFAR-10 evaluation and visualization

1.2.4 Part D - Dataset 3: Diabetes

- 21. Load data
- 22. Explore data
- 23. Preprocess and split data
- 24. MLP from Scratch
- 25. MLP with PyTorch
- 26. MLP with Keras
- 27. Diabetes evaluation and visualization

1.2.5 Part E - Final comparison and report

- 28. Framework comparison
- 29. Performance comparison
- 30. Answers to Lecture 04 questions
- 31. Save models and website-ready outputs
- 32. Final conclusions

2 Part A - Project and data preparation

2.1 1. Project structure

The notebook assumes the following project layout:

```
assignment 4/
|
+-- notebooks/
|   +-- CNN_Assignment.ipynb
|
+-- data/
|   +-- mnist/
|   +-- cifar10/
```

```
|   +-- diabetes/
|
+-- models/
|
+-- outputs/
|   +-- figures/
|   +-- metrics/
|   +-- predictions/
|
+-- web/
```

The notebook is located inside `notebooks/`, so the data folder is normally one level above the notebook.

```
[70]: from pathlib import Path
import os
import json
import time
import math
import random
import warnings
import platform

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import (
    accuracy_score,
    precision_recall_fscore_support,
    classification_report,
    confusion_matrix,
    ConfusionMatrixDisplay,
)
from sklearn.decomposition import PCA
from sklearn.utils.class_weight import compute_class_weight

warnings.filterwarnings("ignore")

# Find the project root instead of hard-coding a Windows path.
current = Path.cwd()
candidates = [
    current,
    current.parent,
    current.parent.parent,
]
```

```

PROJECT_ROOT = None
for candidate in candidates:
    if (candidate / "data").exists():
        PROJECT_ROOT = candidate
        break

if PROJECT_ROOT is None:
    raise FileNotFoundError(
        "Cannot find the project root. Open this notebook from the assignment4/
↳notebooks folder."
    )

DATA_DIR = PROJECT_ROOT / "data"
MODELS_DIR = PROJECT_ROOT / "models"
OUTPUTS_DIR = PROJECT_ROOT / "outputs"
FIGURES_DIR = OUTPUTS_DIR / "figures"
METRICS_DIR = OUTPUTS_DIR / "metrics"
PREDICTIONS_DIR = OUTPUTS_DIR / "predictions"

for folder in [MODELS_DIR, OUTPUTS_DIR, FIGURES_DIR, METRICS_DIR,
↳PREDICTIONS_DIR]:
    folder.mkdir(parents=True, exist_ok=True)

print("Project root:", PROJECT_ROOT)
print("Data folder:", DATA_DIR)
print("Models folder:", MODELS_DIR)
print("Outputs folder:", OUTPUTS_DIR)

```

```

Project root: C:\Users\ADMIN\Downloads\Intelligent System\assignment 4\DATA
Data folder: C:\Users\ADMIN\Downloads\Intelligent System\assignment 4\DATA\data
Models folder: C:\Users\ADMIN\Downloads\Intelligent System\assignment
4\DATA\models
Outputs folder: C:\Users\ADMIN\Downloads\Intelligent System\assignment
4\DATA\outputs

```

2.2 2. Environment

This project uses three implementation levels:

- NumPy for Scratch implementation.
- PyTorch for an explicit modular deep-learning implementation.
- TensorFlow/Keras for a higher-level implementation.

The same learning idea should remain visible even when the programming abstraction changes.

```

[71]: import sys
import numpy as np
import sklearn

```

```

print("Python:", sys.version.split()[0])
print("Platform:", platform.platform())
print("NumPy:", np.__version__)
print("Pandas:", pd.__version__)
print("Scikit-learn:", sklearn.__version__)

try:
    import torch
    print("PyTorch:", torch.__version__)
except Exception as exc:
    print("PyTorch import failed:", exc)

try:
    import tensorflow as tf
    print("TensorFlow:", tf.__version__)
except Exception as exc:
    print("TensorFlow import failed:", exc)

```

```

Python: 3.12.13
Platform: Windows-11-10.0.26200-SP0
NumPy: 2.5.3
Pandas: 3.0.6
Scikit-learn: 1.9.1
PyTorch: 2.14.0+cpu
TensorFlow: 2.21.0

```

2.3 3. Reproducibility and experiment configuration

Reproducibility matters because a fair comparison should use the same dataset split and comparable training settings.

The default configuration uses a deterministic subset for training. This keeps the Scratch implementation practical on a normal CPU while preserving the same selected samples for Scratch, PyTorch and Keras.

Set the limits to `None` later when you want to run on the full training set.

```

[72]: SEED = 42

# Common experiment settings
EPOCHS = 5
BATCH_SIZE = 128
LEARNING_RATE = 1e-3
VALIDATION_SIZE = 0.10

# CPU-friendly default for Scratch.
# All three frameworks use the same selected training/validation indices.
COMPARISON_TRAIN_LIMIT = None

```

```

COMPARISON_VAL_LIMIT = None

# Set to None for full-data training.
USE_COMPARISON_LIMITS = False

random.seed(SEED)
np.random.seed(SEED)

def set_global_seed(seed=SEED):
    random.seed(seed)
    np.random.seed(seed)
    try:
        import torch
        torch.manual_seed(seed)
        if torch.cuda.is_available():
            torch.cuda.manual_seed_all(seed)
    except Exception:
        pass

    try:
        import tensorflow as tf
        tf.random.set_seed(seed)
    except Exception:
        pass

set_global_seed()

print("Seed:", SEED)
print("Epochs:", EPOCHS)
print("Batch size:", BATCH_SIZE)
print("Learning rate:", LEARNING_RATE)

```

```

Seed: 42
Epochs: 5
Batch size: 128
Learning rate: 0.001

```

2.4 4. Dataset inventory

Before building a model, inspect the representation of the data.

For images, we need to know:

- Number of samples.
- Height and width.
- Number of channels.
- Number of classes.
- Class distribution.

For Diabetes, we need to know:

- Number of rows.
- Number of input features.
- Target column.
- Missing values.
- Duplicate rows.
- Class distribution.
- File size.

This follows the first principle of the lecture: the model should match the structure of the data.

```
[73]: def find_first_file(folder, patterns):
    matches = []
    for pattern in patterns:
        matches.extend(folder.rglob(pattern))
    matches = [p for p in matches if p.is_file()]
    return sorted(set(matches))

mnist_npz = DATA_DIR / "mnist" / "mnist.npz"
cifar_npz = DATA_DIR / "cifar10" / "cifar10.npz"

diabetes_candidates = find_first_file(
    DATA_DIR / "diabetes",
    ["*.csv"]
)

print("MNIST file:", mnist_npz, "->", mnist_npz.exists())
print("CIFAR-10 file:", cifar_npz, "->", cifar_npz.exists())
print("Diabetes CSV candidates:", diabetes_candidates)
```

```
MNIST file: C:\Users\ADMIN\Downloads\Intelligent System\assignment
4\DATA\data\mnist\mnist.npz -> True
CIFAR-10 file: C:\Users\ADMIN\Downloads\Intelligent System\assignment
4\DATA\data\cifar10\cifar10.npz -> True
Diabetes CSV candidates: [WindowsPath('C:/Users/ADMIN/Downloads/Intelligent
System/assignment
4/DATA/data/diabetes/diabetes_012_health_indicators_BRFSS2015.csv'),
WindowsPath('C:/Users/ADMIN/Downloads/Intelligent System/assignment 4/DATA/data/
diabetes/diabetes_binary_5050split_health_indicators_BRFSS2015.csv'),
WindowsPath('C:/Users/ADMIN/Downloads/Intelligent System/assignment
4/DATA/data/diabetes/diabetes_binary_health_indicators_BRFSS2015.csv')]
```

2.5 4.1 Automatic fallback for MNIST and CIFAR-10

If the .npz files are missing, the next cell can download them through TensorFlow/Keras and store them inside the project.

This is only data loading. It does not mean the CNN itself is implemented with Keras.

```
[74]: # Download fallback: run if the local .npz files are missing.

if not mnist_npz.exists():
    from tensorflow.keras.datasets import mnist
    mnist_dir = mnist_npz.parent
    mnist_dir.mkdir(parents=True, exist_ok=True)

    (x_train_m, y_train_m), (x_test_m, y_test_m) = mnist.load_data()
    np.savez_compressed(
        mnist_npz,
        x_train=x_train_m,
        y_train=y_train_m,
        x_test=x_test_m,
        y_test=y_test_m,
    )
    print("MNIST downloaded and saved to:", mnist_npz)

if not cifar_npz.exists():
    from tensorflow.keras.datasets import cifar10
    cifar_dir = cifar_npz.parent
    cifar_dir.mkdir(parents=True, exist_ok=True)

    (x_train_c, y_train_c), (x_test_c, y_test_c) = cifar10.load_data()
    np.savez_compressed(
        cifar_npz,
        x_train=x_train_c,
        y_train=y_train_c.reshape(-1),
        x_test=x_test_c,
        y_test=y_test_c.reshape(-1),
    )
    print("CIFAR-10 downloaded and saved to:", cifar_npz)
```

```
[75]: if not mnist_npz.exists():
        raise FileNotFoundError(f"Missing MNIST file: {mnist_npz}")

    if not cifar_npz.exists():
        raise FileNotFoundError(f"Missing CIFAR-10 file: {cifar_npz}")

    if not diabetes_candidates:
        raise FileNotFoundError(
            f"No CSV file found under {DATA_DIR / 'diabetes'}"
        )

    DIABETES_PATH = diabetes_candidates[0]

    print("Diabetes file:", DIABETES_PATH)
    print(
```

```

        f"Diabetes file size: {DIABETES_PATH.stat().st_size / (1024**2):.2f} MB"
    )

if DIABETES_PATH.stat().st_size <= 15 * 1024 * 1024:
    print("WARNING: The selected Diabetes CSV is not larger than 15 MB.")
else:
    print("Diabetes file satisfies the >15 MB requirement.")

```

Diabetes file: C:\Users\ADMIN\Downloads\Intelligent System\assignment
4\DATA\data\diabetes\diabetes_012_health_indicators_BRFSS2015.csv
Diabetes file size: 21.68 MB
Diabetes file satisfies the >15 MB requirement.

2.6 5. Common utilities

The next cells define reusable functions for:

- Dataset summaries.
- Image grids.
- Class distributions.
- Training curves.
- Classification metrics.
- Confusion matrices.
- Correct and incorrect predictions.
- Model parameter counts.
- Saving metrics.

Keeping these functions centralized makes the notebook easier to read and reduces duplicated code.

```

[76]: IMAGE_CLASS_NAMES = {
    "mnist": [str(i) for i in range(10)],
    "cifar10": [
        "airplane", "automobile", "bird", "cat", "deer",
        "dog", "frog", "horse", "ship", "truck"
    ],
}

def plot_image_grid(images, labels, class_names=None, n=12, title="Image_
↪samples"):
    n = min(n, len(images))
    cols = 4
    rows = math.ceil(n / cols)

    plt.figure(figsize=(12, 3 * rows))

    for i in range(n):
        ax = plt.subplot(rows, cols, i + 1)

        image = images[i]

```

```

        if image.ndim == 3 and image.shape[0] in (1, 3):
            image = np.transpose(image, (1, 2, 0))

        if image.ndim == 3 and image.shape[-1] == 1:
            ax.imshow(image.squeeze(), cmap="gray")
        elif image.ndim == 2:
            ax.imshow(image, cmap="gray")
        else:
            ax.imshow(image)

        label = int(labels[i])
        text_label = class_names[label] if class_names else str(label)
        ax.set_title(text_label)
        ax.axis("off")

    plt.suptitle(title)
    plt.tight_layout()
    plt.show()

def plot_class_distribution(labels, class_names=None, title="Class_
↪distribution"):
    values, counts = np.unique(labels, return_counts=True)
    names = [
        class_names[int(v)] if class_names else str(v)
        for v in values
    ]

    plt.figure(figsize=(10, 5))
    plt.bar(names, counts)
    plt.title(title)
    plt.xlabel("Class")
    plt.ylabel("Number of samples")
    plt.xticks(rotation=30, ha="right")
    plt.tight_layout()
    plt.show()

def plot_training_history(history, title="Training history"):
    epochs = range(1, len(history["loss"]) + 1)

    plt.figure(figsize=(9, 5))
    plt.plot(epochs, history["loss"], marker="o", label="Train loss")
    if "val_loss" in history:
        plt.plot(epochs, history["val_loss"], marker="o", label="Validation_
↪loss")
    plt.xlabel("Epoch")

```

```

plt.ylabel("Loss")
plt.title(title + " - Loss")
plt.legend()
plt.tight_layout()
plt.show()

if "accuracy" in history:
    plt.figure(figsize=(9, 5))
    plt.plot(epochs, history["accuracy"], marker="o", label="Train_
↪accuracy")
    if "val_accuracy" in history:
        plt.plot(epochs, history["val_accuracy"], marker="o",
↪label="Validation accuracy")
        plt.xlabel("Epoch")
        plt.ylabel("Accuracy")
        plt.title(title + " - Accuracy")
        plt.legend()
        plt.tight_layout()
        plt.show()

def compute_classification_metrics(y_true, y_pred, y_prob=None):
    precision_macro, recall_macro, f1_macro, _ =
↪precision_recall_fscore_support(
        y_true, y_pred, average="macro", zero_division=0
    )

    precision_weighted, recall_weighted, f1_weighted, _ =
↪precision_recall_fscore_support(
        y_true, y_pred, average="weighted", zero_division=0
    )

    result = {
        "accuracy": float(accuracy_score(y_true, y_pred)),
        "precision_macro": float(precision_macro),
        "recall_macro": float(recall_macro),
        "f1_macro": float(f1_macro),
        "precision_weighted": float(precision_weighted),
        "recall_weighted": float(recall_weighted),
        "f1_weighted": float(f1_weighted),
    }

    if y_prob is not None:
        result["mean_confidence"] = float(np.max(y_prob, axis=1).mean())

    return result

```

```

def show_confusion_matrix(y_true, y_pred, class_names, title):
    cm = confusion_matrix(y_true, y_pred)

    fig, ax = plt.subplots(figsize=(8, 7))
    disp = ConfusionMatrixDisplay(confusion_matrix=cm,
    ↪display_labels=class_names)
    disp.plot(ax=ax, values_format="d", colorbar=False)
    ax.set_title(title)
    plt.tight_layout()
    plt.show()

def show_classification_report(y_true, y_pred, class_names, title):
    report = classification_report(
        y_true,
        y_pred,
        target_names=class_names,
        output_dict=True,
        zero_division=0,
    )

    report_df = pd.DataFrame(report).T.iloc[:3]

    plt.figure(figsize=(10, 6))
    report_df[["precision", "recall", "f1-score"]].plot(kind="bar")
    plt.title(title)
    plt.xlabel("Class")
    plt.ylabel("Score")
    plt.ylim(0, 1.05)
    plt.xticks(rotation=35, ha="right")
    plt.tight_layout()
    plt.show()

    return report_df

def show_confidence_distribution(y_prob, correct, title):
    confidence = np.max(y_prob, axis=1)

    plt.figure(figsize=(10, 5))
    plt.hist(confidence[correct], bins=20, alpha=0.7, label="Correct")
    plt.hist(confidence[~correct], bins=20, alpha=0.7, label="Incorrect")
    plt.xlabel("Predicted confidence")
    plt.ylabel("Number of samples")
    plt.title(title)
    plt.legend()

```

```

plt.tight_layout()
plt.show()

def show_prediction_examples(
    images,
    y_true,
    y_pred,
    y_prob,
    class_names,
    title,
    only_wrong=True,
    n=12,
):
    correct = y_true == y_pred
    indices = np.where(~correct if only_wrong else correct)[0]

    if len(indices) == 0:
        print("No matching examples found.")
        return

    indices = indices[:n]
    cols = 4
    rows = math.ceil(len(indices) / cols)

    plt.figure(figsize=(12, 3 * rows))

    for plot_idx, sample_idx in enumerate(indices):
        ax = plt.subplot(rows, cols, plot_idx + 1)

        image = images[sample_idx]
        if image.ndim == 3 and image.shape[0] in (1, 3):
            image = np.transpose(image, (1, 2, 0))

        if image.ndim == 2:
            ax.imshow(image, cmap="gray")
        elif image.ndim == 3 and image.shape[-1] == 1:
            ax.imshow(image.squeeze(), cmap="gray")
        else:
            ax.imshow(image)

        true_name = class_names[int(y_true[sample_idx])]
        pred_name = class_names[int(y_pred[sample_idx])]
        conf = float(np.max(y_prob[sample_idx]))

        ax.set_title(
            f"True: {true_name}\nPred: {pred_name} | conf={conf:.3f}"

```

```

    )
    ax.axis("off")

plt.suptitle(title)
plt.tight_layout()
plt.show()

def save_metrics(metrics, path):
    path.parent.mkdir(parents=True, exist_ok=True)
    with open(path, "w", encoding="utf-8") as f:
        json.dump(metrics, f, indent=2)

```

```

[77]: def count_numpy_parameters(params):
        return int(sum(np.size(value) for value in params.values()))

def make_stratified_indices(y, val_size=0.10, seed=42, train_limit=None,
↪ val_limit=None):
    all_indices = np.arange(len(y))

    train_idx, val_idx = train_test_split(
        all_indices,
        test_size=val_size,
        stratify=y,
        random_state=seed,
    )

    if train_limit is not None and train_limit < len(train_idx):
        train_idx, _ = train_test_split(
            train_idx,
            train_size=train_limit,
            stratify=y[train_idx],
            random_state=seed,
        )

    if val_limit is not None and val_limit < len(val_idx):
        val_idx, _ = train_test_split(
            val_idx,
            train_size=val_limit,
            stratify=y[val_idx],
            random_state=seed,
        )

    return np.sort(train_idx), np.sort(val_idx)

```

```
def report_split(name, y_train, y_val, y_test):
    print(name)
    print("  train:", len(y_train))
    print("  val  :", len(y_val))
    print("  test :", len(y_test))
    print("  train class counts:", np.bincount(y_train))
    print("  val class counts  :", np.bincount(y_val))
    print("  test class counts :", np.bincount(y_test))
```

3 Part B - Dataset 1: MNIST

3.1 6. Load MNIST

MNIST is a natural CNN example because each sample is a 2D image.

Canonical representation:

N x C x H x W
N x 1 x 28 x 28

Data source: https://www.tensorflow.org/api_docs/python/tf/keras/datasets/mnist/load_data

The important idea is that neighboring pixels have meaningful spatial relationships.

```
[78]: mnist_data = np.load(mnist_npz)

mnist_x_train_raw = mnist_data["x_train"]
mnist_y_train = mnist_data["y_train"].astype(np.int64)

mnist_x_test_raw = mnist_data["x_test"]
mnist_y_test = mnist_data["y_test"].astype(np.int64)

# Convert to a canonical N x C x H x W representation.
mnist_x_train = mnist_x_train_raw.astype(np.float32) / 255.0
mnist_x_test = mnist_x_test_raw.astype(np.float32) / 255.0

mnist_x_train = mnist_x_train[:, None, :, :]
mnist_x_test = mnist_x_test[:, None, :, :]

print("Train:", mnist_x_train.shape, mnist_y_train.shape)
print("Test :", mnist_x_test.shape, mnist_y_test.shape)
print("Value range:", mnist_x_train.min(), "to", mnist_x_train.max())
```

```
Train: (60000, 1, 28, 28) (60000,)
Test : (10000, 1, 28, 28) (10000,)
Value range: 0.0 to 1.0
```


3.2 7. Explore MNIST

Visualization is not decoration. It helps verify whether:

- Images are loaded correctly.
- Labels match the images.
- Classes are balanced.
- Pixel values have a sensible range.
- The input representation is what the CNN expects.

```
[79]: plot_image_grid(
    mnist_x_train[:100],
    mnist_y_train[:100],
    class_names=IMAGE_CLASS_NAMES["mnist"],
    n=12,
    title="MNIST sample images",
)

plot_class_distribution(
    mnist_y_train,
    class_names=IMAGE_CLASS_NAMES["mnist"],
    title="MNIST training class distribution",
)
```







```
[80]: plt.figure(figsize=(10, 5))
plt.hist(mnist_x_train.reshape(-1), bins=40)
plt.title("MNIST pixel intensity distribution after normalization")
plt.xlabel("Pixel intensity")
plt.ylabel("Frequency")
plt.tight_layout()
plt.show()

mean_mnist = mnist_x_train.mean(axis=0).squeeze()

plt.figure(figsize=(5, 5))
plt.imshow(mean_mnist, cmap="gray")
plt.title("MNIST mean image")
plt.axis("off")
plt.tight_layout()
plt.show()
```



3.3 7.1 MNIST split

The three implementations must use the same training and validation samples.

This makes the comparison meaningful because the models are not being tested on different data.

```
[81]: mnist_train_idx, mnist_val_idx = make_stratified_indices(
    mnist_y_train,
    val_size=VALIDATION_SIZE,
    seed=SEED,
    train_limit=COMPARISON_TRAIN_LIMIT if USE_COMPARISON_LIMITS else None,
    val_limit=COMPARISON_VAL_LIMIT if USE_COMPARISON_LIMITS else None,
)

mnist_X_tr = mnist_x_train[mnist_train_idx]
mnist_y_tr = mnist_y_train[mnist_train_idx]

mnist_X_val = mnist_x_train[mnist_val_idx]
mnist_y_val = mnist_y_train[mnist_val_idx]

mnist_X_te = mnist_x_test
mnist_y_te = mnist_y_test

report_split(
    "MNIST split",
    mnist_y_tr,
    mnist_y_val,
    mnist_y_te,
)
```

MNIST split

```
train: 54000
val  : 6000
test : 10000
train class counts: [5331 6068 5362 5518 5258 4879 5326 5638 5266 5354]
val class counts  : [592 674 596 613 584 542 592 627 585 595]
test class counts : [ 980 1135 1032 1010  982  892  958 1028  974 1009]
```

3.4 8. CNN architecture for MNIST

We will intentionally keep one simple architecture across all three implementations:

Input

-> Conv2D

-> ReLU

-> MaxPool

-> Flatten

-> Linear

-> Class logits

The exact API syntax changes, but the computational structure does not.

```
[82]: MNIST_CHANNELS = 1
      MNIST_HEIGHT = 28
      MNIST_WIDTH = 28
      MNIST_CLASSES = 10
      MNIST_CONV_OUT = 16

      MNIST_FLAT_FEATURES = MNIST_CONV_OUT * (MNIST_HEIGHT // 2) * (MNIST_WIDTH // 2)

      print("MNIST classifier input features after Pool + Flatten:",
            ↪MNIST_FLAT_FEATURES)
      print("Expected shape after Conv + Pool:")
      print(f"batch x {MNIST_CONV_OUT} x {MNIST_HEIGHT//2} x {MNIST_WIDTH//2}")
```

MNIST classifier input features after Pool + Flatten: 3136

Expected shape after Conv + Pool:

batch x 16 x 14 x 14

3.5 9. CNN from Scratch with NumPy

The Scratch implementation makes the operations visible.

The forward pass is written explicitly as a composition:

Convolution

-> ReLU

-> MaxPool

-> Flatten

-> Linear

The training loop also makes the learning process visible:

Forward

-> Loss

-> Backward

-> Parameter update

Unlike PyTorch and Keras, no CNN layer class is used to hide the convolution and pooling computations.

```
[83]: def im2col_nchw(x, kernel_size=3, padding=1, stride=1):
      x = np.asarray(x, dtype=np.float32)

      if x.ndim != 4:
          raise ValueError("Expected input with shape N x C x H x W.")

      n, c, h, w = x.shape
      kh = kw = kernel_size

      x_pad = np.pad(
```

```

        x,
        ((0, 0), (0, 0), (padding, padding), (padding, padding)),
        mode="constant",
    )

    windows = np.lib.stride_tricks.sliding_window_view(
        x_pad,
        (kh, kw),
        axis=(2, 3),
    )

    windows = windows[:, :, ::stride, ::stride, :, :]
    h_out = windows.shape[2]
    w_out = windows.shape[3]

    cols = windows.transpose(0, 2, 3, 1, 4, 5).reshape(
        n * h_out * w_out,
        c * kh * kw,
    )

    cache = {
        "input_shape": x.shape,
        "padding": padding,
        "stride": stride,
        "kernel_size": kernel_size,
        "h_out": h_out,
        "w_out": w_out,
    }

    return cols, cache

def col2im_nchw(dcols, cache):
    n, c, h, w = cache["input_shape"]
    padding = cache["padding"]
    stride = cache["stride"]
    kh = kw = cache["kernel_size"]
    h_out = cache["h_out"]
    w_out = cache["w_out"]

    dx_pad = np.zeros(
        (
            n,
            c,
            h + 2 * padding,
            w + 2 * padding,
        ),

```

```

        dtype=np.float32,
    )

    reshaped = dcols.reshape(
        n,
        h_out,
        w_out,
        c,
        kh,
        kw,
    ).transpose(0, 3, 1, 2, 4, 5)

    for i in range(kh):
        for j in range(kw):
            dx_pad[
                :,
                :,
                i : i + h_out * stride : stride,
                j : j + w_out * stride : stride,
            ] += reshaped[:, :, :, :, i, j]

    if padding > 0:
        return dx_pad[:, :, padding:-padding, padding:-padding]

    return dx_pad

def scratch_conv_forward(x, W, b, padding=1, stride=1):
    cols, cache = im2col_nchw(
        x,
        kernel_size=W.shape[2],
        padding=padding,
        stride=stride,
    )

    out = cols @ W.reshape(W.shape[0], -1).T
    out += b.reshape(1, -1)

    n = x.shape[0]
    h_out = cache["h_out"]
    w_out = cache["w_out"]
    filters = W.shape[0]

    out = out.reshape(n, h_out, w_out, filters).transpose(0, 3, 1, 2)

    cache["cols"] = cols
    cache["W"] = W

```

```

cache["b"] = b

return out, cache

def scratch_conv_backward(dout, cache):
    W = cache["W"]
    cols = cache["cols"]
    n, filters, h_out, w_out = dout.shape

    dout_flat = dout.transpose(0, 2, 3, 1).reshape(
        n * h_out * w_out,
        filters,
    )

    dW = dout_flat.T @ cols
    dW = dW.reshape(W.shape)

    db = dout_flat.sum(axis=0).reshape(W.shape[0])

    dcols = dout_flat @ W.reshape(W.shape[0], -1)
    dx = col2im_nchw(dcols, cache)

    return dx, dW, db

def scratch_relu_forward(x):
    out = np.maximum(0.0, x)
    return out, x

def scratch_relu_backward(dout, cache):
    x = cache
    return dout * (x > 0)

def scratch_maxpool_forward(x, pool_size=2, stride=2):
    n, c, h, w = x.shape

    if h % pool_size != 0 or w % pool_size != 0:
        raise ValueError("For this educational implementation, H and W must be
        divisible by pool size.")

    h_out = (h - pool_size) // stride + 1
    w_out = (w - pool_size) // stride + 1

    windows = np.lib.stride_tricks.sliding_window_view(

```

```

        x,
        (pool_size, pool_size),
        axis=(2, 3),
    )

    windows = windows[:, :, ::stride, ::stride, :, :]
    windows = windows[:, :, :h_out, :w_out, :, :]

    out = windows.max(axis=(-1, -2))

    cache = {
        "x_shape": x.shape,
        "windows": windows,
        "out": out,
        "pool_size": pool_size,
        "stride": stride,
    }

    return out, cache

def scratch_maxpool_backward(dout, cache):
    windows = cache["windows"]
    out = cache["out"]
    pool_size = cache["pool_size"]
    stride = cache["stride"]

    n, c, h_out, w_out = out.shape
    dx = np.zeros(cache["x_shape"], dtype=np.float32)

    mask = windows == out[:, :, :, :, None, None]

    grad = dout[:, :, :, :, None, None] * mask
    grad = grad.astype(np.float32)

    for i in range(pool_size):
        for j in range(pool_size):
            dx[
                :,
                :,
                i : i + h_out * stride : stride,
                j : j + w_out * stride : stride,
            ] += grad[:, :, :, :, i, j]

    return dx

```

```

def scratch_flatten_forward(x):
    return x.reshape(x.shape[0], -1), x.shape

def scratch_flatten_backward(dout, cache):
    return dout.reshape(cache)

def scratch_linear_forward(x, W, b):
    out = x @ W + b
    return out, x

def scratch_linear_backward(dout, cache, W):
    x = cache
    dW = x.T @ dout
    db = dout.sum(axis=0)
    dx = dout @ W.T
    return dx, dW, db

```

```

[84]: def softmax(logits):
    shifted = logits - logits.max(axis=1, keepdims=True)
    exp_values = np.exp(shifted)
    return exp_values / exp_values.sum(axis=1, keepdims=True)

def cross_entropy_loss_and_grad(logits, labels):
    probabilities = softmax(logits)
    batch_size = logits.shape[0]

    loss = -np.mean(
        np.log(
            np.clip(
                probabilities[np.arange(batch_size), labels],
                1e-12,
                1.0,
            )
        )
    )

    dlogits = probabilities.copy()
    dlogits[np.arange(batch_size), labels] -= 1.0
    dlogits /= batch_size

    return float(loss), dlogits, probabilities

```

```

def init_scratch_cnn(
    in_channels,
    height,
    width,
    num_classes,
    conv_out=16,
    kernel_size=3,
    seed=42,
):
    rng = np.random.default_rng(seed)

    fan_in = in_channels * kernel_size * kernel_size
    W1 = rng.normal(
        0,
        np.sqrt(2.0 / fan_in),
        size=(conv_out, in_channels, kernel_size, kernel_size),
    ).astype(np.float32)

    b1 = np.zeros(conv_out, dtype=np.float32)

    flat_features = conv_out * (height // 2) * (width // 2)

    W2 = rng.normal(
        0,
        np.sqrt(2.0 / flat_features),
        size=(flat_features, num_classes),
    ).astype(np.float32)

    b2 = np.zeros(num_classes, dtype=np.float32)

    return {
        "W1": W1,
        "b1": b1,
        "W2": W2,
        "b2": b2,
    }

def scratch_cnn_forward(x, params, training=True):
    z1, conv_cache = scratch_conv_forward(
        x,
        params["W1"],
        params["b1"],
        padding=1,
        stride=1,
    )

```

```

a1, relu_cache = scratch_relu_forward(z1)

p1, pool_cache = scratch_maxpool_forward(a1, pool_size=2, stride=2)

flat, flat_cache = scratch_flatten_forward(p1)

logits, linear_cache = scratch_linear_forward(
    flat,
    params["W2"],
    params["b2"],
)

cache = {
    "conv": conv_cache,
    "relu": relu_cache,
    "pool": pool_cache,
    "flatten": flat_cache,
    "linear": linear_cache,
}

return logits, cache

def scratch_cnn_backward(dlogits, params, cache):
    dflat, dW2, db2 = scratch_linear_backward(
        dlogits,
        cache["linear"],
        params["W2"],
    )

    dpool = scratch_flatten_backward(
        dflat,
        cache["flatten"],
    )

    drelu = scratch_maxpool_backward(
        dpool,
        cache["pool"],
    )

    dconv = scratch_relu_backward(
        drelu,
        cache["relu"],
    )

    _, dW1, db1 = scratch_conv_backward(
        dconv,

```

```

        cache["conv"],
    )

    grads = {
        "W1": dW1,
        "b1": db1,
        "W2": dW2,
        "b2": db2,
    }

    return grads

class AdamNumpy:
    def __init__(self, params, lr=1e-3, beta1=0.9, beta2=0.999, eps=1e-8):
        self.lr = lr
        self.beta1 = beta1
        self.beta2 = beta2
        self.eps = eps
        self.t = 0

        self.m = {k: np.zeros_like(v) for k, v in params.items()}
        self.v = {k: np.zeros_like(v) for k, v in params.items()}

    def step(self, params, grads):
        self.t += 1

        for key in params:
            self.m[key] = (
                self.beta1 * self.m[key]
                + (1 - self.beta1) * grads[key]
            )

            self.v[key] = (
                self.beta2 * self.v[key]
                + (1 - self.beta2) * np.square(grads[key])
            )

            m_hat = self.m[key] / (1 - self.beta1 ** self.t)
            v_hat = self.v[key] / (1 - self.beta2 ** self.t)

            params[key] -= (
                self.lr
                * m_hat
                / (np.sqrt(v_hat) + self.eps)
            )

```

```

def minibatches(X, y, batch_size=128, shuffle=True, seed=42):
    indices = np.arange(len(X))

    if shuffle:
        rng = np.random.default_rng(seed)
        rng.shuffle(indices)

    for start in range(0, len(indices), batch_size):
        batch_idx = indices[start:start + batch_size]
        yield X[batch_idx], y[batch_idx]

def train_scratch_cnn(
    X_train,
    y_train,
    X_val,
    y_val,
    in_channels,
    height,
    width,
    num_classes,
    epochs=5,
    batch_size=128,
    lr=1e-3,
    seed=42,
):
    params = init_scratch_cnn(
        in_channels,
        height,
        width,
        num_classes,
        conv_out=16,
        seed=seed,
    )

    optimizer = AdamNumpy(params, lr=lr)

    history = {
        "loss": [],
        "accuracy": [],
        "val_loss": [],
        "val_accuracy": [],
    }

    best_state = None
    best_val = float("inf")

```

```

start_time = time.perf_counter()

for epoch in range(epochs):
    running_loss = 0.0
    correct = 0
    seen = 0

    for batch_X, batch_y in minibatches(
        X_train,
        y_train,
        batch_size=batch_size,
        shuffle=True,
        seed=seed + epoch,
    ):
        logits, cache = scratch_cnn_forward(
            batch_X,
            params,
            training=True,
        )

        loss, dlogits, probabilities = cross_entropy_loss_and_grad(
            logits,
            batch_y,
        )

        grads = scratch_cnn_backward(
            dlogits,
            params,
            cache,
        )

        optimizer.step(params, grads)

        predictions = np.argmax(probabilities, axis=1)

        running_loss += loss * len(batch_X)
        correct += int(np.sum(predictions == batch_y))
        seen += len(batch_X)

    train_loss = running_loss / seen
    train_accuracy = correct / seen

    val_logits, _ = scratch_cnn_forward(
        X_val,
        params,
        training=False,

```

```

    )
    val_loss, _, val_prob = cross_entropy_loss_and_grad(
        val_logits,
        y_val,
    )
    val_pred = np.argmax(val_prob, axis=1)
    val_accuracy = accuracy_score(y_val, val_pred)

    history["loss"].append(float(train_loss))
    history["accuracy"].append(float(train_accuracy))
    history["val_loss"].append(float(val_loss))
    history["val_accuracy"].append(float(val_accuracy))

    print(
        f"Epoch {epoch + 1}/{epochs} | "
        f"loss={train_loss:.4f} | "
        f"acc={train_accuracy:.4f} | "
        f"val_loss={val_loss:.4f} | "
        f"val_acc={val_accuracy:.4f}"
    )

    if val_loss < best_val:
        best_val = val_loss
        best_state = {k: v.copy() for k, v in params.items()}

    elapsed = time.perf_counter() - start_time

    params = best_state

    return params, history, elapsed

def predict_scratch_cnn(X, params, batch_size=256):
    all_prob = []

    for start in range(0, len(X), batch_size):
        batch_X = X[start:start + batch_size]

        logits, _ = scratch_cnn_forward(
            batch_X,
            params,
            training=False,
        )

        all_prob.append(softmax(logits))

    probabilities = np.concatenate(all_prob, axis=0)

```

```
predictions = np.argmax(probabilities, axis=1)

return predictions, probabilities
```

3.6 9.1 Train the Scratch CNN on MNIST

The default settings are intentionally small enough for a normal computer.

The implementation is mathematically explicit:

- convolution uses NumPy tensor operations;
- ReLU is a direct nonlinear function;
- max pooling is implemented explicitly;
- gradients are computed manually;
- Adam updates the parameters.



This is the Scratch interpretation of the phrase “CNN is a composition of functions”.

```
[85]: set_global_seed(SEED)

mnist_scratch_params, mnist_scratch_history, mnist_scratch_time = _
    ↪ train_scratch_cnn(
        mnist_X_tr,
        mnist_y_tr,
        mnist_X_val,
        mnist_y_val,
        in_channels=MNIST_CHANNELS,
        height=MNIST_HEIGHT,
        width=MNIST_WIDTH,
        num_classes=MNIST_CLASSES,
        epochs=EPOCHS,
        batch_size=BATCH_SIZE,
        lr=LEARNING_RATE,
        seed=SEED,
    )

plot_training_history(
    mnist_scratch_history,
    title="MNIST Scratch CNN",
)
```

```
Epoch 1/5 | loss=0.3324 | acc=0.9091 | val_loss=0.1624 | val_acc=0.9567
Epoch 2/5 | loss=0.1197 | acc=0.9663 | val_loss=0.1085 | val_acc=0.9698
Epoch 3/5 | loss=0.0840 | acc=0.9763 | val_loss=0.0931 | val_acc=0.9735
Epoch 4/5 | loss=0.0677 | acc=0.9810 | val_loss=0.0916 | val_acc=0.9723
Epoch 5/5 | loss=0.0584 | acc=0.9833 | val_loss=0.0782 | val_acc=0.9768
```



```
[86]: mnist_scratch_pred, mnist_scratch_prob = predict_scratch_cnn(  
    mnist_X_te,  
    mnist_scratch_params,  
)  
  
mnist_scratch_metrics = compute_classification_metrics(  

```

```

    mnist_y_te,
    mnist_scratch_pred,
    mnist_scratch_prob,
)

mnist_scratch_metrics["parameters"] = _
    ↪ count_numpy_parameters(mnist_scratch_params)
mnist_scratch_metrics["training_time_seconds"] = mnist_scratch_time
mnist_scratch_metrics["epochs"] = EPOCHS

print(pd.Series(mnist_scratch_metrics))

show_confusion_matrix(
    mnist_y_te,
    mnist_scratch_pred,
    IMAGE_CLASS_NAMES["mnist"],
    "MNIST - Scratch CNN confusion matrix",
)

show_classification_report(
    mnist_y_te,
    mnist_scratch_pred,
    IMAGE_CLASS_NAMES["mnist"],
    "MNIST - Scratch CNN per-class metrics",
)


show_confidence_distribution(
    mnist_scratch_prob,
    mnist_y_te == mnist_scratch_pred,
    "MNIST - Scratch CNN confidence",
)

show_prediction_examples(
    mnist_x_test,
    mnist_y_te,
    mnist_scratch_pred,
    mnist_scratch_prob,
    IMAGE_CLASS_NAMES["mnist"],
    "MNIST - Scratch CNN incorrect predictions",
    only_wrong=True,
    n=12,
)



```


accuracy	0.979000
precision_macro	0.978987
recall_macro	0.978883
f1_macro	0.978831
precision_weighted	0.979139

recall_weighted 0.979000
f1_weighted 0.978967
mean_confidence 0.974335
parameters 31530.000000
training_time_seconds 139.493802
epochs 5.000000
dtype: float64



<Figure size 1000x600 with 0 Axes>





3.7 10. CNN with PyTorch

The conceptual model stays the same:

Conv -> ReLU -> Pool -> Flatten -> Linear

PyTorch changes the abstraction:

- `nn.Conv2d` represents convolution.
- `nn.ReLU` represents the nonlinearity.
- `nn.MaxPool2d` represents downsampling.
- `nn.Flatten` converts a tensor into a vector.
- `nn.Linear` performs the final affine transformation.
- `loss.backward()` computes gradients automatically.
- `optimizer.step()` updates parameters.

The `forward()` method makes the data flow explicit.

```
[87]: import torch
import torch.nn as nn
from torch.utils.data import TensorDataset, DataLoader
```

```

DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("PyTorch device:", DEVICE)

class SimpleCNNPyTorch(nn.Module):
    def __init__(self, in_channels, height, width, num_classes):
        super().__init__()

        self.features = nn.Sequential(
            nn.Conv2d(in_channels, 16, kernel_size=3, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2),
        )

        flat_features = 16 * (height // 2) * (width // 2)

        self.classifier = nn.Sequential(
            nn.Flatten(),
            nn.Linear(flat_features, num_classes),
        )

    def forward(self, x):
        x = self.features(x)
        x = self.classifier(x)
        return x

def create_torch_loaders(X_train, y_train, X_val, y_val, X_test, y_test,
    ↪batch_size=128):
    train_ds = TensorDataset(
        torch.from_numpy(X_train),
        torch.from_numpy(y_train),
    )
    val_ds = TensorDataset(
        torch.from_numpy(X_val),
        torch.from_numpy(y_val),
    )
    test_ds = TensorDataset(
        torch.from_numpy(X_test),
        torch.from_numpy(y_test),
    )

    train_loader = DataLoader(train_ds, batch_size=batch_size, shuffle=True)
    val_loader = DataLoader(val_ds, batch_size=batch_size, shuffle=False)
    test_loader = DataLoader(test_ds, batch_size=batch_size, shuffle=False)

    return train_loader, val_loader, test_loader

```

```

def train_torch_cnn(
    X_train, y_train, X_val, y_val, X_test, y_test,
    in_channels, height, width, num_classes,
    epochs=5, batch_size=128, lr=1e-3, seed=42
):
    set_global_seed(seed)

    model = SimpleCNNPyTorch(
        in_channels,
        height,
        width,
        num_classes,
    ).to(DEVICE)

    train_loader, val_loader, test_loader = create_torch_loaders(
        X_train, y_train,
        X_val, y_val,
        X_test, y_test,
        batch_size=batch_size,
    )

    criterion = nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(
        model.parameters(),
        lr=lr,
    )

    history = {
        "loss": [],
        "accuracy": [],
        "val_loss": [],
        "val_accuracy": [],
    }

    best_state = None
    best_val_loss = float("inf")

    start_time = time.perf_counter()

    for epoch in range(epochs):
        model.train()

        train_loss_sum = 0.0
        train_correct = 0
        train_seen = 0

```

```

for images, labels in train_loader:
    images = images.to(DEVICE)
    labels = labels.to(DEVICE)

    optimizer.zero_grad()

    outputs = model(images)
    loss = criterion(outputs, labels)

    loss.backward()
    optimizer.step()

    train_loss_sum += loss.item() * len(images)
    train_correct += int((outputs.argmax(dim=1) == labels).sum().item())
    train_seen += len(images)

model.eval()

val_loss_sum = 0.0
val_correct = 0
val_seen = 0

with torch.no_grad():
    for images, labels in val_loader:
        images = images.to(DEVICE)
        labels = labels.to(DEVICE)

        outputs = model(images)
        loss = criterion(outputs, labels)

        val_loss_sum += loss.item() * len(images)
        val_correct += int((outputs.argmax(dim=1) == labels).sum().
↪item())

        val_seen += len(images)

train_loss = train_loss_sum / train_seen
train_acc = train_correct / train_seen

val_loss = val_loss_sum / val_seen
val_acc = val_correct / val_seen

history["loss"].append(float(train_loss))
history["accuracy"].append(float(train_acc))
history["val_loss"].append(float(val_loss))
history["val_accuracy"].append(float(val_acc))

```

```

        print(
            f"Epoch {epoch + 1}/{epochs} | "
            f"loss={train_loss:.4f} | "
            f"acc={train_acc:.4f} | "
            f"val_loss={val_loss:.4f} | "
            f"val_acc={val_acc:.4f}"
        )

        if val_loss < best_val_loss:
            best_val_loss = val_loss
            best_state = {
                key: value.detach().cpu().clone()
                for key, value in model.state_dict().items()
            }

    elapsed = time.perf_counter() - start_time

    model.load_state_dict(best_state)
    return model, history, elapsed, test_loader

def predict_torch(model, X, batch_size=256):
    loader = DataLoader(
        TensorDataset(torch.from_numpy(X)),
        batch_size=batch_size,
        shuffle=False,
    )

    model.eval()
    all_prob = []

    with torch.no_grad():
        for (images,) in loader:
            images = images.to(DEVICE)
            logits = model(images)
            probabilities = torch.softmax(logits, dim=1)
            all_prob.append(probabilities.cpu().numpy())

    probabilities = np.concatenate(all_prob, axis=0)
    predictions = np.argmax(probabilities, axis=1)

    return predictions, probabilities

```

PyTorch device: cpu

```
[88]: mnist_torch_model, mnist_torch_history, mnist_torch_time, _ = train_torch_cnn(
      mnist_X_tr,
```

```

mnist_y_tr,
mnist_X_val,
mnist_y_val,
mnist_X_te,
mnist_y_te,
in_channels=MNIST_CHANNELS,
height=MNIST_HEIGHT,
width=MNIST_WIDTH,
num_classes=MNIST_CLASSES,
epochs=EPOCHS,
batch_size=BATCH_SIZE,
lr=LEARNING_RATE,
seed=SEED,
)


plot_training_history(
    mnist_torch_history,
    title="MNIST PyTorch CNN",
)


```

```

Epoch 1/5 | loss=0.4405 | acc=0.8808 | val_loss=0.2268 | val_acc=0.9347
Epoch 2/5 | loss=0.1762 | acc=0.9500 | val_loss=0.1610 | val_acc=0.9558
Epoch 3/5 | loss=0.1270 | acc=0.9641 | val_loss=0.1329 | val_acc=0.9647
Epoch 4/5 | loss=0.0998 | acc=0.9717 | val_loss=0.1071 | val_acc=0.9702
Epoch 5/5 | loss=0.0825 | acc=0.9764 | val_loss=0.0927 | val_acc=0.9757

```





```
[89]: mnist_torch_pred, mnist_torch_prob = predict_torch(
    mnist_torch_model,
    mnist_X_te,
)

mnist_torch_metrics = compute_classification_metrics(
    mnist_y_te,
    mnist_torch_pred,
    mnist_torch_prob,
)

mnist_torch_metrics["parameters"] = int(
    sum(p.numel() for p in mnist_torch_model.parameters())
)

mnist_torch_metrics["training_time_seconds"] = mnist_torch_time
mnist_torch_metrics["epochs"] = EPOCHS

print(pd.Series(mnist_torch_metrics))

show_confusion_matrix(
    mnist_y_te,
    mnist_torch_pred,
    IMAGE_CLASS_NAMES["mnist"],
    "MNIST - PyTorch CNN confusion matrix",
)

show_classification_report(
```

```


    mnist_y_te,
    mnist_torch_pred,
    IMAGE_CLASS_NAMES["mnist"],
    "MNIST - PyTorch CNN per-class metrics",
)

show_confidence_distribution(
    mnist_torch_prob,
    mnist_y_te == mnist_torch_pred,
    "MNIST - PyTorch CNN confidence",
)



show_prediction_examples(
    mnist_x_test,
    mnist_y_te,
    mnist_torch_pred,
    mnist_torch_prob,
    IMAGE_CLASS_NAMES["mnist"],
    "MNIST - PyTorch CNN incorrect predictions",
    only_wrong=True,
    n=12,
)


```

accuracy	0.975300
precision_macro	0.975397
recall_macro	0.975098
f1_macro	0.975205
precision_weighted	0.975363
recall_weighted	0.975300
f1_weighted	0.975289
mean_confidence	0.969117
parameters	31530.000000
training_time_seconds	26.547899
epochs	5.000000
dtype: float64	



<Figure size 1000x600 with 0 Axes>





3.8 11. CNN with TensorFlow/Keras

Keras provides a higher level of abstraction.

The same architecture becomes:

Conv2D -> ReLU -> MaxPooling2D -> Flatten -> Dense

The important question from Lecture 04 is:

What computation is hidden inside `model.fit()`?

Keras hides the training loop, automatic differentiation and parameter updates behind the training API.

The mathematics remains the same.

```
[90]: import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers

print("TensorFlow:", tf.__version__)
```

```

def build_keras_cnn(in_channels, height, width, num_classes):
    model = keras.Sequential([
        layers.Input(shape=(height, width, in_channels)),
        layers.Conv2D(16, kernel_size=3, padding="same", activation="relu"),
        layers.MaxPooling2D(pool_size=2),
        layers.Flatten(),
        layers.Dense(num_classes),
    ])

    return model

def train_keras_cnn(
    X_train, y_train, X_val, y_val,
    in_channels, height, width, num_classes,
    epochs=5, batch_size=128, lr=1e-3, seed=42
):
    set_global_seed(seed)

    model = build_keras_cnn(
        in_channels,
        height,
        width,
        num_classes,
    )

    model.compile(
        optimizer=keras.optimizers.Adam(learning_rate=lr),
        loss=keras.losses.SparseCategoricalCrossentropy(from_logits=True),
        metrics=["accuracy"],
    )

    start_time = time.perf_counter()

    history_obj = model.fit(
        np.transpose(X_train, (0, 2, 3, 1)),
        y_train,
        validation_data=(
            np.transpose(X_val, (0, 2, 3, 1)),
            y_val,
        ),
        epochs=epochs,
        batch_size=batch_size,
        verbose=1,
    )

```

```

elapsed = time.perf_counter() - start_time

history = {
    "loss": history_obj.history["loss"],
    "accuracy": history_obj.history["accuracy"],
    "val_loss": history_obj.history["val_loss"],
    "val_accuracy": history_obj.history["val_accuracy"],
}

return model, history, elapsed

```

TensorFlow: 2.21.0

```

[91]: mnist_keras_model, mnist_keras_history, mnist_keras_time = train_keras_cnn(
    mnist_X_tr,
    mnist_y_tr,
    mnist_X_val,
    mnist_y_val,
    in_channels=MNIST_CHANNELS,
    height=MNIST_HEIGHT,
    width=MNIST_WIDTH,
    num_classes=MNIST_CLASSES,
    epochs=EPOCHS,
    batch_size=BATCH_SIZE,
    lr=LEARNING_RATE,
    seed=SEED,
)

plot_training_history(
    mnist_keras_history,
    title="MNIST Keras CNN",
)

```

Epoch 1/5

422/422 5s 10ms/step -

accuracy: 0.8889 - loss: 0.4237 - val_accuracy: 0.9430 - val_loss: 0.2057

Epoch 2/5

422/422 3s 8ms/step -

accuracy: 0.9551 - loss: 0.1566 - val_accuracy: 0.9627 - val_loss: 0.1385

Epoch 3/5

422/422 4s 9ms/step -

accuracy: 0.9691 - loss: 0.1095 - val_accuracy: 0.9693 - val_loss: 0.1102

Epoch 4/5



422/422 4s 9ms/step -

accuracy: 0.9757 - loss: 0.0876 - val_accuracy: 0.9727 - val_loss: 0.0953

Epoch 5/5

422/422 4s 9ms/step -

accuracy: 0.9790 - loss: 0.0749 - val_accuracy: 0.9755 - val_loss: 0.0869



```
[92]: def predict_keras(model, X):  
    logits = model.predict(  
        np.transpose(X, (0, 2, 3, 1)),  
        verbose=0,  
    )
```

```

probabilities = softmax(logits)
predictions = np.argmax(probabilities, axis=1)

return predictions, probabilities

mnist_keras_pred, mnist_keras_prob = predict_keras(
    mnist_keras_model,
    mnist_X_te,
)

mnist_keras_metrics = compute_classification_metrics(
    mnist_y_te,
    mnist_keras_pred,
    mnist_keras_prob,
)

mnist_keras_metrics["parameters"] = int(mnist_keras_model.count_params())
mnist_keras_metrics["training_time_seconds"] = mnist_keras_time
mnist_keras_metrics["epochs"] = EPOCHS

print(pd.Series(mnist_keras_metrics))

show_confusion_matrix(
    mnist_y_te,
    mnist_keras_pred,
    IMAGE_CLASS_NAMES["mnist"],
    "MNIST - Keras CNN confusion matrix",
)

show_classification_report(
    mnist_y_te,
    mnist_keras_pred,
    IMAGE_CLASS_NAMES["mnist"],
    "MNIST - Keras CNN per-class metrics",
)


show_confidence_distribution(
    mnist_keras_prob,
    mnist_y_te == mnist_keras_pred,
    "MNIST - Keras CNN confidence",
)

show_prediction_examples(
    mnist_x_test,
    mnist_y_te,
    mnist_keras_pred,



```


```
mnist_keras_prob,  
IMAGE_CLASS_NAMES["mnist"],  
"MNIST - Keras CNN incorrect predictions",  
only_wrong=True,  
n=12,  
)
```

accuracy	0.978000
precision_macro	0.978049
recall_macro	0.977922
f1_macro	0.977931
precision_weighted	0.978043
recall_weighted	0.978000
f1_weighted	0.977967
mean_confidence	0.970449
parameters	31530.000000
training_time_seconds	19.982689
epochs	5.000000
dtype: float64	



<Figure size 1000x600 with 0 Axes>





3.9 12. MNIST representation visualization

CNNs are useful because they learn visual features.

The following visualization inspects the first convolution layer from PyTorch.

This is directly related to the conceptual hierarchy:

pixels -> edges -> textures -> parts -> objects -> class

The hierarchy is a teaching model, not a strict guarantee for every CNN.

```
[93]: sample_index = 0
sample = torch.from_numpy(mnist_X_te[sample_index:sample_index + 1]).to(DEVICE)


mnist_torch_model.eval()

with torch.no_grad():
    feature_map = mnist_torch_model.features[0](sample)
    feature_map = torch.relu(feature_map).cpu().numpy()[0]

num_maps = min(8, feature_map.shape[0])
```

```
plt.figure(figsize=(12, 6))
for i in range(num_maps):
    ax = plt.subplot(2, 4, i + 1)
    ax.imshow(feature_map[i], cmap="gray")
    ax.set_title(f"Filter {i}")
    ax.axis("off")

plt.suptitle("MNIST - first convolution feature maps")
plt.tight_layout()
plt.show()
```



3.10 13. MNIST framework comparison

The same architecture has now been implemented three times.

Expected conceptual mapping:

Concept	Scratch	Keras	PyTorch
Convolution	manual <code>conv2d</code>	<code>Conv2D</code>	<code>nn.Conv2d</code>
Activation	manual ReLU	ReLU or activation argument	<code>nn.ReLU</code>
Pooling	manual max pool	<code>MaxPooling2D</code>	<code>MaxPool2d</code>
Flatten	manual reshape	<code>Flatten</code>	<code>nn.Flatten</code>
Dense	manual matrix multiplication	<code>Dense</code>	<code>nn.Linear</code>
Loss	manual cross entropy	Keras loss	<code>CrossEntropyLoss</code>
Gradient	manual backpropagation	automatic	autograd
Optimizer	manual parameter update	optimizer API	optimizer API
Training loop	explicit	mostly hidden by <code>fit</code>	explicit

```
[94]: mnist_results = pd.DataFrame([
    {"dataset": "MNIST", "framework": "Scratch", **mnist_scratch_metrics},
    {"dataset": "MNIST", "framework": "PyTorch", **mnist_torch_metrics},
    {"dataset": "MNIST", "framework": "Keras", **mnist_keras_metrics},
])

mnist_results = mnist_results[
    [
        "dataset",
        "framework",
        "accuracy",
        "precision_macro",
        "recall_macro",
        "f1_macro",
        "mean_confidence",
        "parameters",
        "training_time_seconds",
        "epochs",
    ]
]

mnist_results
```

```
[94]:
```

	dataset	framework	accuracy	precision_macro	recall_macro	f1_macro	\
0	MNIST	Scratch	0.9790	0.978987	0.978883	0.978831	
1	MNIST	PyTorch	0.9753	0.975397	0.975098	0.975205	
2	MNIST	Keras	0.9780	0.978049	0.977922	0.977931	

	mean_confidence	parameters	training_time_seconds	epochs
0	0.974335	31530	139.493802	5
1	0.969117	31530	26.547899	5
2	0.970449	31530	19.982689	5

```
[95]: plt.figure(figsize=(10, 5))
plt.bar(
    mnist_results["framework"],
    mnist_results["accuracy"],
)
plt.title("MNIST accuracy by implementation")
plt.xlabel("Implementation")
plt.ylabel("Accuracy")
plt.ylim(0, 1.05)
plt.tight_layout()
plt.show()


plt.figure(figsize=(10, 5))
plt.bar(
```



```

mnist_results["framework"],
mnist_results["training_time_seconds"],
)
plt.title("MNIST training time by implementation")
plt.xlabel("Implementation")
plt.ylabel("Training time in seconds")
plt.tight_layout()
plt.show()

plt.figure(figsize=(10, 5))
plt.bar(
    mnist_results["framework"],
    mnist_results["parameters"],
)
plt.title("MNIST parameter count by implementation")
plt.xlabel("Implementation")
plt.ylabel("Number of parameters")
plt.tight_layout()
plt.show()

```





4 Part C - Dataset 2: CIFAR-10

4.1 14. Load CIFAR-10

CIFAR-10 gives us an important contrast with MNIST:

- MNIST: one grayscale channel.
- CIFAR-10: three RGB channels.
- MNIST: 28 x 28.

- CIFAR-10: 32 x 32.

Data source: <https://www.tensorflow.org/datasets/catalog/cifar10>

This lets us inspect how the number of channels affects tensor shape and parameter count.

```
[96]: cifar_data = np.load(cifar_npz)

cifar_x_train_raw = cifar_data["x_train"]
cifar_y_train = cifar_data["y_train"].astype(np.int64).reshape(-1)

cifar_x_test_raw = cifar_data["x_test"]
cifar_y_test = cifar_data["y_test"].astype(np.int64).reshape(-1)

# Convert from N x H x W x C to N x C x H x W.
cifar_x_train = cifar_x_train_raw.astype(np.float32) / 255.0
cifar_x_test = cifar_x_test_raw.astype(np.float32) / 255.0

cifar_x_train = np.transpose(cifar_x_train, (0, 3, 1, 2))
cifar_x_test = np.transpose(cifar_x_test, (0, 3, 1, 2))

print("Train:", cifar_x_train.shape, cifar_y_train.shape)
print("Test :", cifar_x_test.shape, cifar_y_test.shape)
print("Value range:", cifar_x_train.min(), "to", cifar_x_train.max())
```

Train: (50000, 3, 32, 32) (50000,)



Test : (10000, 3, 32, 32) (10000,)

Value range: 0.0 to 1.0

4.2 15. Explore CIFAR-10

```
[97]: plot_image_grid(
    cifar_x_train[:100],
    cifar_y_train[:100],
    class_names=IMAGE_CLASS_NAMES["cifar10"],
    n=12,
    title="CIFAR-10 sample images",
)

plot_class_distribution(
    cifar_y_train,
    class_names=IMAGE_CLASS_NAMES["cifar10"],
    title="CIFAR-10 training class distribution",
)
```






```
[98]: mean_cifar = cifar_x_train.mean(axis=0)
mean_cifar_display = np.transpose(mean_cifar, (1, 2, 0))

plt.figure(figsize=(5, 5))
plt.imshow(mean_cifar_display)
plt.title("CIFAR-10 mean RGB image")
plt.axis("off")
plt.tight_layout()
plt.show()

for channel, name in enumerate(["Red", "Green", "Blue"]):
    plt.figure(figsize=(9, 4))
    plt.hist(cifar_x_train[:, channel].reshape(-1), bins=40)
    plt.title(f"CIFAR-10 {name} channel distribution")
    plt.xlabel("Normalized intensity")
    plt.ylabel("Frequency")
    plt.tight_layout()
    plt.show()
```

CIFAR-10 mean RGB image





```
[99]: cifar_train_idx, cifar_val_idx = make_stratified_indices(
    cifar_y_train,
    val_size=VALIDATION_SIZE,
    seed=SEED,
    train_limit=COMPARISON_TRAIN_LIMIT if USE_COMPARISON_LIMITS else None,
    val_limit=COMPARISON_VAL_LIMIT if USE_COMPARISON_LIMITS else None,
)

cifar_X_tr = cifar_x_train[cifar_train_idx]
cifar_y_tr = cifar_y_train[cifar_train_idx]

cifar_X_val = cifar_x_train[cifar_val_idx]
cifar_y_val = cifar_y_train[cifar_val_idx]

cifar_X_te = cifar_x_test
cifar_y_te = cifar_y_test

report_split(
    "CIFAR-10 split",
    cifar_y_tr,
    cifar_y_val,
    cifar_y_te,
)
```

```
CIFAR-10 split
train: 45000
val  : 5000
test : 10000
train class counts: [4500 4500 4500 4500 4500 4500 4500 4500 4500 4500]
val class counts  : [500 500 500 500 500 500 500 500 500 500]
test class counts : [1000 1000 1000 1000 1000 1000 1000 1000 1000 1000]
```

4.3 16. CIFAR-10 shape calculation before writing the Dense layer

The lecture requires shape tracking before defining the final dense layer.

For our architecture:

Input: B x 3 x 32 x 32

Conv2d(3, 16, kernel=3, padding=1, stride=1)
-> B x 16 x 32 x 32

MaxPool(2)
-> B x 16 x 16 x 16

Flatten
-> B x 4096

Therefore the final linear layer must receive 4096 input features.

```
[100]: CIFAR_CHANNELS = 3
CIFAR_HEIGHT = 32
CIFAR_WIDTH = 32
CIFAR_CLASSES = 10

CIFAR_FLAT_FEATURES = (
    16
    * (CIFAR_HEIGHT // 2)
    * (CIFAR_WIDTH // 2)
)

print("CIFAR-10 classifier input features:", CIFAR_FLAT_FEATURES)
```

CIFAR-10 classifier input features: 4096

4.4 17. CIFAR-10 Scratch CNN

The same Scratch CNN code is reused with `in_channels=3`.

This demonstrates an important point:



The implementation is the same idea; only the input representation changes.

```
[101]: set_global_seed(SEED)

cifar_scratch_params, cifar_scratch_history, cifar_scratch_time = \
    train_scratch_cnn(
        cifar_X_tr,
        cifar_y_tr,
        cifar_X_val,
        cifar_y_val,
        in_channels=CIFAR_CHANNELS,
        height=CIFAR_HEIGHT,
        width=CIFAR_WIDTH,
        num_classes=CIFAR_CLASSES,
        epochs=EPOCHS,
        batch_size=BATCH_SIZE,
        lr=LEARNING_RATE,
        seed=SEED,
    )

plot_training_history(
    cifar_scratch_history,
    title="CIFAR-10 Scratch CNN",
)
```

Epoch 1/5 | loss=1.6493 | acc=0.4161 | val_loss=1.4475 | val_acc=0.4874
Epoch 2/5 | loss=1.3824 | acc=0.5182 | val_loss=1.3629 | val_acc=0.5206
Epoch 3/5 | loss=1.2800 | acc=0.5567 | val_loss=1.2801 | val_acc=0.5578
Epoch 4/5 | loss=1.2186 | acc=0.5791 | val_loss=1.2406 | val_acc=0.5724
Epoch 5/5 | loss=1.1671 | acc=0.5976 | val_loss=1.2219 | val_acc=0.5798



```

[102]: cifar_scratch_pred, cifar_scratch_prob = predict_scratch_cnn(
        cifar_X_te,
        cifar_scratch_params,
    )

    cifar_scratch_metrics = compute_classification_metrics(
        cifar_y_te,
        cifar_scratch_pred,
        cifar_scratch_prob,
    )

    cifar_scratch_metrics["parameters"] = _
        ↪count_numpy_parameters(cifar_scratch_params)
    cifar_scratch_metrics["training_time_seconds"] = cifar_scratch_time
    cifar_scratch_metrics["epochs"] = EPOCHS

    print(pd.Series(cifar_scratch_metrics))

    show_confusion_matrix(
        cifar_y_te,
        cifar_scratch_pred,
        IMAGE_CLASS_NAMES["cifar10"],
        "CIFAR-10 - Scratch CNN confusion matrix",
    )

    show_classification_report(
        cifar_y_te,
        cifar_scratch_pred,
        IMAGE_CLASS_NAMES["cifar10"],
        "CIFAR-10 - Scratch CNN per-class metrics",
    )


    show_confidence_distribution(
        cifar_scratch_prob,
        cifar_y_te == cifar_scratch_pred,
        "CIFAR-10 - Scratch CNN confidence",
    )

    show_prediction_examples(
        cifar_x_test,
        cifar_y_te,
        cifar_scratch_pred,
        cifar_scratch_prob,
        IMAGE_CLASS_NAMES["cifar10"],
        "CIFAR-10 - Scratch CNN incorrect predictions",
        only_wrong=True,
        n=12,
    )



```


```
)
```

```
accuracy          0.572900
precision_macro    0.588441
recall_macro       0.572900
f1_macro           0.572666
precision_weighted 0.588441
recall_weighted    0.572900
f1_weighted        0.572666
mean_confidence    0.546630
parameters         41418.000000
training_time_seconds 188.094851
epochs            5.000000
dtype: float64
```



<Figure size 1000x600 with 0 Axes>







4.5 18. CIFAR-10 PyTorch CNN

```
[103]: cifar_torch_model, cifar_torch_history, cifar_torch_time, _ = train_torch_cnn(
    cifar_X_tr,
    cifar_y_tr,
    cifar_X_val,
    cifar_y_val,
    cifar_X_te,
    cifar_y_te,
    in_channels=CIFAR_CHANNELS,
    height=CIFAR_HEIGHT,
    width=CIFAR_WIDTH,
    num_classes=CIFAR_CLASSES,
    epochs=EPOCHS,
    batch_size=BATCH_SIZE,
    lr=LEARNING_RATE,
    seed=SEED,
)
```

```
plot_training_history(  
    cifar_torch_history,  
    title="CIFAR-10 PyTorch CNN",  
)
```

Epoch 1/5		loss=1.6746		acc=0.4147		val_loss=1.5004		val_acc=0.4682
Epoch 2/5		loss=1.3811		acc=0.5211		val_loss=1.3404		val_acc=0.5328
Epoch 3/5		loss=1.2753		acc=0.5557		val_loss=1.2555		val_acc=0.5648
Epoch 4/5		loss=1.2100		acc=0.5798		val_loss=1.2175		val_acc=0.5712
Epoch 5/5		loss=1.1673		acc=0.5951		val_loss=1.2012		val_acc=0.5774





```
[104]: cifar_torch_pred, cifar_torch_prob = predict_torch(
        cifar_torch_model,
        cifar_X_te,
    )

    cifar_torch_metrics = compute_classification_metrics(
        cifar_y_te,
        cifar_torch_pred,
        cifar_torch_prob,
    )

    cifar_torch_metrics["parameters"] = int(
        sum(p.numel() for p in cifar_torch_model.parameters())
    )
    cifar_torch_metrics["training_time_seconds"] = cifar_torch_time
    cifar_torch_metrics["epochs"] = EPOCHS

    print(pd.Series(cifar_torch_metrics))

    show_confusion_matrix(
        cifar_y_te,
        cifar_torch_pred,
        IMAGE_CLASS_NAMES["cifar10"],
        "CIFAR-10 - PyTorch CNN confusion matrix",
    )

    show_classification_report(
```

```


    cifar_y_te,
    cifar_torch_pred,
    IMAGE_CLASS_NAMES["cifar10"],
    "CIFAR-10 - PyTorch CNN per-class metrics",
)

show_confidence_distribution(
    cifar_torch_prob,
    cifar_y_te == cifar_torch_pred,
    "CIFAR-10 - PyTorch CNN confidence",
)



show_prediction_examples(
    cifar_x_test,
    cifar_y_te,
    cifar_torch_pred,
    cifar_torch_prob,
    IMAGE_CLASS_NAMES["cifar10"],
    "CIFAR-10 - PyTorch CNN incorrect predictions",
    only_wrong=True,
    n=12,
)

```


accuracy	0.573700
precision_macro	0.576051
recall_macro	0.573700
f1_macro	0.564701
precision_weighted	0.576051
recall_weighted	0.573700
f1_weighted	0.564701
mean_confidence	0.561211
parameters	41418.000000
training_time_seconds	45.944586
epochs	5.000000
dtype: float64	



<Figure size 1000x600 with 0 Axes>



CIFAR-10 - PyTorch CNN incorrect predictions




4.6 19. CIFAR-10 Keras CNN


```
[105]: cifar_keras_model, cifar_keras_history, cifar_keras_time = train_keras_cnn(
    cifar_X_tr,
    cifar_y_tr,
    cifar_X_val,
    cifar_y_val,
    in_channels=CIFAR_CHANNELS,
    height=CIFAR_HEIGHT,
    width=CIFAR_WIDTH,
    num_classes=CIFAR_CLASSES,
    epochs=EPOCHS,
    batch_size=BATCH_SIZE,
    lr=LEARNING_RATE,
    seed=SEED,
)

plot_training_history(
    cifar_keras_history,
    title="CIFAR-10 Keras CNN",
```

)

Epoch 1/5
352/352 5s 11ms/step -
accuracy: 0.4182 - loss: 1.6689 - val_accuracy: 0.4896 - val_loss: 1.4667
Epoch 2/5
352/352 4s 10ms/step -
accuracy: 0.5198 - loss: 1.3804 - val_accuracy: 0.5318 - val_loss: 1.3557
Epoch 3/5
352/352 4s 10ms/step -
accuracy: 0.5539 - loss: 1.2857 - val_accuracy: 0.5474 - val_loss: 1.3006
Epoch 4/5
352/352 4s 10ms/step -
accuracy: 0.5745 - loss: 1.2309 - val_accuracy: 0.5628 - val_loss: 1.2655
Epoch 5/5
352/352 4s 11ms/step -
accuracy: 0.5906 - loss: 1.1907 - val_accuracy: 0.5750 - val_loss: 1.2405





```
[106]: cifar_keras_pred, cifar_keras_prob = predict_keras(
        cifar_keras_model,
        cifar_X_te,
    )

    cifar_keras_metrics = compute_classification_metrics(
        cifar_y_te,
        cifar_keras_pred,
        cifar_keras_prob,
    )

    cifar_keras_metrics["parameters"] = int(cifar_keras_model.count_params())
    cifar_keras_metrics["training_time_seconds"] = cifar_keras_time
    cifar_keras_metrics["epochs"] = EPOCHS

    print(pd.Series(cifar_keras_metrics))

    show_confusion_matrix(
        cifar_y_te,
        cifar_keras_pred,
        IMAGE_CLASS_NAMES["cifar10"],
        "CIFAR-10 - Keras CNN confusion matrix",
    )

    show_classification_report(
        cifar_y_te,
        cifar_keras_pred,
```

```

    IMAGE_CLASS_NAMES["cifar10"],
    "CIFAR-10 - Keras CNN per-class metrics",
)

show_confidence_distribution(
    cifar_keras_prob,
    cifar_y_te == cifar_keras_pred,
    "CIFAR-10 - Keras CNN confidence",
)


show_prediction_examples(
    cifar_x_test,
    cifar_y_te,
    cifar_keras_pred,
    cifar_keras_prob,
    IMAGE_CLASS_NAMES["cifar10"],
    "CIFAR-10 - Keras CNN incorrect predictions",
    only_wrong=True,
    n=12,
)

```



```


accuracy                0.568100
precision_macro          0.574029
recall_macro            0.568100
f1_macro                0.563815
precision_weighted       0.574029
recall_weighted          0.568100
f1_weighted             0.563815
mean_confidence          0.548039
parameters              41418.000000
training_time_seconds    19.185502
epochs                  5.000000
dtype: float64

```



<Figure size 1000x600 with 0 Axes>





4.7 20. CIFAR-10 representation visualization

Because CIFAR-10 is RGB, the learned feature maps are especially useful for showing that a convolution processes all input channels together while producing new feature channels.

```
[107]: sample_index = 0
sample = torch.from_numpy(cifar_X_te[sample_index:sample_index + 1]).to(DEVICE)

cifar_torch_model.eval()

with torch.no_grad():
    feature_map = cifar_torch_model.features[0](sample)
    feature_map = torch.relu(feature_map).cpu().numpy()[0]

num_maps = min(8, feature_map.shape[0])


plt.figure(figsize=(12, 6))
for i in range(num_maps):
    ax = plt.subplot(2, 4, i + 1)
    ax.imshow(feature_map[i], cmap="gray")
```

```

ax.set_title(f"Filter {i}")
ax.axis("off")

plt.suptitle("CIFAR-10 - first convolution feature maps")
plt.tight_layout()
plt.show()

```



```

[108]: cifar_results = pd.DataFrame([
    {"dataset": "CIFAR-10", "framework": "Scratch", **cifar_scratch_metrics},
    {"dataset": "CIFAR-10", "framework": "PyTorch", **cifar_torch_metrics},
    {"dataset": "CIFAR-10", "framework": "Keras", **cifar_keras_metrics},
])

cifar_results = cifar_results[
    [
        "dataset",
        "framework",
        "accuracy",
        "precision_macro",
        "recall_macro",
        "f1_macro",
        "mean_confidence",
        "parameters",
        "training_time_seconds",
        "epochs",
    ]
]

```

```

display(cifar_results)


plt.figure(figsize=(10, 5))
plt.bar(cifar_results["framework"], cifar_results["accuracy"])
plt.title("CIFAR-10 accuracy by implementation")
plt.xlabel("Implementation")
plt.ylabel("Accuracy")
plt.ylim(0, 1.05)
plt.tight_layout()
plt.show()


plt.figure(figsize=(10, 5))
plt.bar(cifar_results["framework"], cifar_results["f1_macro"])
plt.title("CIFAR-10 macro F1 by implementation")
plt.xlabel("Implementation")
plt.ylabel("Macro F1")
plt.ylim(0, 1.05)
plt.tight_layout()
plt.show()

```

	dataset	framework	accuracy	precision_macro	recall_macro	f1_macro	\
0	CIFAR-10	Scratch	0.5729	0.588441	0.5729	0.572666	
1	CIFAR-10	PyTorch	0.5737	0.576051	0.5737	0.564701	
2	CIFAR-10	Keras	0.5681	0.574029	0.5681	0.563815	

	mean_confidence	parameters	training_time_seconds	epochs
0	0.546630	41418	188.094851	5
1	0.561211	41418	45.944586	5
2	0.548039	41418	19.185502	5





5 Part D - Dataset 3: Diabetes

5.1 21. Load and inspect Diabetes

This dataset is tabular data, not image data.

A sample is represented as a feature vector:

$x = [x_1, x_2, \dots, x_d]$

There is normally no meaningful 2D spatial neighborhood between adjacent tabular columns.

Therefore the natural neural-network baseline is an MLP:

Features -> Linear -> ReLU -> Linear -> ReLU -> Linear -> Prediction

Data source: <https://www.kaggle.com/datasets/alexteboul/diabetes-health-indicators-dataset>

This is exactly the distinction emphasized in Lecture 04.

```
[109]: diabetes_df = pd.read_csv(DIABETES_PATH)

print("Shape:", diabetes_df.shape)
print("\nColumns:")
print(diabetes_df.columns.tolist())

display(diabetes_df.head())
```

Shape: (253680, 22)

Columns:

['Diabetes_012', 'HighBP', 'HighChol', 'CholCheck', 'BMI', 'Smoker', 'Stroke', 'HeartDiseaseorAttack', 'PhysActivity', 'Fruits', 'Veggies',

```
'HvyAlcoholConsump', 'AnyHealthcare', 'NoDocbcCost', 'GenHlth', 'MentHlth',
'PhysHlth', 'DiffWalk', 'Sex', 'Age', 'Education', 'Income']
```

	Diabetes_012	HighBP	HighChol	CholCheck	BMI	Smoker	Stroke	\
0	0.0	1.0	1.0	1.0	40.0	1.0	0.0	
1	0.0	0.0	0.0	0.0	25.0	1.0	0.0	
2	0.0	1.0	1.0	1.0	28.0	0.0	0.0	
3	0.0	1.0	0.0	1.0	27.0	0.0	0.0	
4	0.0	1.0	1.0	1.0	24.0	0.0	0.0	

	HeartDiseaseorAttack	PhysActivity	Fruits	...	AnyHealthcare	\
0	0.0	0.0	0.0	...	1.0	
1	0.0	1.0	0.0	...	0.0	
2	0.0	0.0	1.0	...	1.0	
3	0.0	1.0	1.0	...	1.0	
4	0.0	1.0	1.0	...	1.0	

	NoDocbcCost	GenHlth	MentHlth	PhysHlth	DiffWalk	Sex	Age	Education	\
0	0.0	5.0	18.0	15.0	1.0	0.0	9.0	4.0	
1	1.0	3.0	0.0	0.0	0.0	0.0	7.0	6.0	
2	1.0	5.0	30.0	30.0	1.0	0.0	9.0	4.0	
3	0.0	2.0	0.0	0.0	0.0	0.0	11.0	3.0	
4	0.0	2.0	3.0	0.0	0.0	0.0	11.0	5.0	

	Income
0	3.0
1	1.0
2	8.0
3	6.0
4	4.0

[5 rows x 22 columns]

```
[110]: print("Missing values by column:")
missing = diabetes_df.isna().sum().sort_values(ascending=False)
display(missing[missing > 0].to_frame("missing_count"))

print("Duplicate rows:", diabetes_df.duplicated().sum())

print("Data types:")
display(diabetes_df.dtypes.to_frame("dtype"))

print("File size MB:", round(DIABETES_PATH.stat().st_size / (1024**2), 2))
```

Missing values by column:

Empty DataFrame

Columns: [missing_count]

Index: []

Duplicate rows: 23899

Data types:

	dtype
Diabetes_012	float64
HighBP	float64
HighChol	float64
CholCheck	float64
BMI	float64
Smoker	float64
Stroke	float64
HeartDiseaseorAttack	float64
PhysActivity	float64
Fruits	float64
Veggies	float64
HvyAlcoholConsump	float64
AnyHealthcare	float64
NoDocbcCost	float64
GenHlth	float64
MentHlth	float64
PhysHlth	float64
DiffWalk	float64
Sex	float64
Age	float64
Education	float64
Income	float64

File size MB: 21.68

5.2 22. Diabetes target detection

The preferred target column for the selected dataset is `Diabetes_012`.

The code below keeps a small fallback list so the notebook can fail clearly if a different Diabetes file is placed in the folder.

```
[111]: TARGET_CANDIDATES = [
        "Diabetes_012",
        "Diabetes_binary",
        "diabetes",
        "target",
        "Target",
    ]

    target_column = next(
        (name for name in TARGET_CANDIDATES if name in diabetes_df.columns),
        None,
    )
```

```

if target_column is None:
    diabetes_like = [
        col for col in diabetes_df.columns
        if "diabetes" in col.lower()
    ]
    if len(diabetes_like) == 1:
        target_column = diabetes_like[0]

if target_column is None:
    raise ValueError(
        "Cannot identify the Diabetes target column. "
        "Expected Diabetes_012 or another target column."
    )

print("Target column:", target_column)

X_diabetes_df = diabetes_df.drop(columns=[target_column]).copy()
y_diabetes = diabetes_df[target_column].astype(int).to_numpy()

# Keep numeric columns only for this assignment.
non_numeric = X_diabetes_df.select_dtypes(exclude=[np.number]).columns.tolist()

if non_numeric:
    print("Dropping non-numeric columns:", non_numeric)
    X_diabetes_df = X_diabetes_df.drop(columns=non_numeric)

X_diabetes = X_diabetes_df.to_numpy(dtype=np.float32)

print("Input features:", X_diabetes.shape)
print("Target shape:", y_diabetes.shape)
print("Classes:", np.unique(y_diabetes))

```

```

Target column: Diabetes_012
Input features: (253680, 21)
Target shape: (253680,)
Classes: [0 1 2]

```

5.3 22.1 Diabetes data exploration


The visualization set is intentionally broader than a single histogram.

We inspect:

1. Target balance.
2. Feature distributions.
3. Correlation structure.
4. A PCA projection.
5. Example feature distributions by target.

This helps explain both the dataset and the later model behavior.

```
[112]: plot_class_distribution(
        y_diabetes,
        class_names=[str(v) for v in np.unique(y_diabetes)],
        title="Diabetes target distribution",
    )
```



```
[113]: display(X_diabetes_df.describe().T)




numeric_features_for_plot = X_diabetes_df.columns[:8]



for feature in numeric_features_for_plot:
    plt.figure(figsize=(9, 4))
    plt.hist(X_diabetes_df[feature], bins=30)
    plt.title(f"Distribution of {feature}")
    plt.xlabel(feature)
    plt.ylabel("Frequency")
    plt.tight_layout()
    plt.show()
```




	count	mean	std	min	25%	50%	75%	\
HighBP	253680.0	0.429001	0.494934	0.0	0.0	0.0	1.0	
HighChol	253680.0	0.424121	0.494210	0.0	0.0	0.0	1.0	
CholCheck	253680.0	0.962670	0.189571	0.0	1.0	1.0	1.0	
BMI	253680.0	28.382364	6.608694	12.0	24.0	27.0	31.0	
Smoker	253680.0	0.443169	0.496761	0.0	0.0	0.0	1.0	
Stroke	253680.0	0.040571	0.197294	0.0	0.0	0.0	0.0	
HeartDiseaseorAttack	253680.0	0.094186	0.292087	0.0	0.0	0.0	0.0	
PhysActivity	253680.0	0.756544	0.429169	0.0	1.0	1.0	1.0	
Fruits	253680.0	0.634256	0.481639	0.0	0.0	1.0	1.0	

Veggies	253680.0	0.811420	0.391175	0.0	1.0	1.0	1.0
HvyAlcoholConsump	253680.0	0.056197	0.230302	0.0	0.0	0.0	0.0
AnyHealthcare	253680.0	0.951053	0.215759	0.0	1.0	1.0	1.0
NoDocbcCost	253680.0	0.084177	0.277654	0.0	0.0	0.0	0.0
GenHlth	253680.0	2.511392	1.068477	1.0	2.0	2.0	3.0
MentHlth	253680.0	3.184772	7.412847	0.0	0.0	0.0	2.0
PhysHlth	253680.0	4.242081	8.717951	0.0	0.0	0.0	3.0
DiffWalk	253680.0	0.168224	0.374066	0.0	0.0	0.0	0.0
Sex	253680.0	0.440342	0.496429	0.0	0.0	0.0	1.0
Age	253680.0	8.032119	3.054220	1.0	6.0	8.0	10.0
Education	253680.0	5.050434	0.985774	1.0	4.0	5.0	6.0
Income	253680.0	6.053875	2.071148	1.0	5.0	7.0	8.0

	max
HighBP	1.0
HighChol	1.0
CholCheck	1.0
BMI	98.0
Smoker	1.0
Stroke	1.0
HeartDiseaseorAttack	1.0
PhysActivity	1.0
Fruits	1.0
Veggies	1.0
HvyAlcoholConsump	1.0
AnyHealthcare	1.0
NoDocbcCost	1.0
GenHlth	5.0
MentHlth	30.0
PhysHlth	30.0
DiffWalk	1.0
Sex	1.0
Age	13.0
Education	6.0
Income	8.0








```
[114]: corr = X_diabetes_df.corr(numeric_only=True).to_numpy()
corr_labels = X_diabetes_df.corr(numeric_only=True).columns.tolist()

plt.figure(figsize=(12, 10))
plt.imshow(corr, aspect="auto")
plt.colorbar(label="Correlation")
plt.xticks(range(len(corr_labels)), corr_labels, rotation=90, fontsize=7)
plt.yticks(range(len(corr_labels)), corr_labels, fontsize=7)
plt.title("Diabetes feature correlation matrix")
plt.tight_layout()
plt.show()
```



```
[115]: # PCA is used only for visualization.
scaler_for_pca = StandardScaler()
```

```

X_diabetes_scaled_for_pca = scaler_for_pca.fit_transform(X_diabetes)


pca = PCA(n_components=2, random_state=SEED)
X_pca = pca.fit_transform(X_diabetes_scaled_for_pca)

plt.figure(figsize=(9, 6))
for label in np.unique(y_diabetes):
    mask = y_diabetes == label
    plt.scatter(
        X_pca[mask, 0],
        X_pca[mask, 1],
        s=8,
        alpha=0.4,
        label=f"Class {label}",
    )

plt.xlabel("Principal component 1")
plt.ylabel("Principal component 2")
plt.title("Diabetes PCA projection")
plt.legend()
plt.tight_layout()
plt.show()

print("Explained variance ratio:", pca.explained_variance_ratio_)

```




Explained variance ratio: [0.16713002 0.083967]



```
[116]: # A few feature distributions grouped by target.
selected_features = X_diabetes_df.columns[:4]


for feature in selected_features:
    plt.figure(figsize=(10, 5))

    for label in np.unique(y_diabetes):
        values = X_diabetes_df.loc[y_diabetes == label, feature]
        plt.hist(
            values,
            bins=25,
            alpha=0.45,
            label=f"Class {label}",
        )

    plt.title(f"{feature} distribution by target class")
    plt.xlabel(feature)
    plt.ylabel("Frequency")
    plt.legend()
    plt.tight_layout()
    plt.show()
```







5.4 23. Diabetes preprocessing and split

The order matters:

Split -> fit scaler on training data only -> transform train/validation/test

This avoids data leakage.

The same split and same scaler are then used conceptually for Scratch, PyTorch and Keras.

```
[117]: diabetes_train_idx, diabetes_test_idx = train_test_split(
    np.arange(len(y_diabetes)),
    test_size=0.20,
    stratify=y_diabetes,
    random_state=SEED,
)

diabetes_train_idx, diabetes_val_idx = train_test_split(
    diabetes_train_idx,
    test_size=VALIDATION_SIZE,
    stratify=y_diabetes[diabetes_train_idx],
    random_state=SEED,
)

if USE_COMPARISON_LIMITS:
    if COMPARISON_TRAIN_LIMIT is not None and COMPARISON_TRAIN_LIMIT <
    len(diabetes_train_idx):
        diabetes_train_idx, _ = train_test_split(
            diabetes_train_idx,
            train_size=COMPARISON_TRAIN_LIMIT,
```

```

        stratify=y_diabetes[diabetes_train_idx],
        random_state=SEED,
    )

    if COMPARISON_VAL_LIMIT is not None and COMPARISON_VAL_LIMIT <
    len(diabetes_val_idx):
        diabetes_val_idx, _ = train_test_split(
            diabetes_val_idx,
            train_size=COMPARISON_VAL_LIMIT,
            stratify=y_diabetes[diabetes_val_idx],
            random_state=SEED,
        )

X_diabetes_tr_raw = X_diabetes[diabetes_train_idx]
y_diabetes_tr = y_diabetes[diabetes_train_idx]

X_diabetes_val_raw = X_diabetes[diabetes_val_idx]
y_diabetes_val = y_diabetes[diabetes_val_idx]

X_diabetes_te_raw = X_diabetes[diabetes_test_idx]
y_diabetes_te = y_diabetes[diabetes_test_idx]

diabetes_scaler = StandardScaler()
X_diabetes_tr = diabetes_scaler.fit_transform(X_diabetes_tr_raw).astype(np.
    float32)
X_diabetes_val = diabetes_scaler.transform(X_diabetes_val_raw).astype(np.
    float32)
X_diabetes_te = diabetes_scaler.transform(X_diabetes_te_raw).astype(np.float32)

report_split(
    "Diabetes split",
    y_diabetes_tr,
    y_diabetes_val,
    y_diabetes_te,
)

classes = np.unique(y_diabetes_tr)
diabetes_class_weights = compute_class_weight(
    class_weight="balanced",
    classes=classes,
    y=y_diabetes_tr,
)

diabetes_class_weights = np.asarray(
    diabetes_class_weights,
    dtype=np.float32,
)

```

```
print("Class weights:", dict(zip(classes, diabetes_class_weights)))
```

```
Diabetes split
  train: 182649
  val   : 20295
  test  : 50736
  train class counts: [153865   3335  25449]
  val class counts   : [17097    370   2828]
  test class counts  : [42741    926   7069]
Class weights: {np.int64(0): np.float32(0.39569104), np.int64(1):
np.float32(18.255772), np.int64(2): np.float32(2.3923533)}
```

5.5 24. MLP from Scratch for Diabetes

The manual MLP is intentionally simple:

Input

```
-> Linear(64)
-> ReLU
-> Linear(32)
-> ReLU
-> Linear(number_of_classes)
```

This demonstrates the same function-composition idea without forcing an image-specific operation onto tabular data.

```
[118]: def init_scratch_mlp(input_dim, hidden1, hidden2, num_classes, seed=42):
        rng = np.random.default_rng(seed)

        params = {}

        params["W1"] = rng.normal(
            0,
            np.sqrt(2.0 / input_dim),
            size=(input_dim, hidden1),
        ).astype(np.float32)
        params["b1"] = np.zeros(hidden1, dtype=np.float32)

        params["W2"] = rng.normal(
            0,
            np.sqrt(2.0 / hidden1),
            size=(hidden1, hidden2),
        ).astype(np.float32)
        params["b2"] = np.zeros(hidden2, dtype=np.float32)

        params["W3"] = rng.normal(
            0,
            np.sqrt(2.0 / hidden2),
```

```

        size=(hidden2, num_classes),
    ).astype(np.float32)
    params["b3"] = np.zeros(num_classes, dtype=np.float32)

    return params

def weighted_cross_entropy(logits, labels, class_weights):
    probabilities = softmax(logits)
    batch_size = len(labels)

    sample_weights = class_weights[labels]
    normalizer = np.mean(sample_weights)

    true_probs = probabilities[
        np.arange(batch_size),
        labels,
    ]

    loss = -np.mean(
        sample_weights
        * np.log(np.clip(true_probs, 1e-12, 1.0))
    ) / normalizer

    dlogits = probabilities.copy()
    dlogits[np.arange(batch_size), labels] -= 1.0
    dlogits *= sample_weights[:, None]
    dlogits /= (batch_size * normalizer)

    return float(loss), dlogits, probabilities

def train_scratch_mlp(
    X_train, y_train, X_val, y_val,
    input_dim, num_classes,
    class_weights,
    epochs=5,
    batch_size=128,
    lr=1e-3,
    seed=42,
):
    params = init_scratch_mlp(
        input_dim=input_dim,
        hidden1=64,
        hidden2=32,
        num_classes=num_classes,
        seed=seed,

```

```

)

optimizer = AdamNumpy(params, lr=lr)

history = {
    "loss": [],
    "accuracy": [],
    "val_loss": [],
    "val_accuracy": [],
}

best_state = None
best_val_loss = float("inf")

start_time = time.perf_counter()

for epoch in range(epochs):
    train_loss_sum = 0.0
    train_correct = 0
    train_seen = 0

    for batch_X, batch_y in minibatches(
        X_train,
        y_train,
        batch_size=batch_size,
        shuffle=True,
        seed=seed + epoch,
    ):
        z1 = batch_X @ params["W1"] + params["b1"]
        a1 = np.maximum(0.0, z1)

        z2 = a1 @ params["W2"] + params["b2"]
        a2 = np.maximum(0.0, z2)

        logits = a2 @ params["W3"] + params["b3"]

        loss, dlogits, probabilities = weighted_cross_entropy(
            logits,
            batch_y,
            class_weights,
        )

        dW3 = a2.T @ dlogits
        db3 = dlogits.sum(axis=0)
        da2 = dlogits @ params["W3"].T

        dz2 = da2 * (z2 > 0)

```

```

dW2 = a1.T @ dz2
db2 = dz2.sum(axis=0)
da1 = dz2 @ params["W2"].T

dz1 = da1 * (z1 > 0)
dW1 = batch_X.T @ dz1
db1 = dz1.sum(axis=0)

grads = {
    "W1": dW1,
    "b1": db1,
    "W2": dW2,
    "b2": db2,
    "W3": dW3,
    "b3": db3,
}

optimizer.step(params, grads)

predictions = np.argmax(probabilities, axis=1)

train_loss_sum += loss * len(batch_X)
train_correct += int(np.sum(predictions == batch_y))
train_seen += len(batch_X)

train_loss = train_loss_sum / train_seen
train_acc = train_correct / train_seen

z1 = X_val @ params["W1"] + params["b1"]
a1 = np.maximum(0.0, z1)
z2 = a1 @ params["W2"] + params["b2"]
a2 = np.maximum(0.0, z2)
val_logits = a2 @ params["W3"] + params["b3"]

val_loss, _, val_prob = weighted_cross_entropy(
    val_logits,
    y_val,
    class_weights,
)

val_pred = np.argmax(val_prob, axis=1)
val_acc = accuracy_score(y_val, val_pred)

history["loss"].append(float(train_loss))
history["accuracy"].append(float(train_acc))
history["val_loss"].append(float(val_loss))
history["val_accuracy"].append(float(val_acc))

```

```

    print(
        f"Epoch {epoch + 1}/{epochs} | "
        f"loss={train_loss:.4f} | "
        f"acc={train_acc:.4f} | "
        f"val_loss={val_loss:.4f} | "
        f"val_acc={val_acc:.4f}"
    )

    if val_loss < best_val_loss:
        best_val_loss = val_loss
        best_state = {k: v.copy() for k, v in params.items()}

    elapsed = time.perf_counter() - start_time

    return best_state, history, elapsed

def predict_scratch_mlp(X, params):
    z1 = X @ params["W1"] + params["b1"]
    a1 = np.maximum(0.0, z1)

    z2 = a1 @ params["W2"] + params["b2"]
    a2 = np.maximum(0.0, z2)

    logits = a2 @ params["W3"] + params["b3"]

    probabilities = softmax(logits)
    predictions = np.argmax(probabilities, axis=1)

    return predictions, probabilities

```

```

[119]: set_global_seed(SEED)


diabetes_num_classes = len(np.unique(y_diabetes_tr))


diabetes_scratch_params, diabetes_scratch_history, diabetes_scratch_time = ↪ train_scratch_mlp(
    X_diabetes_tr,
    y_diabetes_tr,
    X_diabetes_val,
    y_diabetes_val,
    input_dim=X_diabetes_tr.shape[1],
    num_classes=diabetes_num_classes,
    class_weights=diabetes_class_weights,
    epochs=EPOCHS,
    batch_size=BATCH_SIZE,
    lr=LEARNING_RATE,

```

```
seed=SEED,  
)  
  
plot_training_history(  
    diabetes_scratch_history,  
    title="Diabetes Scratch MLP",  
)
```

Epoch 1/5		loss=0.9845		acc=0.6285		val_loss=0.9587		val_acc=0.5984
Epoch 2/5		loss=0.9425		acc=0.6467		val_loss=0.9529		val_acc=0.6455
Epoch 3/5		loss=0.9358		acc=0.6469		val_loss=0.9564		val_acc=0.6707
Epoch 4/5		loss=0.9317		acc=0.6464		val_loss=0.9561		val_acc=0.6451
Epoch 5/5		loss=0.9280		acc=0.6445		val_loss=0.9601		val_acc=0.6502





```
[120]: diabetes_scratch_pred, diabetes_scratch_prob = predict_scratch_mlp(
        X_diabetes_te,
        diabetes_scratch_params,
    )

    diabetes_scratch_metrics = compute_classification_metrics(
        y_diabetes_te,
        diabetes_scratch_pred,
        diabetes_scratch_prob,
    )

    diabetes_scratch_metrics["parameters"] = _
    ↪count_numpy_parameters(diabetes_scratch_params)
    diabetes_scratch_metrics["training_time_seconds"] = diabetes_scratch_time
    diabetes_scratch_metrics["epochs"] = EPOCHS

    print(pd.Series(diabetes_scratch_metrics))

    show_confusion_matrix(
        y_diabetes_te,
        diabetes_scratch_pred,
        [str(v) for v in sorted(np.unique(y_diabetes_te))],
        "Diabetes - Scratch MLP confusion matrix",
    )

    diabetes_scratch_report = show_classification_report(
```

```

y_diabetes_te,
diabetes_scratch_pred,
[str(v) for v in sorted(np.unique(y_diabetes_te))],
"Diabetes - Scratch MLP per-class metrics",
)


show_confidence_distribution(
    diabetes_scratch_prob,
    y_diabetes_te == diabetes_scratch_pred,
    "Diabetes - Scratch MLP confidence",
)

```



```

accuracy                0.644513
precision_macro          0.436117
recall_macro            0.525018
f1_macro                0.422389
precision_weighted       0.849747
recall_weighted          0.644513
f1_weighted             0.709948
mean_confidence          0.589493
parameters              3587.000000
training_time_seconds    2.660124
epochs                  5.000000
dtype: float64

```



<Figure size 1000x600 with 0 Axes>



5.6 25. Diabetes with PyTorch

The architecture is kept equivalent to the Scratch MLP:

Linear(64) -> ReLU -> Linear(32) -> ReLU -> Linear(3)

The main difference is who performs the low-level work:

- Scratch: the notebook implements the formulas and gradients.
- PyTorch: autograd calculates gradients.

```
[121]: class SimpleMLPPyTorch(nn.Module):
        def __init__(self, input_dim, num_classes):
            super().__init__()

            self.network = nn.Sequential(
                nn.Linear(input_dim, 64),
                nn.ReLU(),
                nn.Linear(64, 32),
                nn.ReLU(),
                nn.Linear(32, num_classes),
            )

        def forward(self, x):
            return self.network(x)

def train_torch_mlp(
    X_train, y_train,
    X_val, y_val,
    input_dim,
    num_classes,
    class_weights,
    epochs=5,
    batch_size=128,
    lr=1e-3,
    seed=42,
):
    set_global_seed(seed)

    model = SimpleMLPPyTorch(
        input_dim,
        num_classes,
    ).to(DEVICE)

    train_ds = TensorDataset(
        torch.from_numpy(X_train),
        torch.from_numpy(y_train),
    )
    val_ds = TensorDataset(
        torch.from_numpy(X_val),
        torch.from_numpy(y_val),
```

```

)

train_loader = DataLoader(
    train_ds,
    batch_size=batch_size,
    shuffle=True,
)

val_loader = DataLoader(
    val_ds,
    batch_size=batch_size,
    shuffle=False,
)

weight_tensor = torch.tensor(
    class_weights,
    dtype=torch.float32,
    device=DEVICE,
)

criterion = nn.CrossEntropyLoss(weight=weight_tensor)
optimizer = torch.optim.Adam(
    model.parameters(),
    lr=lr,
)

history = {
    "loss": [],
    "accuracy": [],
    "val_loss": [],
    "val_accuracy": [],
}

best_state = None
best_val_loss = float("inf")

start_time = time.perf_counter()

for epoch in range(epochs):
    model.train()

    train_loss_sum = 0.0
    train_correct = 0
    train_seen = 0

    for features, labels in train_loader:
        features = features.to(DEVICE)
        labels = labels.to(DEVICE)

```

```

optimizer.zero_grad()

outputs = model(features)
loss = criterion(outputs, labels)

loss.backward()
optimizer.step()

train_loss_sum += loss.item() * len(features)
train_correct += int((outputs.argmax(dim=1) == labels).sum().item())
train_seen += len(features)

model.eval()

val_loss_sum = 0.0
val_correct = 0
val_seen = 0

with torch.no_grad():
    for features, labels in val_loader:
        features = features.to(DEVICE)
        labels = labels.to(DEVICE)

        outputs = model(features)
        loss = criterion(outputs, labels)

        val_loss_sum += loss.item() * len(features)
        val_correct += int((outputs.argmax(dim=1) == labels).sum()).
↪item())
        val_seen += len(features)

train_loss = train_loss_sum / train_seen
train_acc = train_correct / train_seen
val_loss = val_loss_sum / val_seen
val_acc = val_correct / val_seen

history["loss"].append(float(train_loss))
history["accuracy"].append(float(train_acc))
history["val_loss"].append(float(val_loss))
history["val_accuracy"].append(float(val_acc))

print(
    f"Epoch {epoch + 1}/{epochs} | "
    f"loss={train_loss:.4f} | "
    f"acc={train_acc:.4f} | "
    f"val_loss={val_loss:.4f} | "

```

```

        f"val_acc={val_acc:.4f}"
    )

    if val_loss < best_val_loss:
        best_val_loss = val_loss
        best_state = {
            key: value.detach().cpu().clone()
            for key, value in model.state_dict().items()
        }

    elapsed = time.perf_counter() - start_time
    model.load_state_dict(best_state)

    return model, history, elapsed

```

```

[122]: diabetes_torch_model, diabetes_torch_history, diabetes_torch_time =
↳ train_torch_mlp(
    X_diabetes_tr,
    y_diabetes_tr,
    X_diabetes_val,
    y_diabetes_val,
    input_dim=X_diabetes_tr.shape[1],
    num_classes=diabetes_num_classes,
    class_weights=diabetes_class_weights,
    epochs=EPOCHS,
    batch_size=BATCH_SIZE,
    lr=LEARNING_RATE,
    seed=SEED,
)



plot_training_history(
    diabetes_torch_history,
    title="Diabetes PyTorch MLP",
)

```

```

Epoch 1/5 | loss=0.9513 | acc=0.6360 | val_loss=0.9443 | val_acc=0.6926
Epoch 2/5 | loss=0.9388 | acc=0.6558 | val_loss=0.9433 | val_acc=0.6257
Epoch 3/5 | loss=0.9339 | acc=0.6508 | val_loss=0.9411 | val_acc=0.6488
Epoch 4/5 | loss=0.9328 | acc=0.6501 | val_loss=0.9436 | val_acc=0.6791
Epoch 5/5 | loss=0.9307 | acc=0.6464 | val_loss=0.9411 | val_acc=0.6404

```



```
[123]: def predict_torch_mlp(model, X, batch_size=256):  
        loader = DataLoader(  
            TensorDataset(torch.from_numpy(X)),  
            batch_size=batch_size,  
            shuffle=False,  
        )
```

```

model.eval()
all_prob = []

with torch.no_grad():
    for (features,) in loader:
        features = features.to(DEVICE)
        logits = model(features)
        probabilities = torch.softmax(logits, dim=1)
        all_prob.append(probabilities.cpu().numpy())

probabilities = np.concatenate(all_prob, axis=0)
predictions = np.argmax(probabilities, axis=1)

return predictions, probabilities

diabetes_torch_pred, diabetes_torch_prob = predict_torch_mlp(
    diabetes_torch_model,
    X_diabetes_te,
)

diabetes_torch_metrics = compute_classification_metrics(
    y_diabetes_te,
    diabetes_torch_pred,
    diabetes_torch_prob,
)

diabetes_torch_metrics["parameters"] = int(
    sum(p.numel() for p in diabetes_torch_model.parameters())
)

diabetes_torch_metrics["training_time_seconds"] = diabetes_torch_time
diabetes_torch_metrics["epochs"] = EPOCHS

print(pd.Series(diabetes_torch_metrics))

show_confusion_matrix(
    y_diabetes_te,
    diabetes_torch_pred,
    [str(v) for v in sorted(np.unique(y_diabetes_te))],
    "Diabetes - PyTorch MLP confusion matrix",
)

show_classification_report(
    y_diabetes_te,
    diabetes_torch_pred,
    [str(v) for v in sorted(np.unique(y_diabetes_te))],

```

```

    "Diabetes - PyTorch MLP per-class metrics",
)


show_confidence_distribution(
    diabetes_torch_prob,
    y_diabetes_te == diabetes_torch_pred,
    "Diabetes - PyTorch MLP confidence",
)

```



```

accuracy                0.639999
precision_macro         0.442533
recall_macro           0.527373
f1_macro               0.425242
precision_weighted      0.852344
recall_weighted        0.639999
f1_weighted            0.712216
mean_confidence        0.591627
parameters             3587.000000
training_time_seconds   31.134991
epochs                 5.000000
dtype: float64

```



<Figure size 1000x600 with 0 Axes>



5.7 26. Diabetes with Keras

Keras gives the same MLP structure with a higher abstraction:

Dense -> ReLU -> Dense -> ReLU -> Dense

The training loop is compressed into `model.fit()`.

```
[124]: def build_keras_mlp(input_dim, num_classes):
    model = keras.Sequential([
        layers.Input(shape=(input_dim,)),
        layers.Dense(64, activation="relu"),
        layers.Dense(32, activation="relu"),
        layers.Dense(num_classes),
    ])

    return model

def train_keras_mlp(
    X_train, y_train,
    X_val, y_val,
    input_dim,
    num_classes,
    class_weights,
    epochs=5,
    batch_size=128,
    lr=1e-3,
    seed=42,
):
    set_global_seed(seed)

    model = build_keras_mlp(
        input_dim,
        num_classes,
    )

    model.compile(
        optimizer=keras.optimizers.Adam(learning_rate=lr),
        loss=keras.losses.SparseCategoricalCrossentropy(from_logits=True),
        metrics=["accuracy"],
    )

    class_weight_dict = {
        int(cls): float(weight)
        for cls, weight in zip(
            np.unique(y_train),
            class_weights,
        )
    }

    start_time = time.perf_counter()
```

```

history_obj = model.fit(
    X_train,
    y_train,
    validation_data=(X_val, y_val),
    class_weight=class_weight_dict,
    epochs=epochs,
    batch_size=batch_size,
    verbose=1,
)

elapsed = time.perf_counter() - start_time

history = {
    "loss": history_obj.history["loss"],
    "accuracy": history_obj.history["accuracy"],
    "val_loss": history_obj.history["val_loss"],
    "val_accuracy": history_obj.history["val_accuracy"],
}

return model, history, elapsed

```

```

[125]: diabetes_keras_model, diabetes_keras_history, diabetes_keras_time =
↳ train_keras_mlp(
    X_diabetes_tr,
    y_diabetes_tr,
    X_diabetes_val,
    y_diabetes_val,
    input_dim=X_diabetes_tr.shape[1],
    num_classes=diabetes_num_classes,
    class_weights=diabetes_class_weights,
    epochs=EPOCHS,
    batch_size=BATCH_SIZE,
    lr=LEARNING_RATE,
    seed=SEED,
)

plot_training_history(
    diabetes_keras_history,
    title="Diabetes Keras MLP",
)

```

```

Epoch 1/5
1427/1427          3s 2ms/step -
accuracy: 0.6153 - loss: 0.9626 - val_accuracy: 0.6162 - val_loss: 0.8807
Epoch 2/5
1427/1427          2s 1ms/step -
accuracy: 0.6105 - loss: 0.9450 - val_accuracy: 0.6175 - val_loss: 0.8710

```

Epoch 3/5

1427/1427

2s 2ms/step -

accuracy: 0.6101 - loss: 0.9406 - val_accuracy: 0.6130 - val_loss: 0.8693

Epoch 4/5

1427/1427

2s 2ms/step -



accuracy: 0.6087 - loss: 0.9366 - val_accuracy: 0.6102 - val_loss: 0.8720

Epoch 5/5

1427/1427

2s 2ms/step -

accuracy: 0.6086 - loss: 0.9335 - val_accuracy: 0.6093 - val_loss: 0.8666



```
[126]: def predict_keras_mlp(model, X):
        # Diabetes là dữ liệu tabular nên giữ nguyên dạng (N, features)
        logits = model.predict(
            X,
            verbose=0
        )

        probabilities = softmax(logits)
        predictions = np.argmax(probabilities, axis=1)

        return predictions, probabilities
```

```
[127]: diabetes_keras_pred, diabetes_keras_prob = predict_keras_mlp(
        diabetes_keras_model,
        X_diabetes_te,
    )

    diabetes_keras_metrics = compute_classification_metrics(
        y_diabetes_te,
        diabetes_keras_pred,
        diabetes_keras_prob,
    )

    diabetes_keras_metrics["parameters"] = int(
        diabetes_keras_model.count_params()
    )

    diabetes_keras_metrics["training_time_seconds"] = diabetes_keras_time
    diabetes_keras_metrics["epochs"] = EPOCHS

    print(pd.Series(diabetes_keras_metrics))


    show_confusion_matrix(
        y_diabetes_te,
        diabetes_keras_pred,
        [str(v) for v in sorted(np.unique(y_diabetes_te))],
        "Diabetes - Keras MLP confusion matrix",
    )

    show_classification_report(
        y_diabetes_te,
        diabetes_keras_pred,
        [str(v) for v in sorted(np.unique(y_diabetes_te))],
        "Diabetes - Keras MLP per-class metrics",
```



```
)

show_confidence_distribution(
    diabetes_keras_prob,
    y_diabetes_te == diabetes_keras_pred,
    "Diabetes - Keras MLP confidence",
)
```

```
accuracy                0.611440
precision_macro         0.440206
recall_macro           0.517610
f1_macro               0.412757
precision_weighted      0.854727
recall_weighted         0.611440
f1_weighted            0.692423
mean_confidence         0.591156
parameters              3587.000000
training_time_seconds   12.515781
epochs                  5.000000
dtype: float64
```



<Figure size 1000x600 with 0 Axes>



5.8 27. Diabetes visual analysis

For a tabular dataset, an image grid is not meaningful.

Instead, we inspect:

- confusion matrix;
- class-level precision, recall and F1;
- confidence;
- PCA space;
- target distribution.

This is an example of visualization matching the data representation.

```
[128]: diabetes_results = pd.DataFrame([
    {"dataset": "Diabetes", "framework": "Scratch", **diabetes_scratch_metrics},
    {"dataset": "Diabetes", "framework": "PyTorch", **diabetes_torch_metrics},
    {"dataset": "Diabetes", "framework": "Keras", **diabetes_keras_metrics},
])


diabetes_results = diabetes_results[
    [
        "dataset",
        "framework",
        "accuracy",
        "precision_macro",
        "recall_macro",
        "f1_macro",
        "mean_confidence",
        "parameters",
        "training_time_seconds",
        "epochs",
    ]
]

display(diabetes_results)

plt.figure(figsize=(10, 5))
plt.bar(diabetes_results["framework"], diabetes_results["f1_macro"])
plt.title("Diabetes macro F1 by implementation")
plt.xlabel("Implementation")
plt.ylabel("Macro F1")
plt.ylim(0, 1.05)
plt.tight_layout()
plt.show()
```

[illegible]

1	0.591627	3587	31.134991	5
2	0.591156	3587	12.515781	5



6 Part E - Final comparison and report

6.1 28. Framework comparison

The key lesson from Lecture 04 is:

The same learning idea can be implemented at different levels of abstraction.

The comparison is therefore about visibility and programming structure, not only about the score.

```
[129]: all_results = pd.concat(
    [
        mnist_results,
        cifar_results,
        diabetes_results,
    ],
    ignore_index=True,
)

display(all_results)
```

	dataset	framework	accuracy	precision_macro	recall_macro	f1_macro	\
0	MNIST	Scratch	0.979000	0.978987	0.978883	0.978831	
1	MNIST	PyTorch	0.975300	0.975397	0.975098	0.975205	
2	MNIST	Keras	0.978000	0.978049	0.977922	0.977931	
3	CIFAR-10	Scratch	0.572900	0.588441	0.572900	0.572666	
4	CIFAR-10	PyTorch	0.573700	0.576051	0.573700	0.564701	

5	CIFAR-10	Keras	0.568100	0.574029	0.568100	0.563815
6	Diabetes	Scratch	0.644513	0.436117	0.525018	0.422389
7	Diabetes	PyTorch	0.639999	0.442533	0.527373	0.425242
8	Diabetes	Keras	0.611440	0.440206	0.517610	0.412757

	mean_confidence	parameters	training_time_seconds	epochs
0	0.974335	31530	139.493802	5
1	0.969117	31530	26.547899	5
2	0.970449	31530	19.982689	5
3	0.546630	41418	188.094851	5
4	0.561211	41418	45.944586	5
5	0.548039	41418	19.185502	5
6	0.589493	3587	2.660124	5
7	0.591627	3587	31.134991	5
8	0.591156	3587	12.515781	5

6.2 28.1 What remains the same?

Across Scratch, PyTorch and Keras, the following concepts remain:

1. Data representation.
2. The sequence of transformations.
3. The loss function idea.
4. Gradient-based optimization.
5. Evaluation with predictions and metrics.

What changes is the amount of implementation detail visible to the programmer.

6.2.1 Scratch

Most details are explicit:

- convolution arithmetic;
- pooling arithmetic;
- parameter arrays;
- forward pass;
- backward pass;
- parameter update.

6.2.2 PyTorch

The model structure is explicit through modules and `forward()`, while autograd handles the derivative calculations.

6.2.3 Keras

The model definition is compact and the training process is largely hidden inside `compile()` and `fit()`.

6.3 29. Performance comparison

The report should not stop at accuracy.

The lecture asks us to record:

- Number of parameters.
- Training time.
- Number of epochs.
- Training loss.
- Validation/test accuracy.
- Precision.
- Recall.
- F1.
- Confusion matrix.

A fair comparison uses the same:


- dataset;
- split;
- architecture;
- comparable hyperparameters.

The notebook follows that principle by reusing the same selected train/validation indices for the three implementations on each dataset.

```
[130]: # Accuracy comparison across datasets.
plt.figure(figsize=(12, 6))

for dataset_name in all_results["dataset"].unique():
    subset = all_results[all_results["dataset"] == dataset_name]
    plt.plot(
        subset["framework"],
        subset["accuracy"],
        marker="o",
        label=dataset_name,
    )


plt.xlabel("Implementation")
plt.ylabel("Accuracy")
plt.ylim(0, 1.05)
plt.title("Accuracy comparison across datasets")
plt.legend()
plt.tight_layout()
plt.show()
```



```
[131]: # Macro F1 comparison.
plt.figure(figsize=(12, 6))

for dataset_name in all_results["dataset"].unique():
    subset = all_results[all_results["dataset"] == dataset_name]
    plt.plot(
        subset["framework"],
        subset["f1_macro"],
        marker="o",
        label=dataset_name,
    )

plt.xlabel("Implementation")
plt.ylabel("Macro F1")
plt.ylim(0, 1.05)
plt.title("Macro F1 comparison across datasets")
plt.legend()
plt.tight_layout()
plt.show()
```




```
[132]: # Parameter count comparison for the image datasets.
image_results = all_results[
    all_results["dataset"].isin(["MNIST", "CIFAR-10"])
]

pivot_params = image_results.pivot(
    index="framework",
    columns="dataset",
    values="parameters",
)

display(pivot_params)

pivot_params.plot(kind="bar", figsize=(10, 6))
plt.title("Parameter count by framework and image dataset")
plt.xlabel("Framework")
plt.ylabel("Number of parameters")
plt.xticks(rotation=0)
plt.tight_layout()
plt.show()
```

dataset	CIFAR-10	MNIST
framework		
Keras	41418	31530
PyTorch	41418	31530
Scratch	41418	31530




```
[133]: # Training time comparison.
pivot_time = all_results.pivot(
    index="framework",
    columns="dataset",
    values="training_time_seconds",
)

display(pivot_time)

pivot_time.plot(kind="bar", figsize=(10, 6))
plt.title("Training time by framework and dataset")
plt.xlabel("Framework")
plt.ylabel("Training time in seconds")
plt.xticks(rotation=0)
plt.tight_layout()
plt.show()
```

dataset	CIFAR-10	Diabetes	MNIST
framework			
Keras	19.185502	12.515781	19.982689
PyTorch	45.944586	31.134991	26.547899
Scratch	188.094851	2.660124	139.493802



6.4 30. Answers to the important questions from Lecture 04

6.4.1 Question 1. What is the same, and what is different?

The mathematical learning idea is the same:

Input -> transformations -> loss -> gradient -> parameter update

The syntax and abstraction level differ.

Scratch exposes the operations. Keras packages them into high-level layers and training utilities. PyTorch exposes modules and the training loop while providing automatic differentiation.

6.4.2 Question 2. What computational work is hidden inside Keras fit()?

Conceptually, `fit()` performs repeated:

```
Batch loading
-> Forward pass
-> Loss calculation
-> Gradient calculation
-> Parameter update
-> Metric calculation
```

The user does not have to write each of these operations explicitly.

6.4.3 Question 3. Which implementation hides the most details?

Keras hides more low-level implementation details in the student-facing code, especially the training loop and differentiation.

6.4.4 Question 4. Which implementation gives the programmer the most direct control?

Scratch exposes the most operations directly because the programmer writes the mathematical transformations and training mechanics.

PyTorch sits between Scratch and Keras in the notebook's comparison:

Scratch -> maximum visibility

PyTorch -> modular control plus automatic differentiation

Keras -> high-level convenience

6.4.5 Question 5. Why must tensor shapes be calculated before the Dense or Linear layer?

Because the classifier's input dimension must exactly match the flattened feature size.

For MNIST:

B x 1 x 28 x 28

-> Conv

B x 16 x 28 x 28

-> Pool

B x 16 x 14 x 14

-> Flatten

B x 3136

Therefore the final linear layer receives 3136 features.

For CIFAR-10:

B x 3 x 32 x 32

-> Conv

B x 16 x 32 x 32

-> Pool

B x 16 x 16 x 16

-> Flatten

B x 4096

Therefore the final linear layer receives 4096 features.

6.4.6 Question 6. Why is Diabetes modeled with MLP instead of CNN?

Because Diabetes is tabular data represented as a feature vector, not a 2D image.

Lecture 04 explicitly presents:

Diabetes -> MLP

MNIST -> CNN

The model should match the representation of the data.

6.4.7 Question 7. Why should accuracy not be the only metric?

A single accuracy number can hide class-specific problems, especially with imbalanced data.

Precision, recall and F1 show class-level behavior. The confusion matrix shows which classes are being confused. Confidence visualization shows whether errors are made with high or low confidence.

6.5 31. Save trained models and website-ready outputs

The later web demo should not have to retrain a model.

This section saves:

- Scratch parameters as `.joblib`.
- PyTorch weights as `.pth`.
- Keras models as `.keras`.
- Metric summaries as JSON/CSV.
- Test predictions as CSV.

The output directory is therefore also the bridge from the Jupyter notebook to the future web application.

```
[134]: import joblib

# -----
# Save Scratch models
# -----
joblib.dump(
    mnist_scratch_params,
    MODELS_DIR / "mnist_cnn_scratch.joblib",
)

joblib.dump(
    cifar_scratch_params,
    MODELS_DIR / "cifar10_cnn_scratch.joblib",
)

joblib.dump(
    diabetes_scratch_params,
    MODELS_DIR / "diabetes_mlp_scratch.joblib",
)

# Save preprocessing object for Diabetes.
joblib.dump(
    diabetes_scaler,
    MODELS_DIR / "diabetes_standard_scaler.joblib",
)

# -----
# Save PyTorch models
# -----
torch.save(
```

```

    mnist_torch_model.state_dict(),
    MODELS_DIR / "mnist_cnn_pytorch.pth",
)

torch.save(
    cifar_torch_model.state_dict(),
    MODELS_DIR / "cifar10_cnn_pytorch.pth",
)

torch.save(
    diabetes_torch_model.state_dict(),
    MODELS_DIR / "diabetes_mlp_pytorch.pth",
)

# -----
# Save Keras models
# -----
mnist_keras_model.save(
    MODELS_DIR / "mnist_cnn_keras.keras"
)

cifar_keras_model.save(
    MODELS_DIR / "cifar10_cnn_keras.keras"
)

diabetes_keras_model.save(
    MODELS_DIR / "diabetes_mlp_keras.keras"
)

print("Models saved to:", MODELS_DIR)

```

Models saved to: C:\Users\ADMIN\Downloads\Intelligent System\assignment
4\DATA\models

```

[135]: # Save metrics.
save_metrics(
    mnist_scratch_metrics,
    METRICS_DIR / "mnist_scratch_metrics.json",
)
save_metrics(
    mnist_torch_metrics,
    METRICS_DIR / "mnist_pytorch_metrics.json",
)
save_metrics(
    mnist_keras_metrics,
    METRICS_DIR / "mnist_keras_metrics.json",
)

```

```

save_metrics(
    cifar_scratch_metrics,
    METRICS_DIR / "cifar10_scratch_metrics.json",
)
save_metrics(
    cifar_torch_metrics,
    METRICS_DIR / "cifar10_pytorch_metrics.json",
)
save_metrics(
    cifar_keras_metrics,
    METRICS_DIR / "cifar10_keras_metrics.json",
)

save_metrics(
    diabetes_scratch_metrics,
    METRICS_DIR / "diabetes_scratch_metrics.json",
)
save_metrics(
    diabetes_torch_metrics,
    METRICS_DIR / "diabetes_pytorch_metrics.json",
)
save_metrics(
    diabetes_keras_metrics,
    METRICS_DIR / "diabetes_keras_metrics.json",
)

all_results.to_csv(
    METRICS_DIR / "all_model_comparison.csv",
    index=False,
)

print("Metrics saved to:", METRICS_DIR)

```

Metrics saved to: C:\Users\ADMIN\Downloads\Intelligent System\assignment 4\DATA\outputs\metrics

[136]: *# Save website-friendly prediction tables.*

```

mnist_demo_df = pd.DataFrame({
    "true_label": mnist_y_te,
    "scratch_prediction": mnist_scratch_pred,
    "scratch_confidence": np.max(mnist_scratch_prob, axis=1),
    "pytorch_prediction": mnist_torch_pred,
    "pytorch_confidence": np.max(mnist_torch_prob, axis=1),
    "keras_prediction": mnist_keras_pred,
    "keras_confidence": np.max(mnist_keras_prob, axis=1),
})

```

```

})

cifar_demo_df = pd.DataFrame({
    "true_label": cifar_y_te,
    "scratch_prediction": cifar_scratch_pred,
    "scratch_confidence": np.max(cifar_scratch_prob, axis=1),
    "pytorch_prediction": cifar_torch_pred,
    "pytorch_confidence": np.max(cifar_torch_prob, axis=1),
    "keras_prediction": cifar_keras_pred,
    "keras_confidence": np.max(cifar_keras_prob, axis=1),
})

diabetes_demo_df = pd.DataFrame({
    "true_label": y_diabetes_te,
    "scratch_prediction": diabetes_scratch_pred,
    "scratch_confidence": np.max(diabetes_scratch_prob, axis=1),
    "pytorch_prediction": diabetes_torch_pred,
    "pytorch_confidence": np.max(diabetes_torch_prob, axis=1),
    "keras_prediction": diabetes_keras_pred,
    "keras_confidence": np.max(diabetes_keras_prob, axis=1),
})

mnist_demo_df.to_csv(
    PREDICTIONS_DIR / "mnist_predictions.csv",
    index=False,
)

cifar_demo_df.to_csv(
    PREDICTIONS_DIR / "cifar10_predictions.csv",
    index=False,
)

diabetes_demo_df.to_csv(
    PREDICTIONS_DIR / "diabetes_predictions.csv",
    index=False,
)

print("Prediction tables saved to:", PREDICTIONS_DIR)

```

Prediction tables saved to: C:\Users\ADMIN\Downloads\Intelligent System\assignment 4\DATA\outputs\predictions

6.6 31.1 A simple notebook demo

The future website can follow the same idea as this notebook demo:

Input image

-> select model

-> preprocess

-> model prediction
-> class probabilities
-> confidence
-> visualization

The web application should eventually expose:

- predicted class;
- probability for every class;
- confidence;
- model name;
- training metrics;
- confusion matrix and charts;
- example predictions;
- model comparison.

That design is consistent with the current notebook outputs and avoids retraining during a demo.

```
[137]: def demo_mnist_prediction(index=0, model_name="PyTorch"):
        image = mnist_x_test[index:index + 1]

        if model_name.lower() == "scratch":
            pred, prob = predict_scratch_cnn(image, mnist_scratch_params)

        elif model_name.lower() == "pytorch":
            pred, prob = predict_torch(mnist_torch_model, image)

        elif model_name.lower() == "keras":
            pred, prob = predict_keras(mnist_keras_model, image)

        else:
            raise ValueError("model_name must be Scratch, PyTorch or Keras")

        predicted_class = int(pred[0])
        confidence = float(prob[0, predicted_class])

        plt.figure(figsize=(5, 5))
        plt.imshow(image[0, 0], cmap="gray")
        plt.title(
            f"Model: {model_name}\n"
            f"Prediction: {predicted_class}\n"
            f"Confidence: {confidence:.3f}"
        )
        plt.axis("off")
        plt.tight_layout()
        plt.show()

        probability_table = pd.DataFrame({
            "class": IMAGE_CLASS_NAMES["mnist"],
```

```

        "probability": prob[0],
    }).sort_values(
        "probability",
        ascending=False,
    )

    display(probability_table)

    return predicted_class, probability_table

# Example:
# demo_mnist_prediction(index=0, model_name="PyTorch")

```

7 32. Final conclusions

7.1 32.1 Conceptual conclusion

The notebook follows the central conceptual map:

Mathematical function

- > Neuron
- > Nonlinear neuron
- > Layer
- > Function composition
- > CNN
- > Training
- > Evaluation
- > Modern architectural ideas

The implementation comparison follows:

Understand

- > Implement
- > Abstract
- > Compare

7.2 32.2 CNN conclusion

For image data, the key CNN ideas demonstrated are:

- Local connectivity.
- Weight sharing.
- Convolution.
- Nonlinearity.
- Spatial downsampling.
- Shape tracking.
- Classification.
- Gradient-based training.

7.3 32.3 Framework conclusion

The three implementations are not three different mathematical learning systems.

They are three programming abstractions for related computations:

Scratch

-> explicit mathematical operations

Keras

-> high-level layer and training API

PyTorch

-> modular layers and explicit forward/training flow

7.4 32.4 Evaluation conclusion

A serious model report should show more than accuracy.

This notebook therefore records:

- accuracy;
- precision;
- recall;
- macro F1;
- weighted F1;
- confidence;
- parameter count;
- training time;
- training history;
- confusion matrix;
- incorrect predictions.

7.5 32.5 Demo conclusion

The saved models and output tables are deliberately structured so that the later web application can load the trained artifacts instead of retraining them.

AI Vision Lab
Assignment 04

Dashboard

MNIST

CIFAR-10

Diabetes

Comparison

CORE IDEA
Understand the data, model, functions and framework before comparing implementations.

DEEP LEARNING AS FUNCTION COMPOSITION

Models ready

AI Model Demonstration

Prediction, probability and evaluation for the Assignment 04 models.

3 DATASETS

One project, multiple representations

MNIST and CIFAR-10 use CNNs. Diabetes uses an MLP because it is tabular data. The demo loads the trained artifacts created by the Jupyter Notebook.

Input → Features → Classifier → Probability

28 MNIST
Grayscale digit images →

32 CIFAR-10
RGB image classification →

MLP Diabetes
Tabular classification →

RESULTS

Model overview

DATASET	FRAMEWORK	ACCURACY	MACRO F1	PARAMETERS	TIME
MNIST	Scratch	97.9%	97.9%	31,530	139.49 s

AI Vision Lab
Assignment 04

Dashboard

MNIST

CIFAR-10

Diabetes

Comparison

CORE IDEA
Understand the data, model, functions and framework before comparing implementations.

28 MNIST
Grayscale digit images →

32 CIFAR-10
RGB image classification →

MLP Diabetes
Tabular classification →

RESULTS

Model overview

DATASET	FRAMEWORK	ACCURACY	MACRO F1	PARAMETERS	TIME
MNIST	Scratch	97.9%	97.9%	31,530	139.49 s
MNIST	PyTorch	97.5%	97.5%	31,530	26.55 s
MNIST	Keras	97.8%	97.8%	31,530	19.98 s
CIFAR-10	Scratch	57.3%	57.3%	41,418	188.09 s
CIFAR-10	PyTorch	57.4%	56.5%	41,418	45.94 s
CIFAR-10	Keras	56.8%	56.4%	41,418	19.19 s
Diabetes	Scratch	64.5%	42.2%	3,587	2.66 s
Diabetes	PyTorch	64.0%	42.5%	3,587	31.13 s
Diabetes	Keras	61.1%	41.3%	3,587	12.52 s

AI Vision Lab
Assignment 04

Dashboard

MNIST

CIFAR-10

Diabetes

Comparison

CORE IDEA
Understand the data, model, functions and framework before comparing implementations.

DEEP LEARNING AS FUNCTION COMPOSITION

MNIST Demo

CNN classification with grayscale handwritten digits.

IMAGE CLASSIFICATION

MNIST

Input: $1 \times 28 \times 28$. Ten digit classes.

Prediction demo

PyTorch

+

Choose a digit image

Prefer a white digit on a dark background to match MNIST.



Scratch

2

Confidence: 100.0%

2	100.0%
6	0.0%
1	0.0%
3	0.0%
8	0.0%
0	0.0%
5	0.0%
9	0.0%
4	0.0%
7	0.0%

Run prediction

143

AI Vision Lab
Assignment 04

Dashboard

MNIST

CIFAR-10

Diabetes

Comparison

CORE IDEA
Understand the data, model, functions and framework before comparing implementations.

IMAGE CLASSIFICATION

CIFAR-10

Input: $3 \times 32 \times 32$. Ten image classes.

Prediction demo

PyTorch

+

Choose an RGB image

The demo resizes your image to 32×32 .

Run prediction

PyTorch

cat

Confidence: 61.5%

cat	61.5%
dog	23.6%
frog	12.2%
ship	1.2%
bird	0.8%
airplane	0.4%
deer	0.3%
horse	0.0%
truck	0.0%
automobile	0.0%

IMAGE CLASSIFICATION

CIFAR-10

Input: $3 \times 32 \times 32$. Ten image classes.

Prediction demo

Scratch

+

Choose an RGB image

The demo resizes your image to 32×32 .


Run prediction


Scratch



cat


Confidence: 50.8%

cat	50.8%
frog	18.6%
dog	16.4%
ship	8.9%
airplane	2.7%
bird	1.7%
deer	0.8%
automobile	0.1%
horse	0.0%
truck	0.0%









[]: